

Portfolio 2

January 17, 2026

HEALTHCARE PROVIDER FRAUD DETECTION ANALYSIS

<https://www.kaggle.com/datasets/rohitrox/healthcare-provider-fraud-detection-analysis>

Dataset Overview The dataset contains healthcare data related to beneficiaries, inpatient claims, and outpatient claims.

It is intended for fraud detection analysis.

The main files include:

Beneficiary Data – demographic info, chronic diseases, date of death, gender, etc.

Inpatient Data – hospital visits, admission dates, discharge dates, procedures, charges.

Outpatient Data – clinic visits, procedures, medications, charges.

```
[4]: import numpy as np  
import pandas as pd  
import matplotlib.pyplot as plt  
import seaborn as sns
```

```
[5]: df1 = pd.read_csv('D:\Hira\Project 2\Portfolio 2\Train_Beneficiarydata.csv')
```

```
[6]: df1
```

```
[6]:      BeneID        DOB     DOD  Gender  Race RenalDiseaseIndicator \
0    BENE11001  1943-01-01   NaN      1      1                  0
1    BENE11002  1936-09-01   NaN      2      1                  0
2    BENE11003  1936-08-01   NaN      1      1                  0
3    BENE11004  1922-07-01   NaN      1      1                  0
4    BENE11005  1935-09-01   NaN      1      1                  0
...       ...       ...  ...  ...  ...
138551  BENE159194  1939-07-01   NaN      1      1                  0
138552  BENE159195  1938-12-01   NaN      2      1                  0
138553  BENE159196  1916-06-01   NaN      2      1                  0
138554  BENE159197  1930-01-01   NaN      1      1                  0
138555  BENE159198  1952-04-01   NaN      2      1                  0
```

	State	County	NoOfMonths_PartACov	NoOfMonths_PartBCov	...	\
0	39	230	12	12	...	
1	39	280	12	12	...	
2	52	590	12	12	...	
3	39	270	12	12	...	
4	24	680	12	12	...	
...	
138551	39	140	12	12	...	
138552	49	530	12	12	...	
138553	6	150	12	12	...	
138554	16	560	12	12	...	
138555	21	20	12	12	...	
ChronicCond_Depression ChronicCond_Diabetes \						
0		1	1			
1		2	2			
2		2	2			
3		2	1			
4		2	1			
...	
138551		2	2			
138552		2	1			
138553		1	1			
138554		2	2			
138555		1	1			
ChronicCond_IschemicHeart ChronicCond_Osteoporosis \						
0		1	2			
1		2	2			
2		1	2			
3		1	1			
4		2	2			
...	
138551		2	2			
138552		2	2			
138553		1	2			
138554		1	2			
138555		2	2			
ChronicCond_rheumatoidarthritis ChronicCond_stroke \						
0		1	1			
1		2	2			
2		2	2			
3		1	2			
4		2	2			
...	
138551		2	2			

138552	2	2
138553	2	2
138554	2	2
138555	1	2
		\
0	36000	3204
1	0	0
2	0	0
3	0	0
4	0	0
...
138551	0	0
138552	0	0
138553	2000	1068
138554	0	0
138555	0	0
0	OPAnnualReimbursementAmt	OPAnnualDeductibleAmt
1	60	70
2	30	50
3	90	40
4	1810	760
5	1790	1200
...
138551	430	460
138552	880	100
138553	3240	1390
138554	2650	10
138555	5470	1870

[138556 rows x 25 columns]

[7]: df1.shape

[7]: (138556, 25)

[8]: df1.columns

[8]: Index(['BeneID', 'DOB', 'DOD', 'Gender', 'Race', 'RenalDiseaseIndicator',
 'State', 'County', 'NoOfMonths_PartACov', 'NoOfMonths_PartBCov',
 'ChronicCond_Alzheimer', 'ChronicCond_Heartfailure',
 'ChronicCond_KidneyDisease', 'ChronicCond_Cancer',
 'ChronicCond_ObstrPulmonary', 'ChronicCond_Depression',
 'ChronicCond_Diabetes', 'ChronicCond_IschemicHeart',
 'ChronicCond_Osteoporasis', 'ChronicCond_rheumatoidarthritis',
 'ChronicCond_stroke', 'IPAnnualReimbursementAmt',

```
'IPAnnualDeductibleAmt', 'OPAnnualReimbursementAmt',
'OPAnnualDeductibleAmt'],
dtype='object')
```

[9]: df1.info

	BeneID	DOB	DOD	Gender
0	1	1		0
1	2	1		0
2	1	1		0
3	1	1		0
4	1	1		0
...	
138551	1	1		0
138552	2	1		0
138553	2	1		0
138554	1	1		0
138555	2	1		0
...	
0	39	230	12	12 ...
1	39	280	12	12 ...
2	52	590	12	12 ...
3	39	270	12	12 ...
4	24	680	12	12 ...
...
138551	39	140	12	12 ...
138552	49	530	12	12 ...
138553	6	150	12	12 ...
138554	16	560	12	12 ...
138555	21	20	12	12 ...
...	
0	ChronicCond_Depression	ChronicCond_Diabetes	...	
1	1	1		
2	2	2		
3	2	1		
4	2	1		
...	
138551	2	2		
138552	2	1		
138553	1	1		
138554	2	2		
138555	1	1		
...	
0	ChronicCond_IschemicHeart	ChronicCond_Osteoporasis	...	

0	1	2	
1	2	2	
2	1	2	
3	1	1	
4	2	2	
...	
138551	2	2	
138552	2	2	
138553	1	2	
138554	1	2	
138555	2	2	
\\			
0	1	1	
1	2	2	
2	2	2	
3	1	2	
4	2	2	
...	
138551	2	2	
138552	2	2	
138553	2	2	
138554	2	2	
138555	1	2	
\\			
0	36000	3204	
1	0	0	
2	0	0	
3	0	0	
4	0	0	
...	
138551	0	0	
138552	0	0	
138553	2000	1068	
138554	0	0	
138555	0	0	
\\			
0	60	70	
1	30	50	
2	90	40	
3	1810	760	
4	1790	1200	
...	
138551	430	460	
138552	880	100	

138553	3240	1390
138554	2650	10
138555	5470	1870

[138556 rows x 25 columns]>

[10]: df1.describe()

	Gender	Race	State	County	\
count	138556.000000	138556.000000	138556.000000	138556.000000	
mean	1.570932	1.254511	25.666734	374.424745	
std	0.494945	0.717007	15.223443	266.277581	
min	1.000000	1.000000	1.000000	0.000000	
25%	1.000000	1.000000	11.000000	141.000000	
50%	2.000000	1.000000	25.000000	340.000000	
75%	2.000000	1.000000	39.000000	570.000000	
max	2.000000	5.000000	54.000000	999.000000	
	NoOfMonths_PartACov	NoOfMonths_PartBCov	ChronicCond_Alzheimer	\	
count	138556.000000	138556.000000	138556.000000		
mean	11.907727	11.910145	1.667817		
std	1.032332	0.936893	0.470998		
min	0.000000	0.000000	1.000000		
25%	12.000000	12.000000	1.000000		
50%	12.000000	12.000000	2.000000		
75%	12.000000	12.000000	2.000000		
max	12.000000	12.000000	2.000000		
	ChronicCond_Heartfailure	ChronicCond_KidneyDisease	\		
count	138556.000000	138556.000000			
mean	1.506322	1.687643			
std	0.499962	0.463456			
min	1.000000	1.000000			
25%	1.000000	1.000000			
50%	2.000000	2.000000			
75%	2.000000	2.000000			
max	2.000000	2.000000			
	ChronicCond_Cancer	...	ChronicCond_Depression	ChronicCond_Diabetes	\
count	138556.000000	...	138556.000000	138556.000000	
mean	1.880041	...	1.644476	1.398142	
std	0.324914	...	0.478674	0.489517	
min	1.000000	...	1.000000	1.000000	
25%	2.000000	...	1.000000	1.000000	
50%	2.000000	...	2.000000	1.000000	
75%	2.000000	...	2.000000	2.000000	
max	2.000000	...	2.000000	2.000000	

	ChronicCond_IschemicHeart	ChronicCond_Osteoporasis	\
count	138556.000000	138556.000000	
mean	1.324143	1.725317	
std	0.468056	0.446356	
min	1.000000	1.000000	
25%	1.000000	1.000000	
50%	1.000000	2.000000	
75%	2.000000	2.000000	
max	2.000000	2.000000	
	ChronicCond_rheumatoidarthritis	ChronicCond_stroke	\
count	138556.000000	138556.000000	
mean	1.743180	1.920942	
std	0.436881	0.269831	
min	1.000000	1.000000	
25%	1.000000	2.000000	
50%	2.000000	2.000000	
75%	2.000000	2.000000	
max	2.000000	2.000000	
	IPAnnualReimbursementAmt	IPAnnualDeductibleAmt	\
count	138556.000000	138556.000000	
mean	3660.346502	399.847296	
std	9568.621827	956.175202	
min	-8000.000000	0.000000	
25%	0.000000	0.000000	
50%	0.000000	0.000000	
75%	2280.000000	1068.000000	
max	161470.000000	38272.000000	
	OPAnnualReimbursementAmt	OPAnnualDeductibleAmt	
count	138556.000000	138556.000000	
mean	1298.219348	377.718258	
std	2493.901134	645.530187	
min	-70.000000	0.000000	
25%	170.000000	40.000000	
50%	570.000000	170.000000	
75%	1500.000000	460.000000	
max	102960.000000	13840.000000	

[8 rows x 21 columns]

[11]: `print(df1.isnull().sum())`

BeneID	0
DOB	0
DOD	137135

```

Gender          0
Race           0
RenalDiseaseIndicator 0
State          0
County         0
NoOfMonths_PartACov 0
NoOfMonths_PartBCov 0
ChronicCond_Alzheimer 0
ChronicCond_Heartfailure 0
ChronicCond_KidneyDisease 0
ChronicCond_Cancer 0
ChronicCond_ObstrPulmonary 0
ChronicCond_Depression 0
ChronicCond_Diabetes 0
ChronicCond_IschemicHeart 0
ChronicCond_Osteoporasis 0
ChronicCond_rheumatoidarthritis 0
ChronicCond_stroke 0
IPAnnualReimbursementAmt 0
IPAnnualDeductibleAmt 0
OPAnnualReimbursementAmt 0
OPAnnualDeductibleAmt 0
dtype: int64

```

0.0.1 Description

The DOD column contains NaN values.

In healthcare datasets, NaN in DOD usually means the beneficiary is still alive. ##### we need a categorical column that clearly shows Alive or Dead instead of NaN.

```
[12]: import pandas as pd

# Convert DOD to datetime
df1['DOD'] = pd.to_datetime(df1['DOD'], errors='coerce')

# Create Status column: Alive if NaN, Dead if date exists
df1['Status'] = df1['DOD'].apply(lambda x: 'Dead' if pd.notnull(x) else 'Alive')

# Check
print(df1['Status'].value_counts())
```

```

Status
Alive    137135
Dead      1421
Name: count, dtype: int64

```

Visualization Alive VS Dead

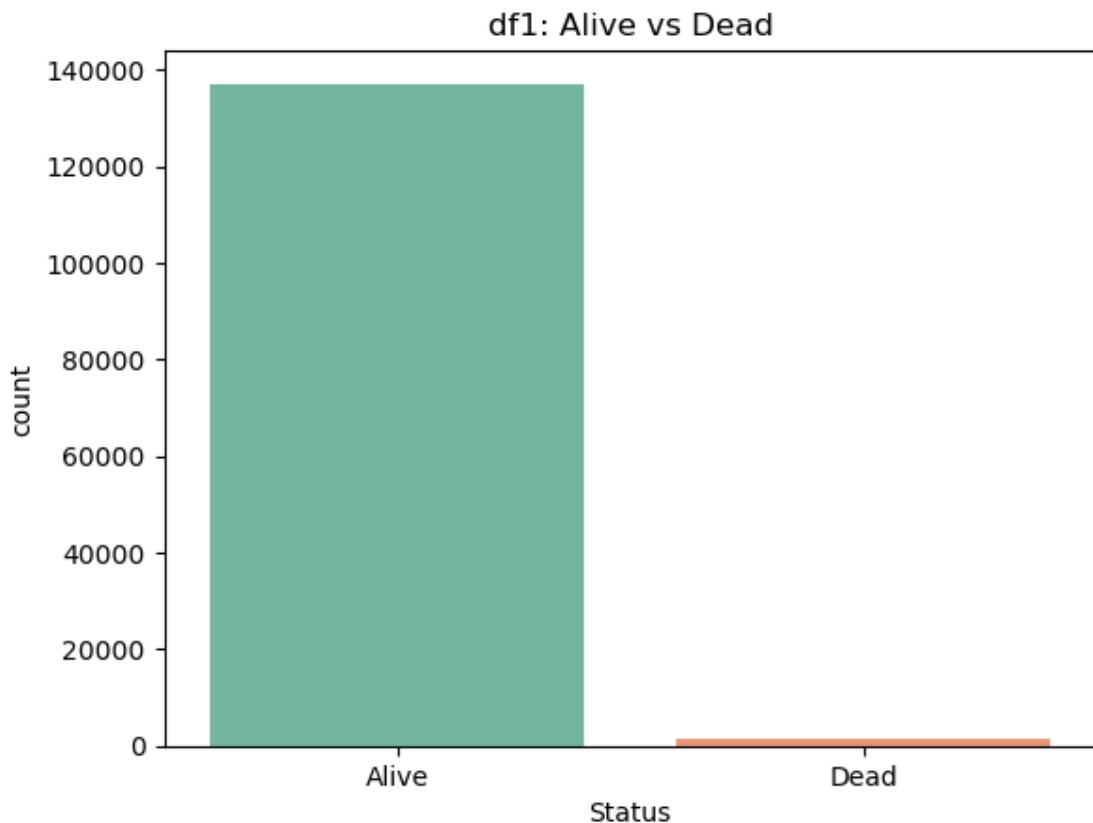
```
[13]: import seaborn as sns
import matplotlib.pyplot as plt

sns.countplot(x='Status', data=df1, palette='Set2')
plt.title("df1: Alive vs Dead")
plt.show()
```

C:\Users\arft\AppData\Local\Temp\ipykernel_20176\1571269568.py:4: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='Status', data=df1, palette='Set2')
```



Handling Description alive

This step filters the dataset to identify alive beneficiaries. with missing values in the Date of Death (DOD) column are assumed to represent beneficiaries who are still alive.

These records are stored in a separate DataFrame for further analysis.

In this step, the dataset is filtered based on the Date of Death (DOD) column.

Beneficiaries with missing DOD values are classified as alive, as no death date is recorded.

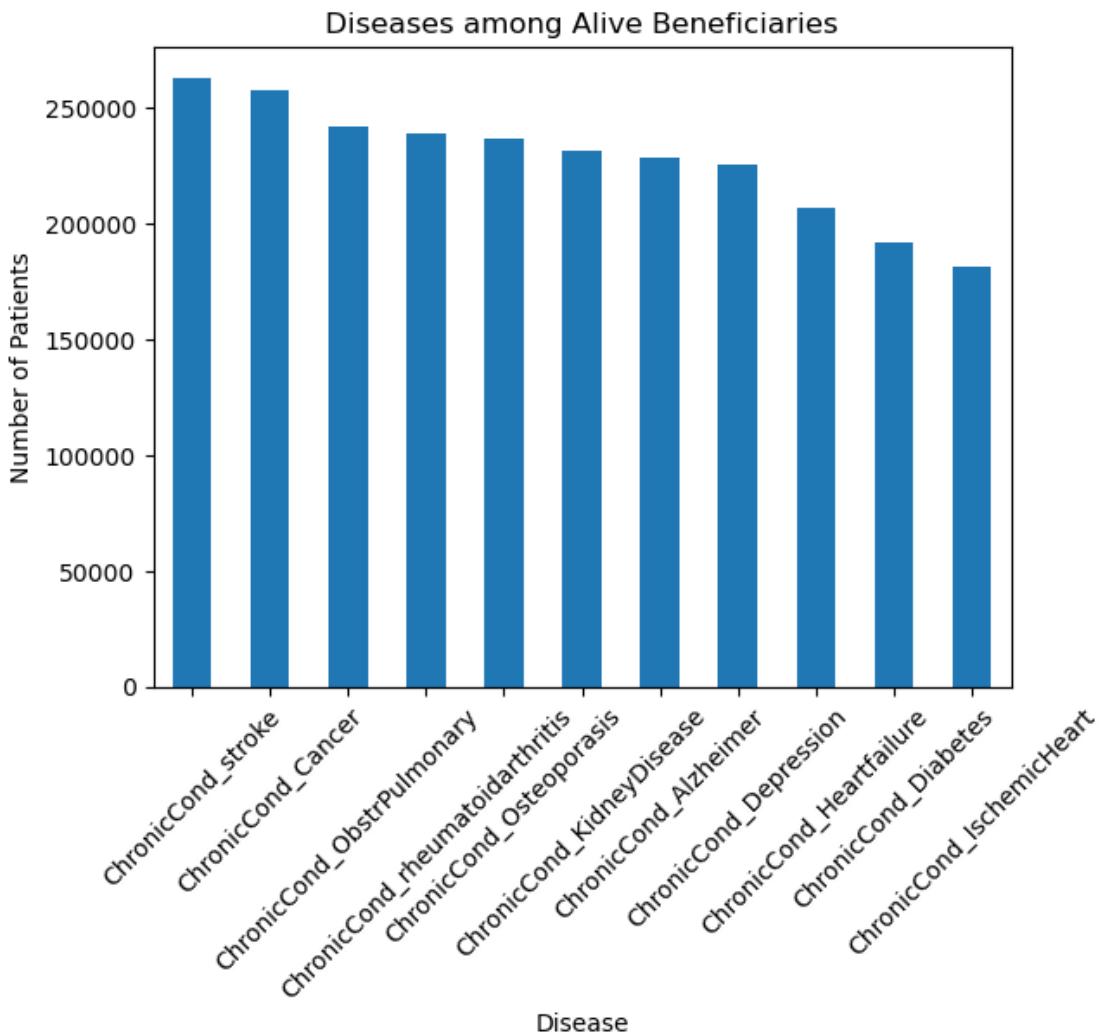
The filtered records are stored in a new DataFrame named alive_df for disease and utilization analysis. This code filters all beneficiaries with missing Date of Death values and classifies them as alive for further analysis

```
[12]: alive_df = df1[df1['DOD'].isna()]
```

```
[13]: alive_diseases = alive_df[  
    ['ChronicCond_Alzheimer',  
     'ChronicCond_Heartfailure',  
     'ChronicCond_KidneyDisease',  
     'ChronicCond_Cancer',  
     'ChronicCond_ObstrPulmonary',  
     'ChronicCond_Depression',  
     'ChronicCond_Diabetes',  
     'ChronicCond_IschemicHeart',  
     'ChronicCond_Osteoporasis',  
     'ChronicCond_rheumatoidarthritis',  
     'ChronicCond_stroke']  
].sum().sort_values(ascending=False)  
  
print(alive_diseases)
```

```
ChronicCond_stroke          263420  
ChronicCond_Cancer           257860  
ChronicCond_ObstrPulmonary   241803  
ChronicCond_rheumatoidarthritis 239087  
ChronicCond_Osteoporasis     236613  
ChronicCond_KidneyDisease    231469  
ChronicCond_Alzheimer        228799  
ChronicCond_Depression       225521  
ChronicCond_Heartfailure     206640  
ChronicCond_Diabetes          191816  
ChronicCond_IschemicHeart    181676  
dtype: int64
```

```
[14]: import matplotlib.pyplot as plt  
  
alive_diseases.plot(kind='bar')  
plt.title("Diseases among Alive Beneficiaries")  
plt.ylabel("Number of Patients")  
plt.xlabel("Disease")  
plt.xticks(rotation=45)  
plt.show()
```



Handling of Deceased Beneficiaries This step identifies deceased beneficiaries in the dataset by filtering records with non-missing values in the Date of Death (DOD) column.

A non-null DOD indicates that the beneficiary has passed away.

These records are stored in a separate DataFrame (dead_df) to analyze disease prevalence and healthcare patterns among deceased patients.

This step filters beneficiaries with a recorded Date of Death (DOD) and classifies them as deceased for further analysis.

Beneficiaries with available Date of Death information are considered deceased.

The dataset is filtered accordingly to study disease distribution and mortality-related trends

```
[15]: df1['DOD'].notna().sum()
```

```
[15]: np.int64(1421)
```

```
[16]: dead_df1 = df1[df1['DOD'].notna()]
```

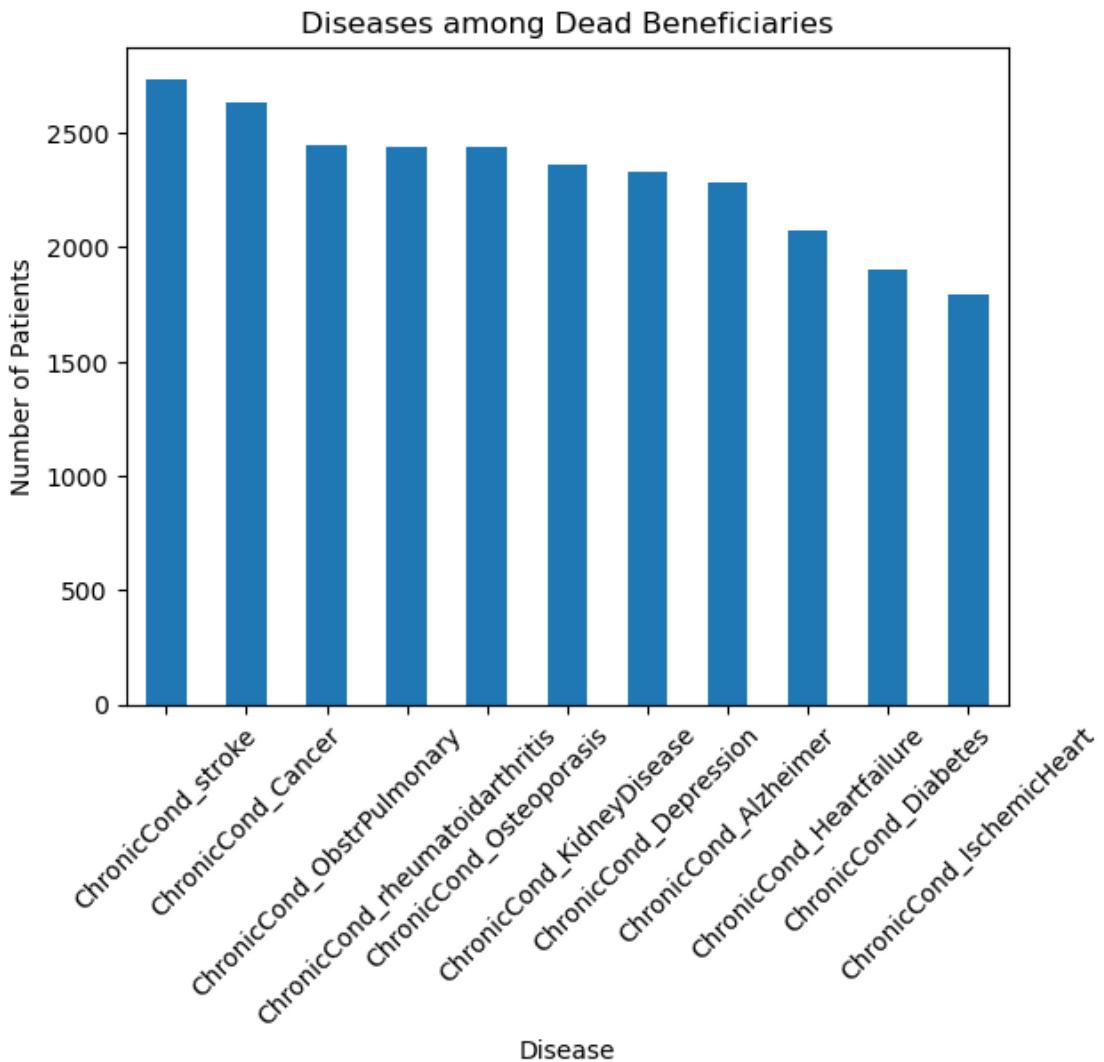
```
[17]: dead_diseases = dead_df1[
    'ChronicCond_Alzheimer',
    'ChronicCond_Heartfailure',
    'ChronicCond_KidneyDisease',
    'ChronicCond_Cancer',
    'ChronicCond_ObstrPulmonary',
    'ChronicCond_Depression',
    'ChronicCond_Diabetes',
    'ChronicCond_IschemicHeart',
    'ChronicCond_Osteoporasis',
    'ChronicCond_rheumatoidarthritis',
    'ChronicCond_stroke']
].sum().sort_values(ascending=False)

print(dead_diseases)
```

```
ChronicCond_stroke          2738
ChronicCond_Cancer           2631
ChronicCond_ObstrPulmonary   2450
ChronicCond_rheumatoidarthritis 2441
ChronicCond_Osteoporasis     2440
ChronicCond_KidneyDisease   2364
ChronicCond_Depression       2331
ChronicCond_Alzheimer        2287
ChronicCond_Heartfailure     2070
ChronicCond_Diabetes          1905
ChronicCond_IschemicHeart    1792
dtype: int64
```

```
[18]: import matplotlib.pyplot as plt

dead_diseases.plot(kind='bar')
plt.title("Diseases among Dead Beneficiaries")
plt.ylabel("Number of Patients")
plt.xlabel("Disease")
plt.xticks(rotation=45)
plt.show()
```



Distribution of Alive and Deceased Beneficiaries The distribution shows that the majority of beneficiaries are classified as alive, as most records do not contain a Date of Death.

A smaller proportion of beneficiaries are marked as deceased, indicating fewer recorded death events in the dataset.

```
[19]: alive_df = df1[df1['DOD'].isna()]
dead_df  = df1[df1['DOD'].notna()]
```

```
[20]: disease_cols = [
    'ChronicCond_Alzheimer',
    'ChronicCond_Heartfailure',
    'ChronicCond_KidneyDisease',
    'ChronicCond_Cancer',
```

```

'ChronicCond_ObstrPulmonary',
'ChronicCond_Depression',
'ChronicCond_Diabetes',
'ChronicCond_IschemicHeart',
'ChronicCond_Osteoporosis',
'ChronicCond_rheumatoidarthritis',
'ChronicCond_stroke'
]

```

```
[21]: alive_counts = alive_df[disease_cols].sum()
dead_counts = dead_df[disease_cols].sum()
```

Description: Alive vs Dead Chronic Disease Comparison compare_df compares the prevalence of chronic diseases between alive and deceased beneficiaries.

Each row represents a specific chronic condition

Columns show: Alive → Number of alive beneficiaries having that disease

Dead → Number of deceased beneficiaries having that disease

```
[22]: compare_df = pd.DataFrame({
    'Alive': alive_counts,
    'Dead': dead_counts
})
```

```
[23]: alive_counts = alive_df[disease_cols].sum()
dead_counts = dead_df[disease_cols].sum()

print("Alive counts:\n", alive_counts)
print("Dead counts:\n", dead_counts)
```

```

Alive counts:
ChronicCond_Alzheimer          228799
ChronicCond_Heartfailure        206640
ChronicCond_KidneyDisease      231469
ChronicCond_Cancer              257860
ChronicCond_ObstrPulmonary      241803
ChronicCond_Depression          225521
ChronicCond_Diabetes            191816
ChronicCond_IschemicHeart       181676
ChronicCond_Osteoporosis        236613
ChronicCond_rheumatoidarthritis 239087
ChronicCond_stroke              263420
dtype: int64
Dead counts:
ChronicCond_Alzheimer          2287
ChronicCond_Heartfailure        2070
ChronicCond_KidneyDisease      2364

```

```

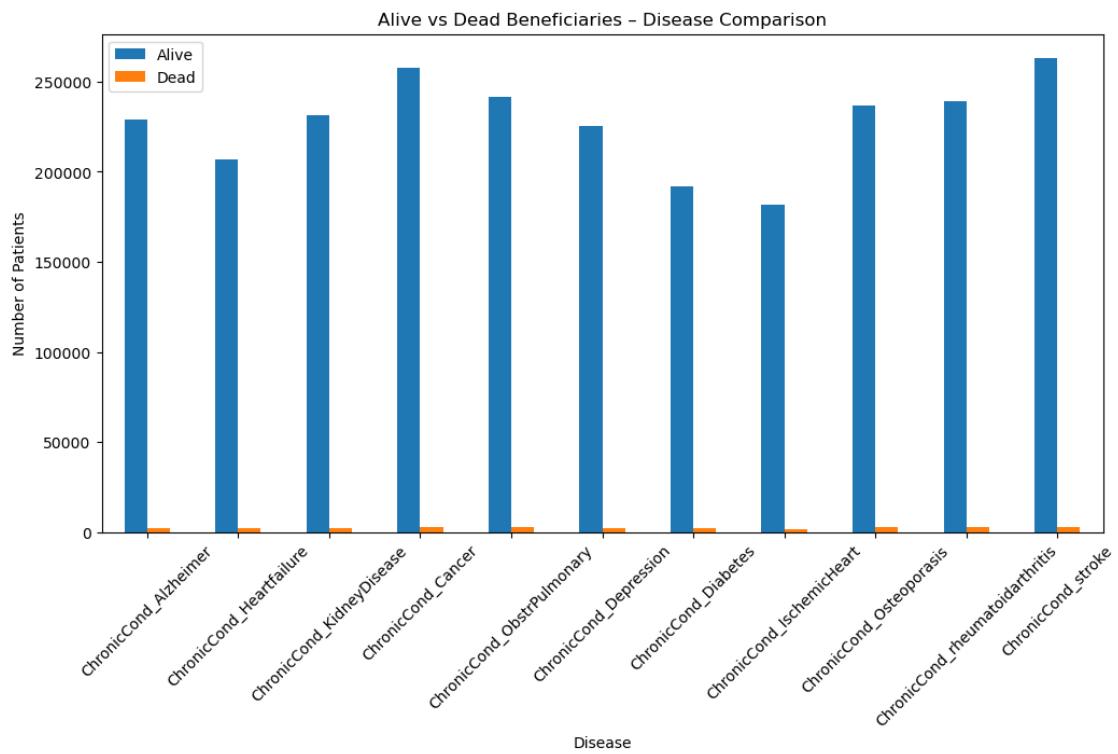
ChronicCond_Cancer           2631
ChronicCond_ObstrPulmonary   2450
ChronicCond_Depression       2331
ChronicCond_Diabetes         1905
ChronicCond_IschemicHeart    1792
ChronicCond_Osteoporasis     2440
ChronicCond_rheumatoidarthritis 2441
ChronicCond_stroke          2738
dtype: int64

```

```
[24]: compare_df = pd.DataFrame({
    'Alive': alive_counts,
    'Dead': dead_counts
})
```

```
[25]: import matplotlib.pyplot as plt

compare_df.plot(kind='bar', figsize=(12,6))
plt.title("Alive vs Dead Beneficiaries - Disease Comparison")
plt.ylabel("Number of Patients")
plt.xlabel("Disease")
plt.xticks(rotation=45)
plt.legend()
plt.show()
```



```
[26]: df1.columns
```

```
[26]: Index(['BeneID', 'DOB', 'DOD', 'Gender', 'Race', 'RenalDiseaseIndicator',
       'State', 'County', 'NoOfMonths_PartACov', 'NoOfMonths_PartBCov',
       'ChronicCond_Alzheimer', 'ChronicCond_Heartfailure',
       'ChronicCond_KidneyDisease', 'ChronicCond_Cancer',
       'ChronicCond_ObstrPulmonary', 'ChronicCond_Depression',
       'ChronicCond_Diabetes', 'ChronicCond_IschemicHeart',
       'ChronicCond_Osteoporasis', 'ChronicCond_rheumatoidarthritis',
       'ChronicCond_stroke', 'IPAnnualReimbursementAmt',
       'IPAnnualDeductibleAmt', 'OPAnnualReimbursementAmt',
       'OPAnnualDeductibleAmt', 'Status'],
      dtype='object')
```

Handling and Distribution of Date of Birth (DOB) ##### The DOB (Date of Birth) column is converted to a datetime format using pd.to_datetime(). ##### Any invalid or incorrectly formatted entries are coerced to NaT (Not a Time) to handle errors and missing data gracefully. ##### This standardization allows for accurate age calculation, chronological analysis, and demographic distribution of beneficiaries.

```
[27]: import pandas as pd
```

```
df1['DOB'] = pd.to_datetime(
    df1['DOB'], errors='coerce'
)
```

```
[28]: from datetime import datetime
```

```
df1['Age'] = (
    (pd.to_datetime('today') - df1['DOB']).dt.days // 365
)
```

```
[29]: df1['DOB'] = pd.to_datetime(
    df1['DOB'], errors='coerce'
)
```

```
print(df1['DOB'].head())
```

```
0    1943-01-01
1    1936-09-01
2    1936-08-01
3    1922-07-01
4    1935-09-01
Name: DOB, dtype: datetime64[ns]
```

```
[30]: df1.head
```

```
[30]: <bound method NDFrame.head of
RenalDiseaseIndicator  State \
0      BENE11001 1943-01-01 NaT      1      1      0      39
1      BENE11002 1936-09-01 NaT      2      1      0      39
2      BENE11003 1936-08-01 NaT      1      1      0      52
3      BENE11004 1922-07-01 NaT      1      1      0      39
4      BENE11005 1935-09-01 NaT      1      1      0      24
...
138551  BENE159194 1939-07-01 NaT      1      1      0      39
138552  BENE159195 1938-12-01 NaT      2      1      0      49
138553  BENE159196 1916-06-01 NaT      2      1      0      6
138554  BENE159197 1930-01-01 NaT      1      1      0      16
138555  BENE159198 1952-04-01 NaT      2      1      0      21

County  NoOfMonths_PartACov  NoOfMonths_PartBCov ... \
0        230                  12                  12 ...
1        280                  12                  12 ...
2        590                  12                  12 ...
3        270                  12                  12 ...
4        680                  12                  12 ...
...
138551  140                  12                  12 ...
138552  530                  12                  12 ...
138553  150                  12                  12 ...
138554  560                  12                  12 ...
138555  20                   12                  12 ...

ChronicCond_IschemicHeart  ChronicCond_Osteoporosis \
0                          1                          2
1                          2                          2
2                          1                          2
3                          1                          1
4                          2                          2
...
138551                      2                      2
138552                      2                      2
138553                      1                      2
138554                      1                      2
138555                      2                      2

ChronicCond_rheumatoidarthritis  ChronicCond_stroke \
0                           1                           1
1                           2                           2
2                           2                           2
3                           1                           2
4                           2                           2
...
```

```

138551          2          2
138552          2          2
138553          2          2
138554          2          2
138555          1          2

```

	IPAnnualReimbursementAmt	IPAnnualDeductibleAmt	\
0	36000	3204	
1	0	0	
2	0	0	
3	0	0	
4	0	0	
...	
138551	0	0	
138552	0	0	
138553	2000	1068	
138554	0	0	
138555	0	0	

	OPAnnualReimbursementAmt	OPAnnualDeductibleAmt	Status	Age
0	60	70	Alive	83
1	30	50	Alive	89
2	90	40	Alive	89
3	1810	760	Alive	103
4	1790	1200	Alive	90
...
138551	430	460	Alive	86
138552	880	100	Alive	87
138553	3240	1390	Alive	109
138554	2650	10	Alive	96
138555	5470	1870	Alive	73

[138556 rows x 27 columns]>

[31]: `print(df1['Age'].head())`

```

0    83
1    89
2    89
3   103
4    90
Name: Age, dtype: int64

```

[32]: `df1[['DOB', 'Age']].head()`

```

[32]:      DOB  Age
0  1943-01-01   83
1  1936-09-01   89

```

```
2 1936-08-01    89
3 1922-07-01   103
4 1935-09-01    90
```

Descriptive Gender ##### Both male and female beneficiaries are present in the dataset.

Gender helps understand utilization patterns, not fraud by itself.

Certain chronic diseases and claim behaviors may vary by gender.

```
[33]: print(df1['Gender'].unique())
```

```
[1 2]
```

```
[34]: print(df1['Gender'].value_counts())
```

```
Gender
2    79106
1    59450
Name: count, dtype: int64
```

```
[35]: df1[['Gender']].head()
```

```
[35]:   Gender
0      1
1      2
2      1
3      1
4      1
```

```
[36]: print(df1['Gender'].value_counts())
```

```
Gender
2    79106
1    59450
Name: count, dtype: int64
```

```
[37]: # Example 1: 'M'/'F' to 'Male'/'Female'
df1['Gender'] = df1['Gender'].replace({'M':'Male', 'F':'Female'})

# Example 2: 1/2 to 'Male'/'Female'
df1['Gender'] = df1['Gender'].replace({1:'Male', 2:'Female'})
```

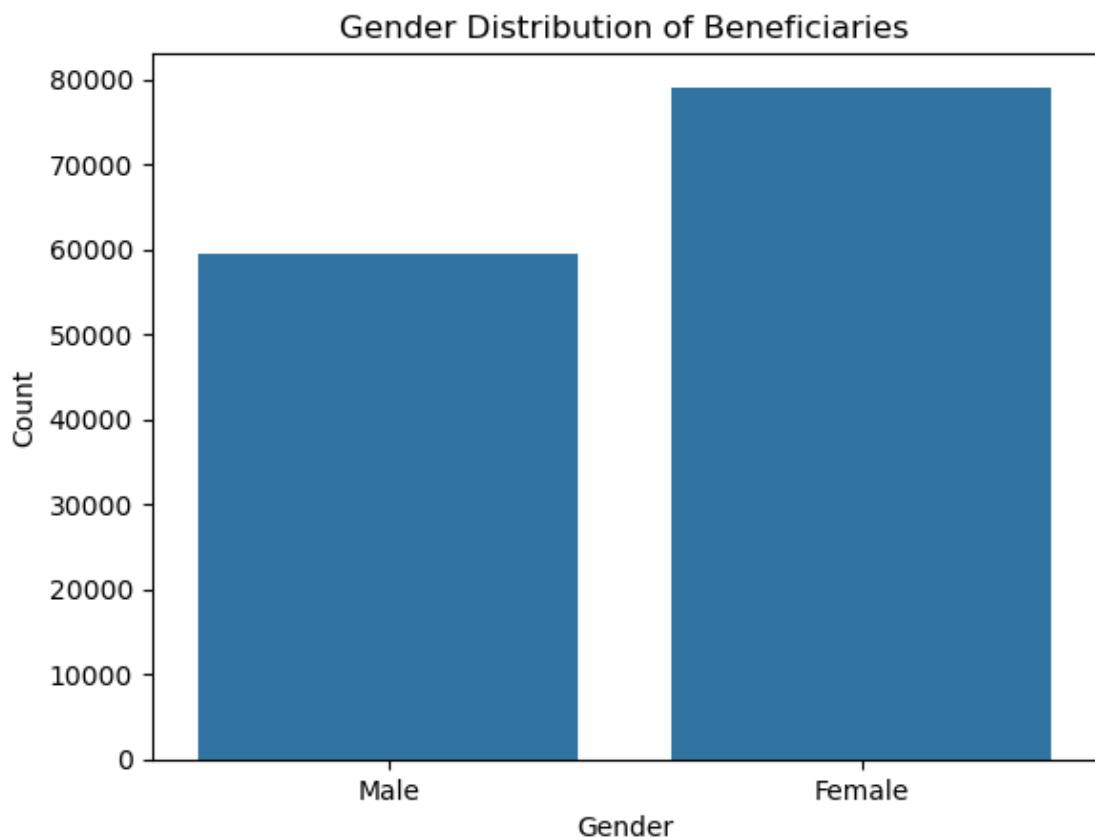
```
[38]: print(df1['Gender'].value_counts())
print(df1[['Gender']].head())
```

```
Gender
Female    79106
Male      59450
Name: count, dtype: int64
Gender
```

```
0    Male
1  Female
2    Male
3    Male
4    Male
```

```
[39]: import seaborn as sns
import matplotlib.pyplot as plt

sns.countplot(x='Gender', data=df1)
plt.title("Gender Distribution of Beneficiaries")
plt.xlabel("Gender")
plt.ylabel("Count")
plt.show()
```



```
[40]: df1.isna().sum()
```

```
[40]: BeneID          0
DOB            0
DOD        137135
```

```

Gender          0
Race           0
RenalDiseaseIndicator 0
State          0
County         0
NoOfMonths_PartACov 0
NoOfMonths_PartBCov 0
ChronicCond_Alzheimer 0
ChronicCond_Heartfailure 0
ChronicCond_KidneyDisease 0
ChronicCond_Cancer 0
ChronicCond_ObstrPulmonary 0
ChronicCond_Depression 0
ChronicCond_Diabetes 0
ChronicCond_IschemicHeart 0
ChronicCond_Osteoporasis 0
ChronicCond_rheumatoidarthritis 0
ChronicCond_stroke 0
IPAnnualReimbursementAmt 0
IPAnnualDeductibleAmt 0
OPAnnualReimbursementAmt 0
OPAnnualDeductibleAmt 0
Status          0
Age            0
dtype: int64

```

[41]: `print(df1.columns)`

```

Index(['BeneID', 'DOB', 'DOD', 'Gender', 'Race', 'RenalDiseaseIndicator',
       'State', 'County', 'NoOfMonths_PartACov', 'NoOfMonths_PartBCov',
       'ChronicCond_Alzheimer', 'ChronicCond_Heartfailure',
       'ChronicCond_KidneyDisease', 'ChronicCond_Cancer',
       'ChronicCond_ObstrPulmonary', 'ChronicCond_Depression',
       'ChronicCond_Diabetes', 'ChronicCond_IschemicHeart',
       'ChronicCond_Osteoporasis', 'ChronicCond_rheumatoidarthritis',
       'ChronicCond_stroke', 'IPAnnualReimbursementAmt',
       'IPAnnualDeductibleAmt', 'OPAnnualReimbursementAmt',
       'OPAnnualDeductibleAmt', 'Status', 'Age'],
      dtype='object')

```

[42]: `df1.head()`

	BeneID	DOB	DOD	Gender	Race	RenalDiseaseIndicator	State	\
0	BENE11001	1943-01-01	NaT	Male	1		0	39
1	BENE11002	1936-09-01	NaT	Female	1		0	39
2	BENE11003	1936-08-01	NaT	Male	1		0	52
3	BENE11004	1922-07-01	NaT	Male	1		0	39
4	BENE11005	1935-09-01	NaT	Male	1		0	24

```

County  NoOfMonths_PartACov  NoOfMonths_PartBCov  ... \
0      230                  12                  12    ...
1      280                  12                  12    ...
2      590                  12                  12    ...
3      270                  12                  12    ...
4      680                  12                  12    ...

ChronicCond_IschemicHeart  ChronicCond_Osteoporasis  \
0                      1                      2
1                      2                      2
2                      1                      2
3                      1                      1
4                      2                      2

ChronicCond_rheumatoidarthritis  ChronicCond_stroke  \
0                         1                         1
1                         2                         2
2                         2                         2
3                         1                         2
4                         2                         2

IPAnnualReimbursementAmt  IPAnnualDeductibleAmt  OPAnnualReimbursementAmt  \
0             36000                  3204                  60
1                 0                  0                  30
2                 0                  0                  90
3                 0                  0                1810
4                 0                  0                1790

OPAnnualDeductibleAmt  Status  Age
0          70  Alive   83
1          50  Alive   89
2          40  Alive   89
3         760  Alive  103
4        1200  Alive   90

[5 rows x 27 columns]

```

Handling Missing Values in Race Column The Race column contains missing values which are filled using forward fill (ffill) method.

This method replaces each missing value with the last valid observation above it in the dataset.

Forward filling ensures that there are no null values in the Race column, enabling accurate demographic analysis and visualization.

```
[43]: df1['Race'] = df1['Race'].ffill()
```

```
[44]: print(df1['Race'].isna().sum())
print(df1['Race'].value_counts())
```

```
0
Race
1    117057
2     13538
3      5059
5      2902
Name: count, dtype: int64
```

```
[45]: race_map = {
    1: 'White',
    2: 'Black',
    3: 'Other',
    5: 'Unknown'
}

df1['Race'] = df1['Race'].map(race_map)
```

```
[46]: print(df1[['Race']].head()) # Top 5 rows
```

```
Race
0  White
1  White
2  White
3  White
4  White
```

```
[47]: print(df1['Race'].value_counts())
```

```
Race
White      117057
Black       13538
Other        5059
Unknown      2902
Name: count, dtype: int64
```

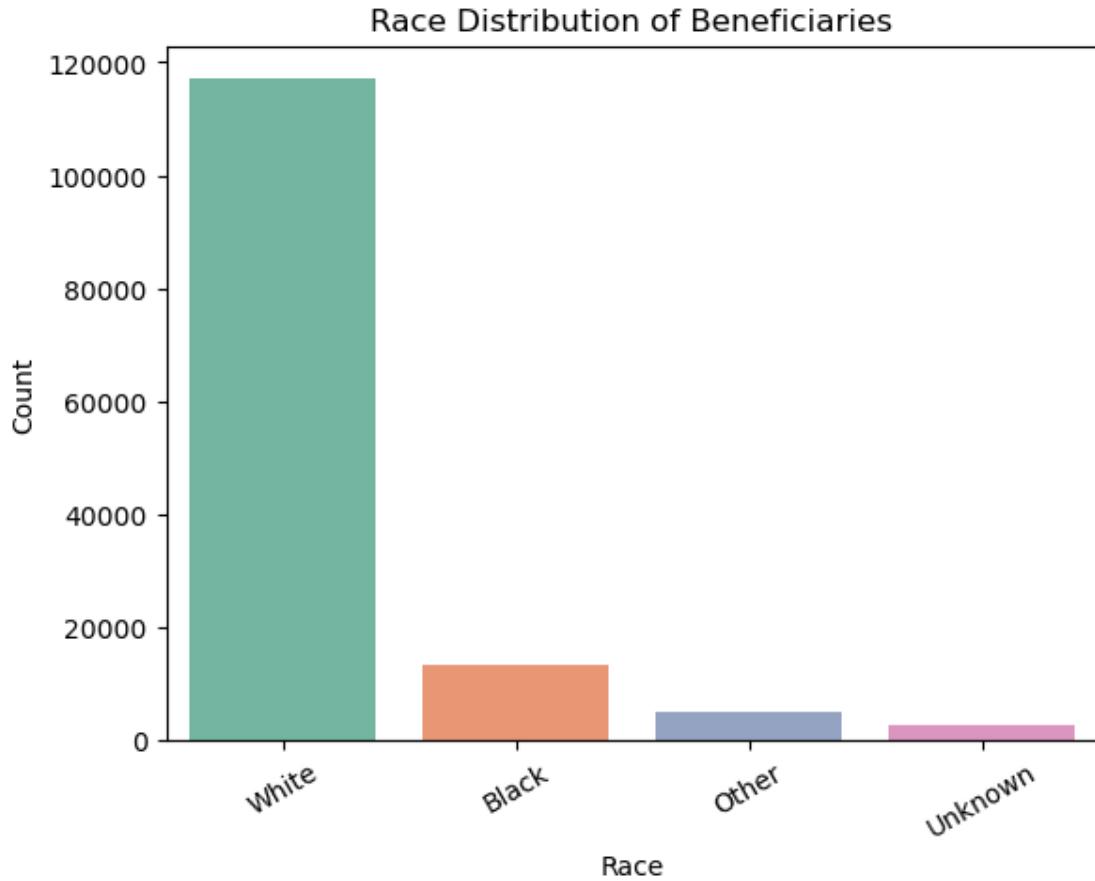
```
[48]: import seaborn as sns
import matplotlib.pyplot as plt

sns.countplot(x='Race', data=df1, palette='Set2')
plt.title("Race Distribution of Beneficiaries")
plt.xlabel("Race")
plt.ylabel("Count")
plt.xticks(rotation=30)
plt.show()
```

```
C:\Users\arft\AppData\Local\Temp\ipykernel_6528\1672263856.py:4: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='Race', data=df1, palette='Set2')
```



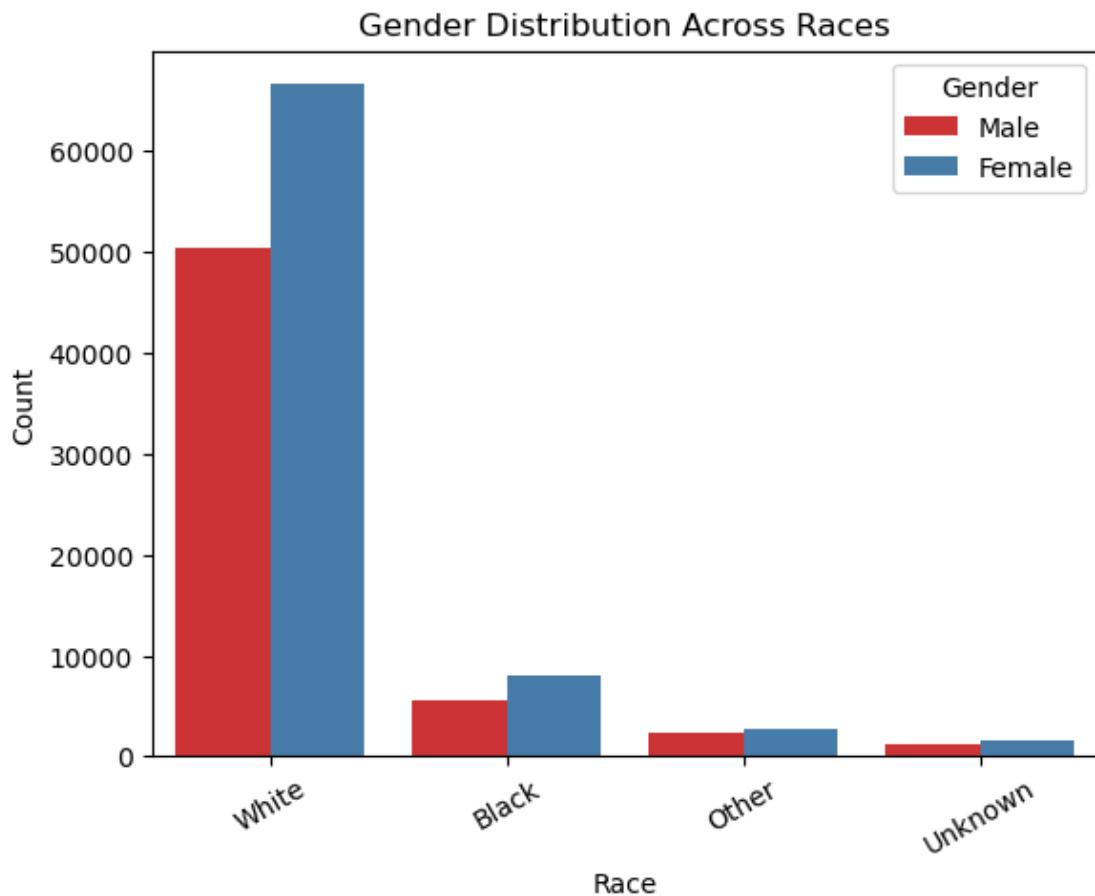
Count Plot (Race Distribution) A count plot is used to visualize the distribution of beneficiaries across different racial categories.

Before plotting, missing values in the Race column were handled using forward fill (ffill) to ensure completeness of data.

The count plot shows the number of beneficiaries in each race, providing insights into the demographic composition of the dataset.

This analysis helps identify if certain racial groups are overrepresented or underrepresented, which can be important for healthcare fraud detection and disease prevalence studies.

```
[49]: sns.countplot(x='Race', hue='Gender', data=df1, palette='Set1')
plt.title("Gender Distribution Across Races")
plt.xlabel("Race")
plt.ylabel("Count")
plt.xticks(rotation=30)
plt.show()
```



```
[50]: df1.columns
```

```
[50]: Index(['BeneID', 'DOB', 'DOD', 'Gender', 'Race', 'RenalDiseaseIndicator',
       'State', 'County', 'NoOfMonths_PartACov', 'NoOfMonths_PartBCov',
       'ChronicCond_Alzheimer', 'ChronicCond_Heartfailure',
       'ChronicCond_KidneyDisease', 'ChronicCond_Cancer',
       'ChronicCond_ObstrPulmonary', 'ChronicCond_Depression',
       'ChronicCond_Diabetes', 'ChronicCond_IschemicHeart',
       'ChronicCond_Osteoporasis', 'ChronicCond_rheumatoidarthritis',
       'ChronicCond_stroke', 'IPAnnualReimbursementAmt',
       'IPAnnualDeductibleAmt', 'OPAnnualReimbursementAmt',
```

```
'OPAnnualDeductibleAmt', 'Status', 'Age'],
dtype='object')
```

```
[51]: df1.head()
```

```
[51]:      BeneID      DOB DOD  Gender  Race RenalDiseaseIndicator  State  \
0  BENE11001  1943-01-01 NaT    Male  White                      0     39
1  BENE11002  1936-09-01 NaT  Female  White                      0     39
2  BENE11003  1936-08-01 NaT    Male  White                      0     52
3  BENE11004  1922-07-01 NaT    Male  White                      0     39
4  BENE11005  1935-09-01 NaT    Male  White                      0     24

      County  NoOfMonths_PartACov  NoOfMonths_PartBCov  ...  \
0        230                 12                  12   ...
1        280                 12                  12   ...
2        590                 12                  12   ...
3        270                 12                  12   ...
4        680                 12                  12   ...

  ChronicCond_IschemicHeart  ChronicCond_Osteoporosis  \
0                     1                         2
1                     2                         2
2                     1                         2
3                     1                         1
4                     2                         2

  ChronicCond_rheumatoidarthritis  ChronicCond_stroke  \
0                           1                         1
1                           2                         2
2                           2                         2
3                           1                         2
4                           2                         2

  IPAnnualReimbursementAmt  IPAnnualDeductibleAmt  OPAnnualReimbursementAmt  \
0                   36000                  3204                          60
1                     0                  0                            30
2                     0                  0                            90
3                     0                  0                          1810
4                     0                  0                          1790

  OPAnnualDeductibleAmt  Status  Age
0             70  Alive   83
1             50  Alive   89
2             40  Alive   89
3            760  Alive  103
4           1200  Alive   90
```

```
[5 rows x 27 columns]
```

Calculation and Distribution of Age A new Age column is derived from the DOB (Date of Birth) column by subtracting the year of birth from 2010.

This creates a numeric representation of each beneficiary's age, enabling demographic and statistical analysis.

Any invalid or missing DOB entries (NaT) are automatically excluded from this calculation.

```
[52]: df1['DOB'] = pd.to_datetime(df1['DOB'], errors='coerce')
```

```
[53]: df1['Age'] = 2010 - df1['DOB'].dt.year
```

```
[54]: print(df1[['DOB', 'Age']].head())
```

	DOB	Age
0	1943-01-01	67
1	1936-09-01	74
2	1936-08-01	74
3	1922-07-01	88
4	1935-09-01	75

```
[56]: df1 = df1[(df1['Age'] > 0) & (df1['Age'] < 120)]
```

Plot Description: Age Distribution Across Races

Age distributions vary across race groups Some race categories have higher median ages, while others show a wider age spread, indicating differences in healthcare utilization patterns.

Presence of outliers in certain races

Extreme age values (very old beneficiaries) appear in some race groups, which often correlates with: Higher medical complexity

Increased claim frequency and costs

Age is a strong cost driver, not a fraud indicator by itself

Older age groups typically require more medical services, which may naturally lead to higher claims.

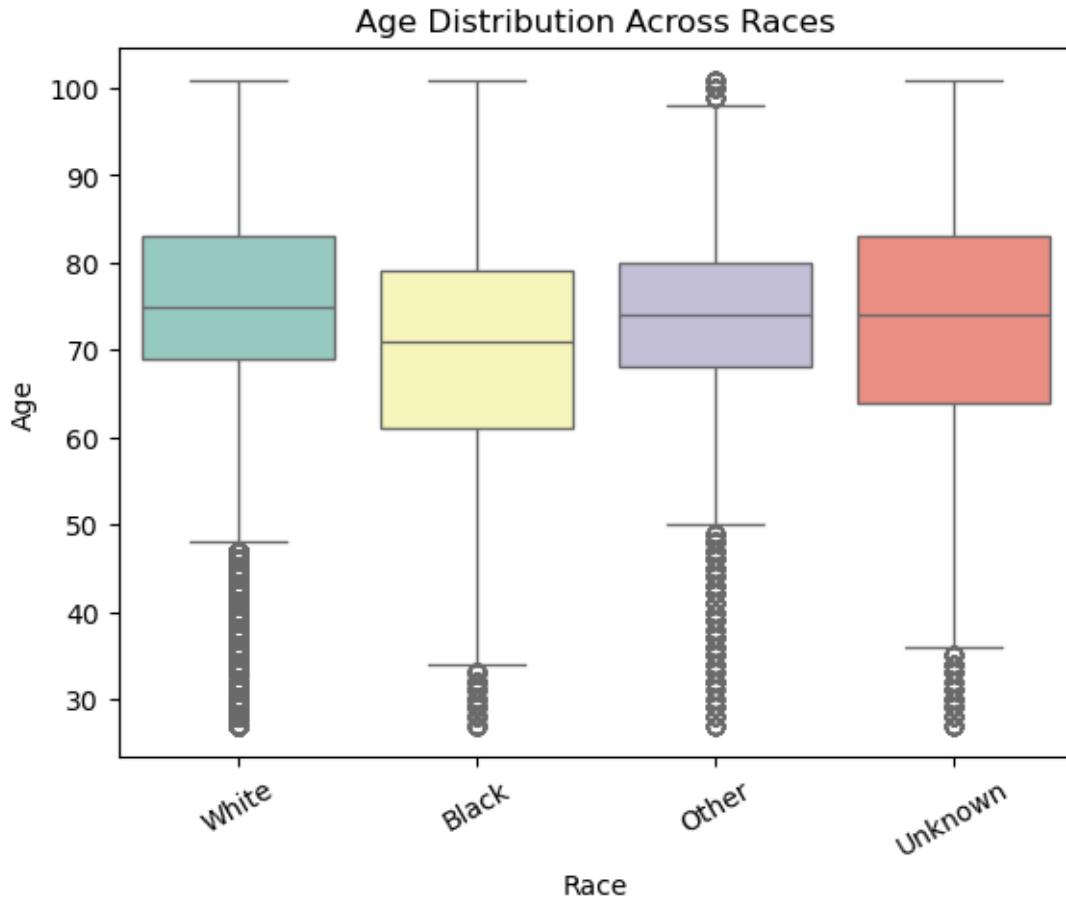
```
[57]: sns.boxplot(x='Race', y='Age', data=df1, palette='Set3')
plt.title("Age Distribution Across Races")
plt.xlabel("Race")
plt.ylabel("Age")
plt.xticks(rotation=30)
plt.show()
```

```
C:\Users\arft\AppData\Local\Temp\ipykernel_6528\3959494853.py:1: FutureWarning:
```

```
Passing `palette` without assigning `hue` is deprecated and will be removed in
```

v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(x='Race', y='Age', data=df1, palette='Set3')
```



```
[58]: print(df1.columns)
```

```
Index(['BeneID', 'DOB', 'DOD', 'Gender', 'Race', 'RenalDiseaseIndicator',
       'State', 'County', 'NoOfMonths_PartACov', 'NoOfMonths_PartBCov',
       'ChronicCond_Alzheimer', 'ChronicCond_Heartfailure',
       'ChronicCond_KidneyDisease', 'ChronicCond_Cancer',
       'ChronicCond_ObstrPulmonary', 'ChronicCond_Depression',
       'ChronicCond_Diabetes', 'ChronicCond_IschemicHeart',
       'ChronicCond_Osteoporosis', 'ChronicCond_rheumatoidarthritis',
       'ChronicCond_stroke', 'IPAnnualReimbursementAmt',
       'IPAnnualDeductibleAmt', 'OPAnnualReimbursementAmt',
       'OPAnnualDeductibleAmt', 'Status', 'Age'],
      dtype='object')
```

```
[59]: df1.head
```

```
[59]: <bound method NDFrame.head of
RenalDiseaseIndicator  State \
0      BENE11001 1943-01-01 NaT    Male  White          0    39
1      BENE11002 1936-09-01 NaT  Female  White          0    39
2      BENE11003 1936-08-01 NaT    Male  White          0    52
3      BENE11004 1922-07-01 NaT    Male  White          0    39
4      BENE11005 1935-09-01 NaT    Male  White          0    24
...
138551   ...     ...     ...     ...     ...
138551   BENE159194 1939-07-01 NaT    Male  White          0    39
138552   BENE159195 1938-12-01 NaT  Female  White          0    49
138553   BENE159196 1916-06-01 NaT  Female  White          0     6
138554   BENE159197 1930-01-01 NaT    Male  White          0    16
138555   BENE159198 1952-04-01 NaT  Female  White          0    21

      County  NoOfMonths_PartACov  NoOfMonths_PartBCov  ...  \
0        230            12             12    ... 
1        280            12             12    ... 
2        590            12             12    ... 
3        270            12             12    ... 
4        680            12             12    ... 
...
138551   ...     ...     ...     ...     ...
138551   140            12             12    ... 
138552   530            12             12    ... 
138553   150            12             12    ... 
138554   560            12             12    ... 
138555   20             12             12    ... 

      ChronicCond_IschemicHeart  ChronicCond_Osteoporosis  \
0                      1                  2
1                      2                  2
2                      1                  2
3                      1                  1
4                      2                  2
...
138551   ...     ...
138551   2                  2
138552   2                  2
138553   1                  2
138554   1                  2
138555   2                  2

      ChronicCond_rheumatoidarthritis  ChronicCond_stroke  \
0                      1                  1
1                      2                  2
2                      2                  2
3                      1                  2
```

```

4                               2                               2
...
...                               ...                               ...
138551                           2                               2
138552                           2                               2
138553                           2                               2
138554                           2                               2
138555                           1                               2

IPAnnualReimbursementAmt  IPAnnualDeductibleAmt \
0                         36000                          3204
1                         0                             0
2                         0                             0
3                         0                             0
4                         0                             0
...
...                               ...                               ...
138551                           0                               0
138552                           0                               0
138553                         2000                          1068
138554                           0                             0
138555                           0                             0

OPAnnualReimbursementAmt  OPAnnualDeductibleAmt  Status  Age
0                         60                            70  Alive   67
1                         30                            50  Alive   74
2                         90                            40  Alive   74
3                        1810                          760  Alive   88
4                        1790                          1200  Alive   75
...
...                               ...                               ...
138551                           430                          460  Alive   71
138552                           880                          100  Alive   72
138553                         3240                          1390  Alive   94
138554                         2650                           10  Alive   80
138555                         5470                          1870  Alive   58

```

[138556 rows x 27 columns]>

[60]: `print(df1.columns.tolist())`

```

['BeneID', 'DOB', 'DOD', 'Gender', 'Race', 'RenalDiseaseIndicator', 'State',
'County', 'NoOfMonths_PartACov', 'NoOfMonths_PartBCov', 'ChronicCond_Alzheimer',
'ChronicCond_Heartfailure', 'ChronicCond_KidneyDisease', 'ChronicCond_Cancer',
'ChronicCond_ObstrPulmonary', 'ChronicCond_Depression', 'ChronicCond_Diabetes',
'ChronicCond_IschemicHeart', 'ChronicCond_Osteoporasis',
'ChronicCond_rheumatoidarthritis', 'ChronicCond_stroke',
'IPAnnualReimbursementAmt', 'IPAnnualDeductibleAmt', 'OPAnnualReimbursementAmt',
'OPAnnualDeductibleAmt', 'Status', 'Age']

```

```
[61]: chronic_cols = [col for col in df1.columns if 'ChronicCond' in col]
print(chronic_cols)

['ChronicCond_Alzheimer', 'ChronicCond_Heartfailure',
 'ChronicCond_KidneyDisease', 'ChronicCond_Cancer', 'ChronicCond_ObstrPulmonary',
 'ChronicCond_Depression', 'ChronicCond_Diabetes', 'ChronicCond_IschemicHeart',
 'ChronicCond_Osteoporasis', 'ChronicCond_rheumatoidarthritis',
 'ChronicCond_stroke']

[62]: df1[chronic_cols].isna().sum()

[62]: ChronicCond_Alzheimer      0
ChronicCond_Heartfailure       0
ChronicCond_KidneyDisease     0
ChronicCond_Cancer             0
ChronicCond_ObstrPulmonary     0
ChronicCond_Depression         0
ChronicCond_Diabetes           0
ChronicCond_IschemicHeart      0
ChronicCond_Osteoporasis       0
ChronicCond_rheumatoidarthritis 0
ChronicCond_stroke             0
dtype: int64
```

Unknown Values in Chronic Condition Columns This step sums up the values in all chronic condition columns to calculate the total number of patients affected by each disease.

Before plotting, missing or unknown values should be properly handled (e.g., replaced with 0 or cleaned) to ensure accuracy.

A bar plot is then created to visualize the prevalence of chronic diseases across the dataset

```
[63]: for col in chronic_cols:
    print(col, ":", (df1[col] == '').sum(), "'empty strings'")
    print(col, ":", (df1[col] == 'Unknown').sum(), "'Unknown' values")

ChronicCond_Alzheimer : 0 empty strings
ChronicCond_Alzheimer : 0 'Unknown' values
ChronicCond_Heartfailure : 0 empty strings
ChronicCond_Heartfailure : 0 'Unknown' values
ChronicCond_KidneyDisease : 0 empty strings
ChronicCond_KidneyDisease : 0 'Unknown' values
ChronicCond_Cancer : 0 empty strings
ChronicCond_Cancer : 0 'Unknown' values
ChronicCond_ObstrPulmonary : 0 empty strings
ChronicCond_ObstrPulmonary : 0 'Unknown' values
ChronicCond_Depression : 0 empty strings
ChronicCond_Depression : 0 'Unknown' values
ChronicCond_Diabetes : 0 empty strings
```

```
ChronicCond_Diabetes : 0 'Unknown' values
ChronicCond_IschemicHeart : 0 empty strings
ChronicCond_IschemicHeart : 0 'Unknown' values
ChronicCond_Osteoporosis : 0 empty strings
ChronicCond_Osteoporosis : 0 'Unknown' values
ChronicCond_rheumatoidarthritis : 0 empty strings
ChronicCond_rheumatoidarthritis : 0 'Unknown' values
ChronicCond_stroke : 0 empty strings
ChronicCond_stroke : 0 'Unknown' values
```

```
Basic statistics for chronic disease columns desc_stats = df1[chronic_cols].describe()
print(desc_stats)
```

Plot Description: Number of Patients with Chronic Diseases

Most common chronic diseases are identified

Diseases with the tallest bars are more prevalent among beneficiaries.

Example: Diabetes or Heart Failure may affect a large portion of the population.

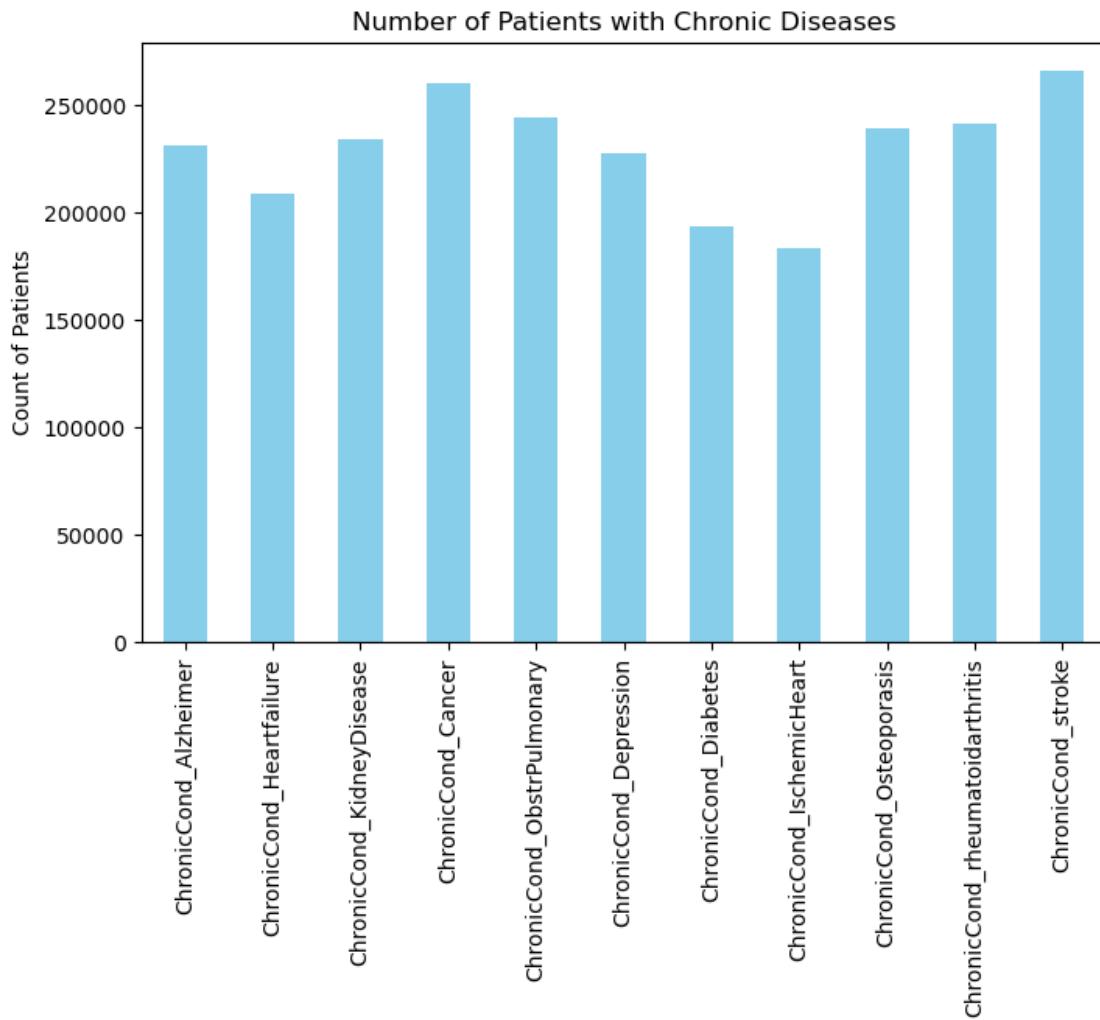
High-prevalence diseases indicate higher healthcare costs

Chronic diseases often require frequent hospital visits, medications, and treatments.

These patients generate higher claim amounts, which could impact insurer finances.

```
[64]: import matplotlib.pyplot as plt

df1[chronic_cols].sum().plot(kind='bar', color='skyblue', figsize=(8,5))
plt.title("Number of Patients with Chronic Diseases")
plt.ylabel("Count of Patients")
plt.show()
```



Plot Description: Correlation Between Chronic Diseases Certain chronic diseases co-occur frequently

Positive correlations indicate that patients with one condition are more likely to have another.

Example: Diabetes and Heart Failure often show moderate to strong correlation.

Multimorbidity is common Many patients have multiple chronic conditions, which leads to:

Higher total claims

Longer hospital stays

Increased complexity of care

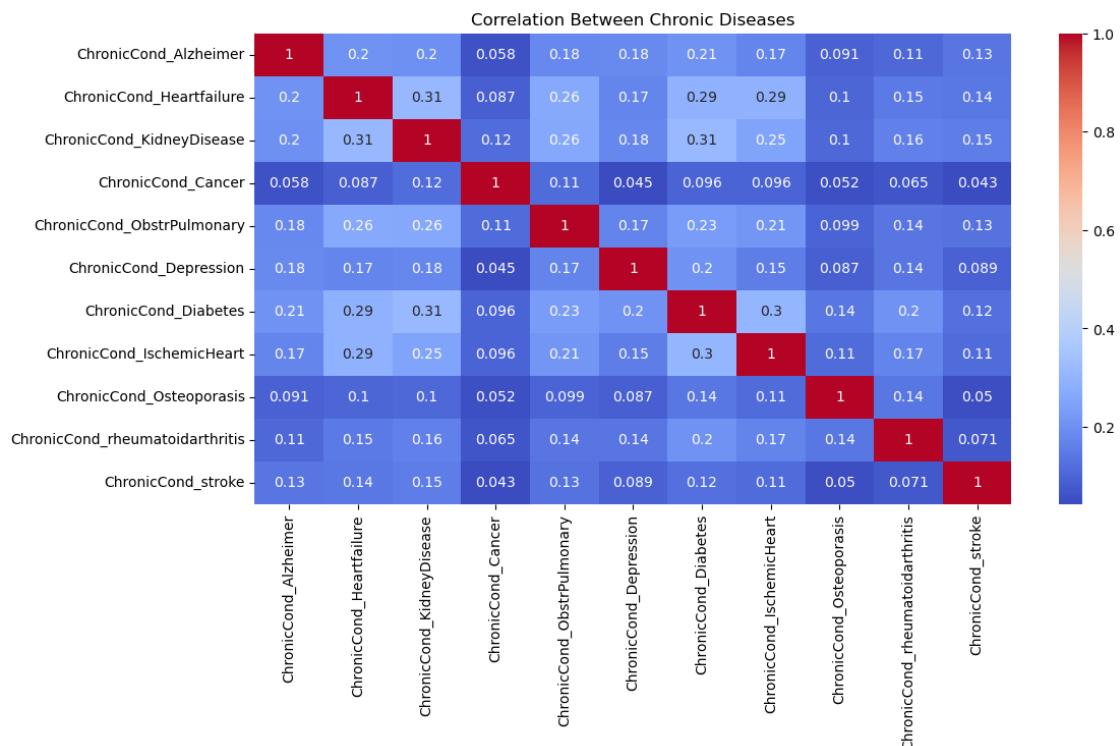
Implications for fraud detection

Providers serving multimorbid patients may naturally have higher claim amounts.

Fraud monitoring should account for co-occurring conditions to avoid false positives.

```
[65]: import seaborn as sns

plt.figure(figsize=(12,6))
sns.heatmap(df1[chronic_cols].corr(), annot=True, cmap='coolwarm')
plt.title("Correlation Between Chronic Diseases")
plt.show()
```



Description: Patient Count per Chronic Disease

```
[66]: disease_counts = df1[chronic_cols].sum().reset_index()
disease_counts.columns = ['Chronic Disease', 'Patient Count']
print(disease_counts)
```

	Chronic Disease	Patient Count
0	ChronicCond_Alzheimer	231086
1	ChronicCond_Heartfailure	208710
2	ChronicCond_KidneyDisease	233833
3	ChronicCond_Cancer	260491
4	ChronicCond_ObstrPulmonary	244253
5	ChronicCond_Depression	227852
6	ChronicCond_Diabetes	193721
7	ChronicCond_IschemicHeart	183468

```
8      ChronicCond_Osteoporasis      239053
9  ChronicCond_rheumatoidarthritis  241528
10     ChronicCond_stroke          266158
```

```
[67]: df1 = df1.drop_duplicates()
print(df1.shape)
```

```
(138556, 27)
```

```
[68]: # Percentage of patients with each disease
disease_percent = (df1[chronic_cols].mean() * 100).round(2).reset_index()
disease_percent.columns = ['Chronic Disease', 'Patient %']
print(disease_percent)
```

	Chronic Disease	Patient %
0	ChronicCond_Alzheimer	166.78
1	ChronicCond_Heartfailure	150.63
2	ChronicCond_KidneyDisease	168.76
3	ChronicCond_Cancer	188.00
4	ChronicCond_ObstrPulmonary	176.28
5	ChronicCond_Depression	164.45
6	ChronicCond_Diabetes	139.81
7	ChronicCond_IschemicHeart	132.41
8	ChronicCond_Osteoporasis	172.53
9	ChronicCond_rheumatoidarthritis	174.32
10	ChronicCond_stroke	192.09

Plot Desc Chronic disease prevalence differs between Alive and Dead groups Some diseases are more common among deceased beneficiaries.

Example: Heart Failure, Cancer, or Stroke may show higher percentages in the Dead group.

High-risk conditions for mortality Diseases with large differences between Dead vs Alive indicate conditions that contribute to higher mortality.

Impact on claim amounts Deceased patients with chronic diseases often generate higher and more complex claims, potentially increasing fraud risk if claims are abnormal.

Fraud monitoring / operational insight Providers treating patients with high percentages of serious chronic conditions should be closely monitored for:

Duplicate claims

Inflated billing

Post-mortem claims

```
[70]: import pandas as pd
import matplotlib.pyplot as plt
```

```

# List of chronic disease columns
chronic_cols = [
    'ChronicCond_Alzheimer',
    'ChronicCond_Heartfailure',
    'ChronicCond_KidneyDisease',
    'ChronicCond_Cancer',
    'ChronicCond_ObstrPulmonary',
    'ChronicCond_Depression',
    'ChronicCond_Diabetes',
    'ChronicCond_IschemicHeart',
    'ChronicCond_Osteoporosis',
    'ChronicCond_rheumatoidarthritis',
    'ChronicCond_stroke'
]

# Split dataset into Alive and Dead
alive_df = df1[df1['DOD'].isna()]
dead_df = df1[df1['DOD'].notna()]

# Count patients with each disease for Alive and Dead
alive_counts = alive_df[chronic_cols].sum() / len(alive_df) * 100
dead_counts = dead_df[chronic_cols].sum() / len(dead_df) * 100

# Combine into a DataFrame
disease_percentage = pd.DataFrame({
    'Chronic Disease': chronic_cols,
    'Alive (%)': alive_counts.values,
    'Dead (%)': dead_counts.values
})

print(disease_percentage)

# Plotting percentage comparison
plt.figure(figsize=(12,6))
bar_width = 0.4
index = range(len(chronic_cols))

plt.bar(index, disease_percentage['Alive (%)'], bar_width, label='Alive', color='skyblue')
plt.bar([i + bar_width for i in index], disease_percentage['Dead (%)'], bar_width, label='Dead', color='salmon')

plt.xticks([i + bar_width/2 for i in index], disease_percentage['Chronic Disease'], rotation=45)
plt.xlabel('Chronic Disease')
plt.ylabel('Percentage of Patients')

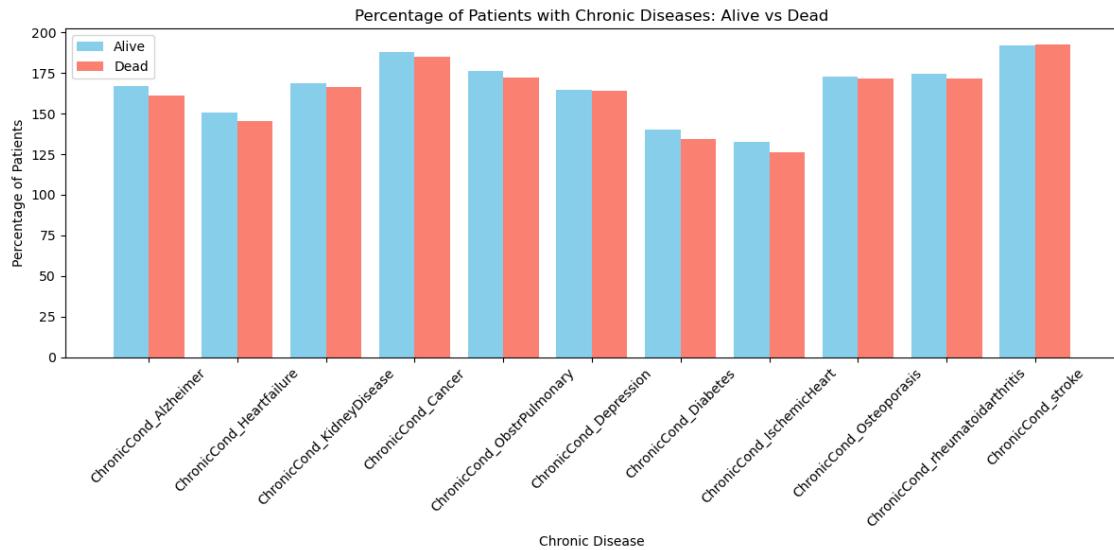
```

```

plt.title('Percentage of Patients with Chronic Diseases: Alive vs Dead')
plt.legend()
plt.tight_layout()
plt.show()

```

	Chronic Disease	Alive (%)	Dead (%)
0	ChronicCond_Alzheimer	166.842163	160.942998
1	ChronicCond_Heartfailure	150.683633	145.672062
2	ChronicCond_KidneyDisease	168.789149	166.361717
3	ChronicCond_Cancer	188.033689	185.151302
4	ChronicCond_ObstrPulmonary	176.324789	172.413793
5	ChronicCond_Depression	164.451818	164.039409
6	ChronicCond_Diabetes	139.873847	134.060521
7	ChronicCond_IschemicHeart	132.479673	126.108374
8	ChronicCond_Osteoporasis	172.540198	171.710063
9	ChronicCond_rheumatoidarthritis	174.344259	171.780436
10	ChronicCond_stroke	192.088088	192.681210



Data Handling & Descriptive Overview – Inpatient & Outpatient Data

```

[14]: import pandas as pd

# Load inpatient and outpatient data
df2 = pd.read_csv(r'D:\Hira\Project 2\Portfolio 2\Train_Inpatientdata.csv')
df3 = pd.read_csv(r'D:\Hira\Project 2\Portfolio 2\Train_Outpatientdata.csv')

# Quick check
print(df2.head())
print(df2.info())

```

```

print(df3.head())
print(df3.info())

      BeneID   ClaimID ClaimStartDt  ClaimEndDt  Provider \
0  BENE11001  CLM46614    2009-04-12  2009-04-18  PRV55912
1  BENE11001  CLM66048    2009-08-31  2009-09-02  PRV55907
2  BENE11001  CLM68358    2009-09-17  2009-09-20  PRV56046
3  BENE11011  CLM38412    2009-02-14  2009-02-22  PRV52405
4  BENE11014  CLM63689    2009-08-13  2009-08-30  PRV56614

      InscClaimAmtReimbursed AttendingPhysician OperatingPhysician \
0                  26000             PHY390922           NaN
1                  5000              PHY318495          PHY318495
2                  5000              PHY372395           NaN
3                  5000              PHY369659          PHY392961
4                 10000             PHY379376          PHY398258

      OtherPhysician AdmissionDt ... ClmDiagnosisCode_7 ClmDiagnosisCode_8 \
0            NaN     2009-04-12 ...          2724          19889
1            NaN     2009-08-31 ...           NaN           NaN
2  PHY324689     2009-09-17 ...           NaN           NaN
3  PHY349768     2009-02-14 ...          25062          40390
4            NaN     2009-08-13 ...          5119          29620

      ClmDiagnosisCode_9 ClmDiagnosisCode_10 ClmProcedureCode_1 \
0                  5849                NaN           NaN
1                  NaN                 NaN          7092.0
2                  NaN                 NaN           NaN
3                  4019                NaN          331.0
4                 20300               NaN          3893.0

      ClmProcedureCode_2 ClmProcedureCode_3 ClmProcedureCode_4 ClmProcedureCode_5 \
0                  NaN                 NaN           NaN           NaN
1                  NaN                 NaN           NaN           NaN
2                  NaN                 NaN           NaN           NaN
3                  NaN                 NaN           NaN           NaN
4                  NaN                 NaN           NaN           NaN

      ClmProcedureCode_6
0                  NaN
1                  NaN
2                  NaN
3                  NaN
4                  NaN

[5 rows x 30 columns]
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 40474 entries, 0 to 40473

```

```

Data columns (total 30 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   BeneID           40474 non-null    object 
 1   ClaimID          40474 non-null    object 
 2   ClaimStartDt     40474 non-null    object 
 3   ClaimEndDt       40474 non-null    object 
 4   Provider          40474 non-null    object 
 5   InscClaimAmtReimbursed  40474 non-null    int64  
 6   AttendingPhysician 40362 non-null    object 
 7   OperatingPhysician 23830 non-null    object 
 8   OtherPhysician     4690 non-null     object 
 9   AdmissionDt       40474 non-null    object 
 10  ClmAdmitDiagnosisCode 40474 non-null    object 
 11  DeductibleAmtPaid 39575 non-null    float64 
 12  DischargeDt       40474 non-null    object 
 13  DiagnosisGroupCode 40474 non-null    object 
 14  ClmDiagnosisCode_1 40474 non-null    object 
 15  ClmDiagnosisCode_2 40248 non-null    object 
 16  ClmDiagnosisCode_3 39798 non-null    object 
 17  ClmDiagnosisCode_4 38940 non-null    object 
 18  ClmDiagnosisCode_5 37580 non-null    object 
 19  ClmDiagnosisCode_6 35636 non-null    object 
 20  ClmDiagnosisCode_7 33216 non-null    object 
 21  ClmDiagnosisCode_8 30532 non-null    object 
 22  ClmDiagnosisCode_9 26977 non-null    object 
 23  ClmDiagnosisCode_10 3927 non-null    object 
 24  ClmProcedureCode_1 23148 non-null    float64 
 25  ClmProcedureCode_2 5454 non-null     float64 
 26  ClmProcedureCode_3 965 non-null      float64 
 27  ClmProcedureCode_4 116 non-null      float64 
 28  ClmProcedureCode_5 9 non-null       float64 
 29  ClmProcedureCode_6 0 non-null       float64 

dtypes: float64(7), int64(1), object(22)
memory usage: 9.3+ MB
None
      BeneID   ClaimID ClaimStartDt ClaimEndDt Provider \
0  BENE11002  CLM624349  2009-10-11  2009-10-11  PRV56011
1  BENE11003  CLM189947  2009-02-12  2009-02-12  PRV57610
2  BENE11003  CLM438021  2009-06-27  2009-06-27  PRV57595
3  BENE11004  CLM121801  2009-01-06  2009-01-06  PRV56011
4  BENE11004  CLM150998  2009-01-22  2009-01-22  PRV56011

      InscClaimAmtReimbursed AttendingPhysician OperatingPhysician \
0                      30             PHY326117             NaN
1                      80             PHY362868             NaN
2                      10             PHY328821             NaN
3                      40             PHY334319             NaN

```

```

4          200        PHY403831       NaN
OtherPhysician ClmDiagnosisCode_1 ... ClmDiagnosisCode_9 \
0           NaN      78943 ...           NaN
1           NaN      6115 ...           NaN
2           NaN      2723 ...           NaN
3           NaN      71988 ...           NaN
4           NaN      82382 ...           NaN

ClmDiagnosisCode_10 ClmProcedureCode_1 ClmProcedureCode_2 \
0           NaN           NaN           NaN
1           NaN           NaN           NaN
2           NaN           NaN           NaN
3           NaN           NaN           NaN
4           NaN           NaN           NaN

ClmProcedureCode_3 ClmProcedureCode_4 ClmProcedureCode_5 ClmProcedureCode_6 \
0           NaN           NaN           NaN           NaN
1           NaN           NaN           NaN           NaN
2           NaN           NaN           NaN           NaN
3           NaN           NaN           NaN           NaN
4           NaN           NaN           NaN           NaN

DeductibleAmtPaid ClmAdmitDiagnosisCode
0             0      56409
1             0      79380
2             0           NaN
3             0           NaN
4             0      71947

[5 rows x 27 columns]
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 517737 entries, 0 to 517736
Data columns (total 27 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   BeneID            517737 non-null   object 
 1   ClaimID           517737 non-null   object 
 2   ClaimStartDt      517737 non-null   object 
 3   ClaimEndDt        517737 non-null   object 
 4   Provider          517737 non-null   object 
 5   InscClaimAmtReimbursed  517737 non-null   int64  
 6   AttendingPhysician 516341 non-null   object 
 7   OperatingPhysician 90617 non-null    object 
 8   OtherPhysician     195046 non-null   object 
 9   ClmDiagnosisCode_1 507284 non-null   object 
 10  ClmDiagnosisCode_2 322357 non-null   object 
 11  ClmDiagnosisCode_3 203257 non-null   object 

```

```

12 ClmDiagnosisCode_4      125596 non-null  object
13 ClmDiagnosisCode_5      74344 non-null   object
14 ClmDiagnosisCode_6      48756 non-null   object
15 ClmDiagnosisCode_7      32961 non-null   object
16 ClmDiagnosisCode_8      22912 non-null   object
17 ClmDiagnosisCode_9      14838 non-null   object
18 ClmDiagnosisCode_10     1083 non-null    object
19 ClmProcedureCode_1       162 non-null    float64
20 ClmProcedureCode_2       36 non-null    float64
21 ClmProcedureCode_3       4 non-null    float64
22 ClmProcedureCode_4       2 non-null    float64
23 ClmProcedureCode_5       0 non-null    float64
24 ClmProcedureCode_6       0 non-null    float64
25 DeductibleAmtPaid      517737 non-null int64
26 ClmAdmitDiagnosisCode  105425 non-null  object
dtypes: float64(6), int64(2), object(19)
memory usage: 106.7+ MB
None

```

```
[15]: import pandas as pd

# Load inpatient and outpatient data
df2 = pd.read_csv(r'D:\Hira\Project 2\Portfolio 2\Train_Inpatientdata.csv')
df3 = pd.read_csv(r'D:\Hira\Project 2\Portfolio 2\Train_Outpatientdata.csv')

# Quick check
print(df2.head())
print(df2.info())
print(df3.head())
print(df3.info())
```

	BeneID	ClaimID	ClaimStartDt	ClaimEndDt	Provider	\
0	BENE11001	CLM46614	2009-04-12	2009-04-18	PRV55912	
1	BENE11001	CLM66048	2009-08-31	2009-09-02	PRV55907	
2	BENE11001	CLM68358	2009-09-17	2009-09-20	PRV56046	
3	BENE11011	CLM38412	2009-02-14	2009-02-22	PRV52405	
4	BENE11014	CLM63689	2009-08-13	2009-08-30	PRV56614	

	InscClaimAmtReimbursed	AttendingPhysician	OperatingPhysician	\
0	26000	PHY390922		NaN
1	5000	PHY318495		PHY318495
2	5000	PHY372395		NaN
3	5000	PHY369659		PHY392961
4	10000	PHY379376		PHY398258

	OtherPhysician	AdmissionDt	...	ClmDiagnosisCode_7	ClmDiagnosisCode_8	\
0	NaN	2009-04-12	...	2724	19889	
1	NaN	2009-08-31	...	NaN	NaN	

```

2      PHY324689  2009-09-17 ...           NaN          NaN
3      PHY349768  2009-02-14 ...           25062        40390
4          NaN    2009-08-13 ...           5119         29620

   ClmDiagnosisCode_9  ClmDiagnosisCode_10  ClmProcedureCode_1  \
0            5849                 NaN          NaN
1            NaN                  NaN        7092.0
2            NaN                  NaN          NaN
3            4019                 NaN        331.0
4            20300                NaN        3893.0

   ClmProcedureCode_2  ClmProcedureCode_3  ClmProcedureCode_4  ClmProcedureCode_5  \
0            NaN                  NaN          NaN          NaN
1            NaN                  NaN          NaN          NaN
2            NaN                  NaN          NaN          NaN
3            NaN                  NaN          NaN          NaN
4            NaN                  NaN          NaN          NaN

   ClmProcedureCode_6
0            NaN
1            NaN
2            NaN
3            NaN
4            NaN

[5 rows x 30 columns]
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 40474 entries, 0 to 40473
Data columns (total 30 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   BeneID          40474 non-null   object 
 1   ClaimID         40474 non-null   object 
 2   ClaimStartDt    40474 non-null   object 
 3   ClaimEndDt     40474 non-null   object 
 4   Provider        40474 non-null   object 
 5   InscClaimAmtReimbursed  40474 non-null   int64  
 6   AttendingPhysician  40362 non-null   object 
 7   OperatingPhysician  23830 non-null   object 
 8   OtherPhysician    4690 non-null    object 
 9   AdmissionDt     40474 non-null   object 
 10  ClmAdmitDiagnosisCode 40474 non-null   object 
 11  DeductibleAmtPaid  39575 non-null   float64 
 12  DischargeDt     40474 non-null   object 
 13  DiagnosisGroupCode 40474 non-null   object 
 14  ClmDiagnosisCode_1 40474 non-null   object 
 15  ClmDiagnosisCode_2 40248 non-null   object 
 16  ClmDiagnosisCode_3 39798 non-null   object 

```

```

17  ClmDiagnosisCode_4      38940 non-null  object
18  ClmDiagnosisCode_5      37580 non-null  object
19  ClmDiagnosisCode_6      35636 non-null  object
20  ClmDiagnosisCode_7      33216 non-null  object
21  ClmDiagnosisCode_8      30532 non-null  object
22  ClmDiagnosisCode_9      26977 non-null  object
23  ClmDiagnosisCode_10     3927 non-null   object
24  ClmProcedureCode_1       23148 non-null  float64
25  ClmProcedureCode_2       5454 non-null   float64
26  ClmProcedureCode_3       965 non-null   float64
27  ClmProcedureCode_4       116 non-null   float64
28  ClmProcedureCode_5       9 non-null    float64
29  ClmProcedureCode_6       0 non-null    float64
dtypes: float64(7), int64(1), object(22)
memory usage: 9.3+ MB
None
BeneID      ClaimID ClaimStartDt  ClaimEndDt Provider \
0  BENE11002  CLM624349  2009-10-11  2009-10-11 PRV56011
1  BENE11003  CLM189947  2009-02-12  2009-02-12 PRV57610
2  BENE11003  CLM438021  2009-06-27  2009-06-27 PRV57595
3  BENE11004  CLM121801  2009-01-06  2009-01-06 PRV56011
4  BENE11004  CLM150998  2009-01-22  2009-01-22 PRV56011

InscClaimAmtReimbursed AttendingPhysician OperatingPhysician \
0                  30          PHY326117           NaN
1                  80          PHY362868           NaN
2                  10          PHY328821           NaN
3                  40          PHY334319           NaN
4                 200          PHY403831           NaN

OtherPhysician ClmDiagnosisCode_1 ... ClmDiagnosisCode_9 \
0            NaN        78943 ...           NaN
1            NaN        6115 ...           NaN
2            NaN        2723 ...           NaN
3            NaN        71988 ...           NaN
4            NaN        82382 ...           NaN

ClmDiagnosisCode_10 ClmProcedureCode_1 ClmProcedureCode_2 \
0            NaN        NaN           NaN
1            NaN        NaN           NaN
2            NaN        NaN           NaN
3            NaN        NaN           NaN
4            NaN        NaN           NaN

ClmProcedureCode_3 ClmProcedureCode_4 ClmProcedureCode_5 ClmProcedureCode_6 \
0            NaN        NaN           NaN           NaN
1            NaN        NaN           NaN           NaN
2            NaN        NaN           NaN           NaN

```

```

3           NaN          NaN          NaN          NaN
4           NaN          NaN          NaN          NaN

DeductibleAmtPaid  ClmAdmitDiagnosisCode
0                  0             56409
1                  0             79380
2                  0             NaN
3                  0             NaN
4                  0             71947

[5 rows x 27 columns]
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 517737 entries, 0 to 517736
Data columns (total 27 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   BeneID            517737 non-null   object 
 1   ClaimID           517737 non-null   object 
 2   ClaimStartDt      517737 non-null   object 
 3   ClaimEndDt        517737 non-null   object 
 4   Provider          517737 non-null   object 
 5   InscClaimAmtReimbursed  517737 non-null   int64  
 6   AttendingPhysician 516341 non-null   object 
 7   OperatingPhysician 90617 non-null    object 
 8   OtherPhysician     195046 non-null   object 
 9   ClmDiagnosisCode_1 507284 non-null   object 
 10  ClmDiagnosisCode_2 322357 non-null   object 
 11  ClmDiagnosisCode_3 203257 non-null   object 
 12  ClmDiagnosisCode_4 125596 non-null   object 
 13  ClmDiagnosisCode_5 74344 non-null    object 
 14  ClmDiagnosisCode_6 48756 non-null    object 
 15  ClmDiagnosisCode_7 32961 non-null    object 
 16  ClmDiagnosisCode_8 22912 non-null    object 
 17  ClmDiagnosisCode_9 14838 non-null    object 
 18  ClmDiagnosisCode_10 1083 non-null    object 
 19  ClmProcedureCode_1 162 non-null     float64 
 20  ClmProcedureCode_2 36 non-null     float64 
 21  ClmProcedureCode_3 4 non-null     float64 
 22  ClmProcedureCode_4 2 non-null     float64 
 23  ClmProcedureCode_5 0 non-null     float64 
 24  ClmProcedureCode_6 0 non-null     float64 
 25  DeductibleAmtPaid 517737 non-null   int64  
 26  ClmAdmitDiagnosisCode 105425 non-null   object 

dtypes: float64(6), int64(2), object(19)
memory usage: 106.7+ MB
None

```

Check missing values ##### Physician column has lots of NaN values ##### Need proper handling before analysis

```
[16]: # Count NaN values per column
print("Inpatient NaNs:\n", df2.isna().sum())
print("Outpatient NaNs:\n", df3.isna().sum())
```

Inpatient NaNs:

```
BeneID          0
ClaimID         0
ClaimStartDt    0
ClaimEndDt     0
Provider        0
InscClaimAmtReimbursed 0
AttendingPhysician 112
OperatingPhysician 16644
OtherPhysician   35784
AdmissionDt     0
ClmAdmitDiagnosisCode 0
DeductibleAmtPaid 899
DischargeDt     0
DiagnosisGroupCode 0
ClmDiagnosisCode_1 0
ClmDiagnosisCode_2 226
ClmDiagnosisCode_3 676
ClmDiagnosisCode_4 1534
ClmDiagnosisCode_5 2894
ClmDiagnosisCode_6 4838
ClmDiagnosisCode_7 7258
ClmDiagnosisCode_8 9942
ClmDiagnosisCode_9 13497
ClmDiagnosisCode_10 36547
ClmProcedureCode_1 17326
ClmProcedureCode_2 35020
ClmProcedureCode_3 39509
ClmProcedureCode_4 40358
ClmProcedureCode_5 40465
ClmProcedureCode_6 40474
dtype: int64
```

Outpatient NaNs:

```
BeneID          0
ClaimID         0
ClaimStartDt    0
ClaimEndDt     0
Provider        0
InscClaimAmtReimbursed 0
AttendingPhysician 1396
OperatingPhysician 427120
```

```
OtherPhysician           322691
ClmDiagnosisCode_1      10453
ClmDiagnosisCode_2      195380
ClmDiagnosisCode_3      314480
ClmDiagnosisCode_4      392141
ClmDiagnosisCode_5      443393
ClmDiagnosisCode_6      468981
ClmDiagnosisCode_7      484776
ClmDiagnosisCode_8      494825
ClmDiagnosisCode_9      502899
ClmDiagnosisCode_10     516654
ClmProcedureCode_1       517575
ClmProcedureCode_2       517701
ClmProcedureCode_3       517733
ClmProcedureCode_4       517735
ClmProcedureCode_5       517737
ClmProcedureCode_6       517737
DeductibleAmtPaid        0
ClmAdmitDiagnosisCode   412312
dtype: int64
```

```
[17]: print(df2.columns)
      print(df3.columns)
```

```
Index(['BeneID', 'ClaimID', 'ClaimStartDt', 'ClaimEndDt', 'Provider',
       'InscClaimAmtReimbursed', 'AttendingPhysician', 'OperatingPhysician',
       'OtherPhysician', 'AdmissionDt', 'ClmAdmitDiagnosisCode',
       'DeductibleAmtPaid', 'DischargeDt', 'DiagnosisGroupCode',
       'ClmDiagnosisCode_1', 'ClmDiagnosisCode_2', 'ClmDiagnosisCode_3',
       'ClmDiagnosisCode_4', 'ClmDiagnosisCode_5', 'ClmDiagnosisCode_6',
       'ClmDiagnosisCode_7', 'ClmDiagnosisCode_8', 'ClmDiagnosisCode_9',
       'ClmDiagnosisCode_10', 'ClmProcedureCode_1', 'ClmProcedureCode_2',
       'ClmProcedureCode_3', 'ClmProcedureCode_4', 'ClmProcedureCode_5',
       'ClmProcedureCode_6'],
      dtype='object')
Index(['BeneID', 'ClaimID', 'ClaimStartDt', 'ClaimEndDt', 'Provider',
       'InscClaimAmtReimbursed', 'AttendingPhysician', 'OperatingPhysician',
       'OtherPhysician', 'ClmDiagnosisCode_1', 'ClmDiagnosisCode_2',
       'ClmDiagnosisCode_3', 'ClmDiagnosisCode_4', 'ClmDiagnosisCode_5',
       'ClmDiagnosisCode_6', 'ClmDiagnosisCode_7', 'ClmDiagnosisCode_8',
       'ClmDiagnosisCode_9', 'ClmDiagnosisCode_10', 'ClmProcedureCode_1',
       'ClmProcedureCode_2', 'ClmProcedureCode_3', 'ClmProcedureCode_4',
       'ClmProcedureCode_5', 'ClmProcedureCode_6', 'DeductibleAmtPaid',
       'ClmAdmitDiagnosisCode'],
      dtype='object')
```

```
[18]: df2.head()
```

```
[18]:      BeneID  ClaimID ClaimStartDt  ClaimEndDt  Provider \
0  BENE11001  CLM46614    2009-04-12  2009-04-18  PRV55912
1  BENE11001  CLM66048    2009-08-31  2009-09-02  PRV55907
2  BENE11001  CLM68358    2009-09-17  2009-09-20  PRV56046
3  BENE11011  CLM38412    2009-02-14  2009-02-22  PRV52405
4  BENE11014  CLM63689    2009-08-13  2009-08-30  PRV56614

      InscClaimAmtReimbursed AttendingPhysician OperatingPhysician \
0                  26000          PHY390922           NaN
1                  5000           PHY318495          PHY318495
2                  5000           PHY372395           NaN
3                  5000           PHY369659          PHY392961
4                 10000          PHY379376          PHY398258

      OtherPhysician AdmissionDt ... ClmDiagnosisCode_7  ClmDiagnosisCode_8 \
0            NaN  2009-04-12 ...          2724           19889
1            NaN  2009-08-31 ...           NaN           NaN
2  PHY324689  2009-09-17 ...           NaN           NaN
3  PHY349768  2009-02-14 ...          25062          40390
4            NaN  2009-08-13 ...          5119           29620

      ClmDiagnosisCode_9  ClmDiagnosisCode_10  ClmProcedureCode_1 \
0                  5849             NaN           NaN
1                  NaN             NaN          7092.0
2                  NaN             NaN           NaN
3                  4019             NaN          331.0
4                 20300            NaN          3893.0

      ClmProcedureCode_2  ClmProcedureCode_3  ClmProcedureCode_4  ClmProcedureCode_5 \
0                  NaN             NaN           NaN           NaN
1                  NaN             NaN           NaN           NaN
2                  NaN             NaN           NaN           NaN
3                  NaN             NaN           NaN           NaN
4                  NaN             NaN           NaN           NaN

      ClmProcedureCode_6
0                  NaN
1                  NaN
2                  NaN
3                  NaN
4                  NaN

[5 rows x 30 columns]
```

```
[19]: df3.head()
```

```
[19]:      BeneID    ClaimID ClaimStartDt  ClaimEndDt  Provider \
0  BENE11002  CLM624349    2009-10-11  2009-10-11  PRV56011
1  BENE11003  CLM189947    2009-02-12  2009-02-12  PRV57610
2  BENE11003  CLM438021    2009-06-27  2009-06-27  PRV57595
3  BENE11004  CLM121801    2009-01-06  2009-01-06  PRV56011
4  BENE11004  CLM150998    2009-01-22  2009-01-22  PRV56011

      InscClaimAmtReimbursed AttendingPhysician OperatingPhysician \
0                  30          PHY326117           NaN
1                  80          PHY362868           NaN
2                  10          PHY328821           NaN
3                  40          PHY334319           NaN
4                 200          PHY403831           NaN

      OtherPhysician ClmDiagnosisCode_1 ... ClmDiagnosisCode_9 \
0            NaN          78943   ...
1            NaN          6115    ...
2            NaN          2723   ...
3            NaN          71988   ...
4            NaN          82382   ...

      ClmDiagnosisCode_10 ClmProcedureCode_1 ClmProcedureCode_2 \
0            NaN           NaN           NaN
1            NaN           NaN           NaN
2            NaN           NaN           NaN
3            NaN           NaN           NaN
4            NaN           NaN           NaN

      ClmProcedureCode_3 ClmProcedureCode_4 ClmProcedureCode_5 ClmProcedureCode_6 \
0            NaN           NaN           NaN           NaN
1            NaN           NaN           NaN           NaN
2            NaN           NaN           NaN           NaN
3            NaN           NaN           NaN           NaN
4            NaN           NaN           NaN           NaN

      DeductibleAmtPaid  ClmAdmitDiagnosisCode
0                  0           56409
1                  0           79380
2                  0           NaN
3                  0           NaN
4                  0           71947

[5 rows x 27 columns]
```

```
[20]: # Check total NaNs per column
print("Inpatient NaNs:\n", df2.isna().sum())
print("Outpatient NaNs:\n", df3.isna().sum())
```

Inpatient NaNs:

BeneID	0
ClaimID	0
ClaimStartDt	0
ClaimEndDt	0
Provider	0
InscClaimAmtReimbursed	0
AttendingPhysician	112
OperatingPhysician	16644
OtherPhysician	35784
AdmissionDt	0
ClmAdmitDiagnosisCode	0
DeductibleAmtPaid	899
DischargeDt	0
DiagnosisGroupCode	0
ClmDiagnosisCode_1	0
ClmDiagnosisCode_2	226
ClmDiagnosisCode_3	676
ClmDiagnosisCode_4	1534
ClmDiagnosisCode_5	2894
ClmDiagnosisCode_6	4838
ClmDiagnosisCode_7	7258
ClmDiagnosisCode_8	9942
ClmDiagnosisCode_9	13497
ClmDiagnosisCode_10	36547
ClmProcedureCode_1	17326
ClmProcedureCode_2	35020
ClmProcedureCode_3	39509
ClmProcedureCode_4	40358
ClmProcedureCode_5	40465
ClmProcedureCode_6	40474

dtype: int64

Outpatient NaNs:

BeneID	0
ClaimID	0
ClaimStartDt	0
ClaimEndDt	0
Provider	0
InscClaimAmtReimbursed	0
AttendingPhysician	1396
OperatingPhysician	427120
OtherPhysician	322691
ClmDiagnosisCode_1	10453
ClmDiagnosisCode_2	195380
ClmDiagnosisCode_3	314480
ClmDiagnosisCode_4	392141
ClmDiagnosisCode_5	443393
ClmDiagnosisCode_6	468981

```

ClmDiagnosisCode_7      484776
ClmDiagnosisCode_8      494825
ClmDiagnosisCode_9      502899
ClmDiagnosisCode_10     516654
ClmProcedureCode_1       517575
ClmProcedureCode_2       517701
ClmProcedureCode_3       517733
ClmProcedureCode_4       517735
ClmProcedureCode_5       517737
ClmProcedureCode_6       517737
DeductibleAmtPaid        0
ClmAdmitDiagnosisCode   412312
dtype: int64

```

Attending Physicians (Inpatient Claims) Highly skewed distribution: Top physicians handle majority of inpatient claims

Helps identify physicians with high workload or potential outlier behavior

Useful for fraud detection → unusually high number of visits by some physicians

Quick comparison among top 10 physicians → visually easy to interpret

Inspect Columns

```
[21]: print(df2[['AttendingPhysician', 'OperatingPhysician']].isna().sum())
```

```

AttendingPhysician      112
OperatingPhysician     16644
dtype: int64

```

```
[22]: df2['AttendingPhysician'] = df2['AttendingPhysician'].fillna('Unknown')
df2['OperatingPhysician'] = df2['OperatingPhysician'].fillna('Unknown')
```

```
df3['AttendingPhysician'] = df3['AttendingPhysician'].fillna('Unknown')
df3['OperatingPhysician'] = df3['OperatingPhysician'].fillna('Unknown')
```

```
[23]: df2['AttendingPhysician'] = df2.groupby('BeneID')['AttendingPhysician'].ffill()
      ↵ bfill()
df2['OperatingPhysician'] = df2.groupby('BeneID')['OperatingPhysician'].ffill()
      ↵ bfill()

df3['AttendingPhysician'] = df3.groupby('BeneID')['AttendingPhysician'].ffill()
      ↵ bfill()
df3['OperatingPhysician'] = df3.groupby('BeneID')['OperatingPhysician'].ffill()
      ↵ bfill()
```

Plot Description: Top 10 Attending Physicians (Inpatient) ##### The top 10 attending physicians by inpatient visits are responsible for a large portion of patient care. While high patient

volume is normal, these physicians' claims should be monitored closely for anomalies, as high-volume providers are more likely to generate high-value claims that may require fraud scrutiny.

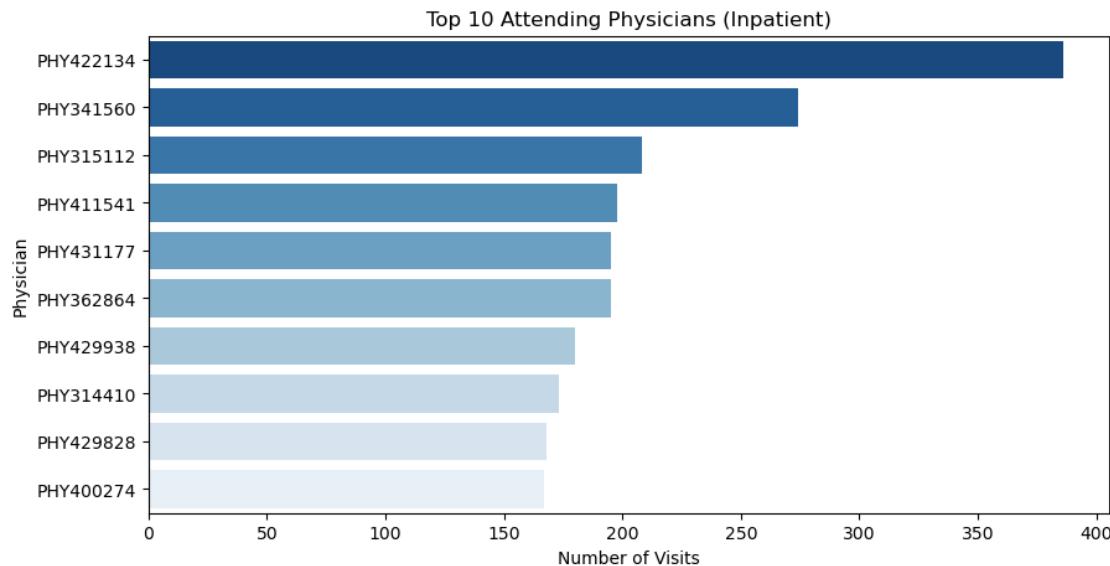
```
[24]: import matplotlib.pyplot as plt
import seaborn as sns

top_attending = df2['AttendingPhysician'].value_counts().head(10)
plt.figure(figsize=(10,5))
sns.barplot(x=top_attending.values, y=top_attending.index, palette='Blues_r')
plt.title("Top 10 Attending Physicians (Inpatient)")
plt.xlabel("Number of Visits")
plt.ylabel("Physician")
plt.show()
```

C:\Users\arft\AppData\Local\Temp\ipykernel_20176\3345240337.py:6: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=top_attending.values, y=top_attending.index, palette='Blues_r')
```



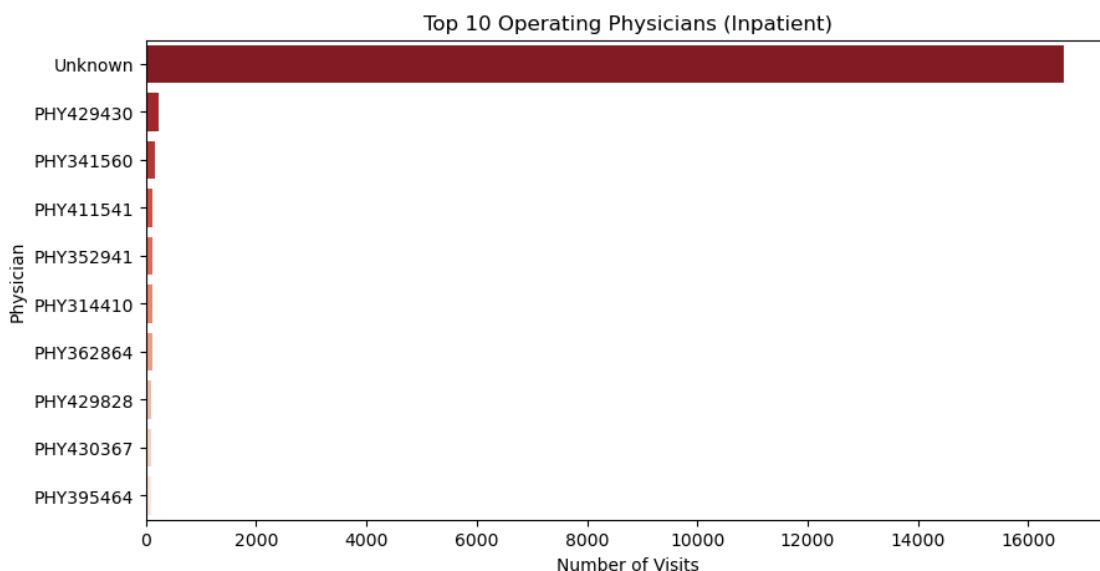
Plot Description: Top 10 Operating Physicians (Inpatient) The top 10 operating physicians account for a significant share of inpatient surgical activity. While high procedure volume is expected for experienced surgeons, their claims should be monitored for anomalies, as frequent surgeries may result in high-value claims and potential fraud risk.

```
[26]: top_operating = df2['OperatingPhysician'].value_counts().head(10)
plt.figure(figsize=(10,5))
sns.barplot(x=top_operating.values, y=top_operating.index, palette='Reds_r')
plt.title("Top 10 Operating Physicians (Inpatient)")
plt.xlabel("Number of Visits")
plt.ylabel("Physician")
plt.show()
```

C:\Users\arft\AppData\Local\Temp\ipykernel_20176\2401160898.py:3: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=top_operating.values, y=top_operating.index, palette='Reds_r')
```



```
[27]: df2.head()
```

```
[27]:      BeneID  ClaimID ClaimStartDt  ClaimEndDt  Provider \
0  BENE11001  CLM46614  2009-04-12  2009-04-18  PRV55912
1  BENE11001  CLM66048  2009-08-31  2009-09-02  PRV55907
2  BENE11001  CLM68358  2009-09-17  2009-09-20  PRV56046
3  BENE11011  CLM38412  2009-02-14  2009-02-22  PRV52405
4  BENE11014  CLM63689  2009-08-13  2009-08-30  PRV56614

      InscClaimAmtReimbursed  AttendingPhysician  OperatingPhysician \
0                      26000          PHY390922           Unknown
1                      5000          PHY318495          PHY318495
```

```

2           5000      PHY372395        Unknown
3           5000      PHY369659      PHY392961
4          10000      PHY379376      PHY398258

  OtherPhysician AdmissionDt ... ClmDiagnosisCode_7 ClmDiagnosisCode_8 \
0           NaN 2009-04-12 ...       2724        19889
1           NaN 2009-08-31 ...        NaN         NaN
2      PHY324689 2009-09-17 ...        NaN         NaN
3      PHY349768 2009-02-14 ...       25062       40390
4           NaN 2009-08-13 ...       5119        29620

  ClmDiagnosisCode_9 ClmDiagnosisCode_10 ClmProcedureCode_1 \
0           5849        NaN        NaN
1           NaN        NaN     7092.0
2           NaN        NaN        NaN
3           4019        NaN     331.0
4           20300       NaN     3893.0

  ClmProcedureCode_2 ClmProcedureCode_3 ClmProcedureCode_4 ClmProcedureCode_5 \
0           NaN        NaN        NaN        NaN
1           NaN        NaN        NaN        NaN
2           NaN        NaN        NaN        NaN
3           NaN        NaN        NaN        NaN
4           NaN        NaN        NaN        NaN

  ClmProcedureCode_6
0           NaN
1           NaN
2           NaN
3           NaN
4           NaN

[5 rows x 30 columns]

```

```
[28]: phys_overlap = df2.groupby(['AttendingPhysician', 'OperatingPhysician']).size() .
    ↪reset_index(name='Count')
phys_overlap.sort_values('Count', ascending=False).head(10)
```

```
[28]:   AttendingPhysician OperatingPhysician  Count
16398      PHY422134      PHY429430    225
16399      PHY422134        Unknown    161
4495       PHY341560      PHY341560    153
14885       PHY411541      PHY411541    121
4496       PHY341560        Unknown    121
17779       PHY431177      PHY352941    110
437        PHY314410      PHY314410    109
544        PHY315112        Unknown    108
```

7766	PHY362864	PHY362864	107
18274	Unknown	Unknown	106

```
[29]: physician_cols = [
    'Physician',
    'AttendingPhysician',
    'OperatingPhysician'
]
```

```
[30]: for col in physician_cols:
    if col in df2.columns:
        df2[col] = df2[col].fillna('Unknown')
    if col in df3.columns:
        df3[col] = df3[col].fillna('Unknown')
```

```
[31]: print(df2['AttendingPhysician'].isna().sum())
print(df3['OperatingPhysician'].isna().sum())
```

0
0

```
[32]: print(df2.columns)
print(df3.columns)
```

```
Index(['BeneID', 'ClaimID', 'ClaimStartDt', 'ClaimEndDt', 'Provider',
       'InscClaimAmtReimbursed', 'AttendingPhysician', 'OperatingPhysician',
       'OtherPhysician', 'AdmissionDt', 'ClmAdmitDiagnosisCode',
       'DeductibleAmtPaid', 'DischargeDt', 'DiagnosisGroupCode',
       'ClmDiagnosisCode_1', 'ClmDiagnosisCode_2', 'ClmDiagnosisCode_3',
       'ClmDiagnosisCode_4', 'ClmDiagnosisCode_5', 'ClmDiagnosisCode_6',
       'ClmDiagnosisCode_7', 'ClmDiagnosisCode_8', 'ClmDiagnosisCode_9',
       'ClmDiagnosisCode_10', 'ClmProcedureCode_1', 'ClmProcedureCode_2',
       'ClmProcedureCode_3', 'ClmProcedureCode_4', 'ClmProcedureCode_5',
       'ClmProcedureCode_6'],
      dtype='object')
Index(['BeneID', 'ClaimID', 'ClaimStartDt', 'ClaimEndDt', 'Provider',
       'InscClaimAmtReimbursed', 'AttendingPhysician', 'OperatingPhysician',
       'OtherPhysician', 'ClmDiagnosisCode_1', 'ClmDiagnosisCode_2',
       'ClmDiagnosisCode_3', 'ClmDiagnosisCode_4', 'ClmDiagnosisCode_5',
       'ClmDiagnosisCode_6', 'ClmDiagnosisCode_7', 'ClmDiagnosisCode_8',
       'ClmDiagnosisCode_9', 'ClmDiagnosisCode_10', 'ClmProcedureCode_1',
       'ClmProcedureCode_2', 'ClmProcedureCode_3', 'ClmProcedureCode_4',
       'ClmProcedureCode_5', 'ClmProcedureCode_6', 'DeductibleAmtPaid',
       'ClmAdmitDiagnosisCode'],
      dtype='object')
```

```
[33]: for col in ['AttendingPhysician', 'OperatingPhysician']:
    if col in df2.columns:
```

```

    print(col, "df2 NaNs:", df2[col].isna().sum())
if col in df3.columns:
    print(col, "df3 NaNs:", df3[col].isna().sum())

```

AttendingPhysician df2 NaNs: 0
 AttendingPhysician df3 NaNs: 0
 OperatingPhysician df2 NaNs: 0
 OperatingPhysician df3 NaNs: 0

[34]: df2['AttendingPhysician'].value_counts().head(10)
 df3['OperatingPhysician'].value_counts().head(10)

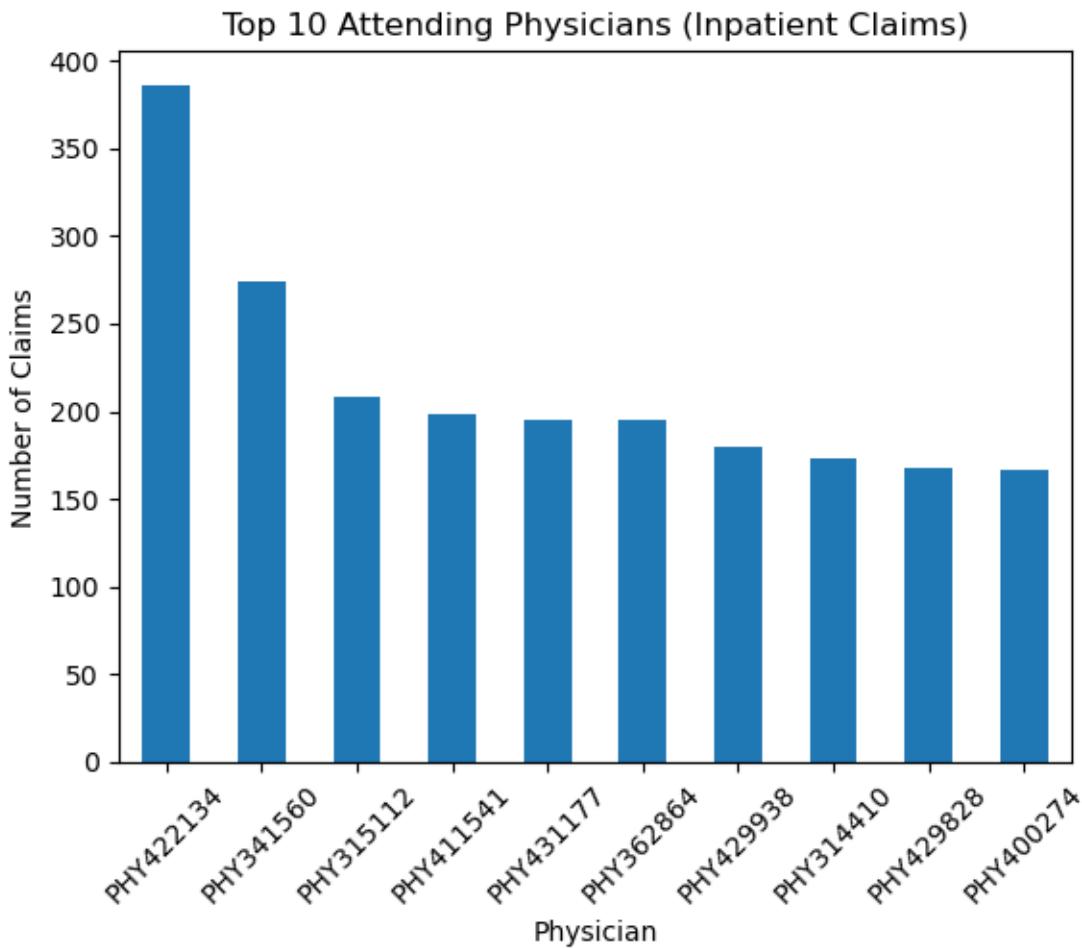
[34]: OperatingPhysician
 Unknown 427120
 PHY330576 424
 PHY424897 293
 PHY314027 256
 PHY423534 250
 PHY357120 249
 PHY412132 245
 PHY327046 236
 PHY333735 232
 PHY381249 231
 Name: count, dtype: int64

Attending Physicians

Helps identify provider-level patterns, possible fraud investigation or resource allocation
 Distribution shows skewed workload among physicians (top 10 handle majority claims)

[35]: import matplotlib.pyplot as plt

 df2['AttendingPhysician'].value_counts().head(10).plot(kind='bar')
 plt.title("Top 10 Attending Physicians (Inpatient Claims)")
 plt.ylabel("Number of Claims")
 plt.xlabel("Physician")
 plt.xticks(rotation=45)
 plt.show()



Plot Description: Top 10 Operating Physicians (Outpatient Claims)

A small group of physicians accounts for a disproportionately high number of outpatient claims.

While a high claim count may indicate a large patient volume, unusually high values compared to peers can be a potential red flag.

Fraud Detection Perspective: Physicians with exceptionally high claim volumes may be associated with:

Overutilization of services

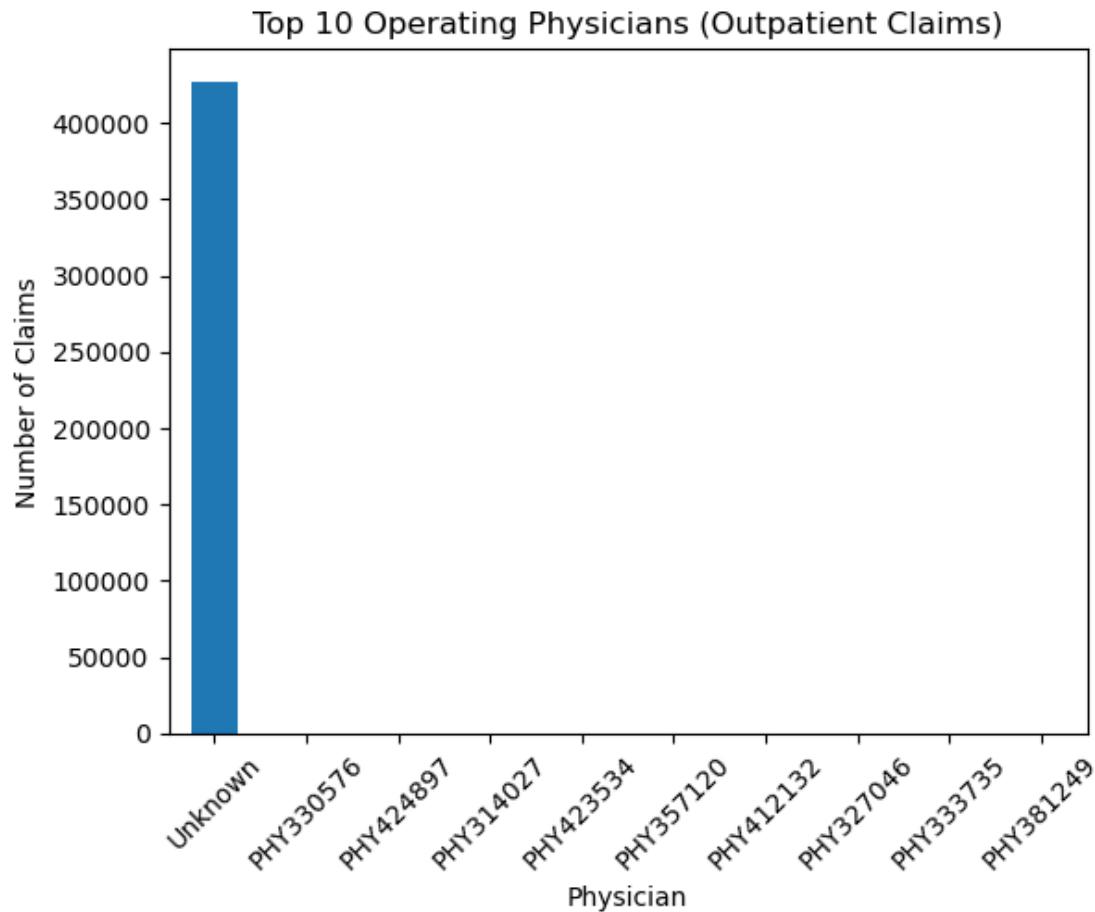
Unnecessary or repetitive procedures

Upcoding or unbundling practices

These physicians should be prioritized for further investigation.

```
[36]: df3['OperatingPhysician'].value_counts().head(10).plot(kind='bar')
plt.title("Top 10 Operating Physicians (Outpatient Claims)")
```

```
plt.ylabel("Number of Claims")
plt.xlabel("Physician")
plt.xticks(rotation=45)
plt.show()
```



```
[37]: cat_cols_df2 = df2.select_dtypes(include='object').columns
cat_cols_df3 = df3.select_dtypes(include='object').columns

df2[cat_cols_df2] = df2[cat_cols_df2].fillna('Unknown')
df3[cat_cols_df3] = df3[cat_cols_df3].fillna('Unknown')
```

```
[38]: num_cols_df2 = df2.select_dtypes(include=['int64','float64']).columns
num_cols_df3 = df3.select_dtypes(include=['int64','float64']).columns

df2[num_cols_df2] = df2[num_cols_df2].fillna(0)
df3[num_cols_df3] = df3[num_cols_df3].fillna(0)
```

```
[39]: date_cols = ['AdmissionDt', 'DischargeDt', 'ClaimStartDt', 'ClaimEndDt']
```

```
for col in date_cols:  
    if col in df2.columns:  
        df2[col] = pd.to_datetime(df2[col], errors='coerce')  
    if col in df3.columns:  
        df3[col] = pd.to_datetime(df3[col], errors='coerce')
```

```
[40]: print("DF2 Remaining NaNs:\n", df2.isna().sum().sum())  
print("DF3 Remaining NaNs:\n", df3.isna().sum().sum())
```

DF2 Remaining NaNs:

0

DF3 Remaining NaNs:

0

```
[41]: (df2 == 'Unknown').sum().sort_values(ascending=False)
```

ClmDiagnosisCode_10	36547
OtherPhysician	35784
OperatingPhysician	16644
ClmDiagnosisCode_9	13497
ClmDiagnosisCode_8	9942
ClmDiagnosisCode_7	7258
ClmDiagnosisCode_6	4838
ClmDiagnosisCode_5	2894
ClmDiagnosisCode_4	1534
ClmDiagnosisCode_3	676
ClmDiagnosisCode_2	226
AttendingPhysician	112
BeneID	0
ClaimStartDt	0
DiagnosisGroupCode	0
DischargeDt	0
DeductibleAmtPaid	0
ClmAdmitDiagnosisCode	0
AdmissionDt	0
ClaimEndDt	0
Provider	0
InscClaimAmtReimbursed	0
ClaimID	0
ClmDiagnosisCode_1	0
ClmProcedureCode_1	0
ClmProcedureCode_2	0
ClmProcedureCode_3	0
ClmProcedureCode_4	0
ClmProcedureCode_5	0
ClmProcedureCode_6	0

```
dtype: int64
```

Handling Missing Claim Information Missing values in claim-related fields were handled based on data semantics.

Reimbursement and deductible amounts were imputed with 0, indicating no payment.

Diagnosis and procedure codes were filled with ‘Not Reported’, while claim date fields were converted to datetime format with invalid entries coerced to NaT

```
[42]: df2.isna().sum().sort_values(ascending=False).head(15)
```

```
[42]: BeneID          0
      ClaimID         0
      ClaimStartDt    0
      ClaimEndDt      0
      Provider         0
      InscClaimAmtReimbursed 0
      AttendingPhysician 0
      OperatingPhysician 0
      OtherPhysician    0
      AdmissionDt      0
      ClmAdmitDiagnosisCode 0
      DeductibleAmtPaid 0
      DischargeDt       0
      DiagnosisGroupCode 0
      ClmDiagnosisCode_1 0
      dtype: int64
```

```
[44]: claim_amount_cols = [
        'InscClaimAmtReimbursed',
        'DeductibleAmtPaid'
    ]

    for col in claim_amount_cols:
        if col in df2.columns:
            df2[col] = df2[col].fillna(0)
        if col in df3.columns:
            df3[col] = df3[col].fillna(0)
```

```
[45]: claim_cat_cols = [
        'ClmAdmitDiagnosisCode',
        'ClmProcedureCode'
    ]

    for col in claim_cat_cols:
        if col in df2.columns:
            df2[col] = df2[col].fillna('Not Reported')
```

```

        if col in df3.columns:
            df3[col] = df3[col].fillna('Not Reported')

[46]: date_cols = ['ClaimStartDt', 'ClaimEndDt']

for col in date_cols:
    if col in df2.columns:
        df2[col] = pd.to_datetime(df2[col], errors='coerce')
    if col in df3.columns:
        df3[col] = pd.to_datetime(df3[col], errors='coerce')

[47]: print("DF2 remaining NaNs:\n", df2.isna().sum().sort_values(ascending=False).
           head())
print("DF3 remaining NaNs:\n", df3.isna().sum().sort_values(ascending=False).
           head())

DF2 remaining NaNs:
BeneID      0
ClaimID     0
ClaimStartDt 0
ClaimEndDt   0
Provider     0
dtype: int64
DF3 remaining NaNs:
BeneID      0
ClaimID     0
ClaimStartDt 0
ClaimEndDt   0
Provider     0
dtype: int64

[48]: df2.head()

[48]:      BeneID  ClaimID ClaimStartDt ClaimEndDt  Provider \
0  BENE11001  CLM46614  2009-04-12 2009-04-18  PRV55912
1  BENE11001  CLM66048  2009-08-31 2009-09-02  PRV55907
2  BENE11001  CLM68358  2009-09-17 2009-09-20  PRV56046
3  BENE11011  CLM38412  2009-02-14 2009-02-22  PRV52405
4  BENE11014  CLM63689  2009-08-13 2009-08-30  PRV56614

      InscClaimAmtReimbursed AttendingPhysician OperatingPhysician \
0                  26000          PHY390922           Unknown
1                  5000          PHY318495           PHY318495
2                  5000          PHY372395           Unknown
3                  5000          PHY369659           PHY392961
4                 10000          PHY379376           PHY398258

OtherPhysician AdmissionDt ... ClmDiagnosisCode_7  ClmDiagnosisCode_8 \

```

```

0      Unknown 2009-04-12 ...
1      Unknown 2009-08-31 ...
2  PHY324689 2009-09-17 ...
3  PHY349768 2009-02-14 ...
4      Unknown 2009-08-13 ...

ClmDiagnosisCode_9 ClmDiagnosisCode_10 ClmProcedureCode_1 \
0          5849      Unknown      0.0
1      Unknown      Unknown  7092.0
2      Unknown      Unknown      0.0
3          4019      Unknown   331.0
4        20300      Unknown  3893.0

ClmProcedureCode_2 ClmProcedureCode_3 ClmProcedureCode_4 ClmProcedureCode_5 \
0          0.0      0.0      0.0      0.0
1          0.0      0.0      0.0      0.0
2          0.0      0.0      0.0      0.0
3          0.0      0.0      0.0      0.0
4          0.0      0.0      0.0      0.0

ClmProcedureCode_6
0          0.0
1          0.0
2          0.0
3          0.0
4          0.0

[5 rows x 30 columns]

```

```
[112]: (df2['InscClaimAmtReimbursed'] == 0).sum()
(df3['InscClaimAmtReimbursed'] == 0).sum()
```

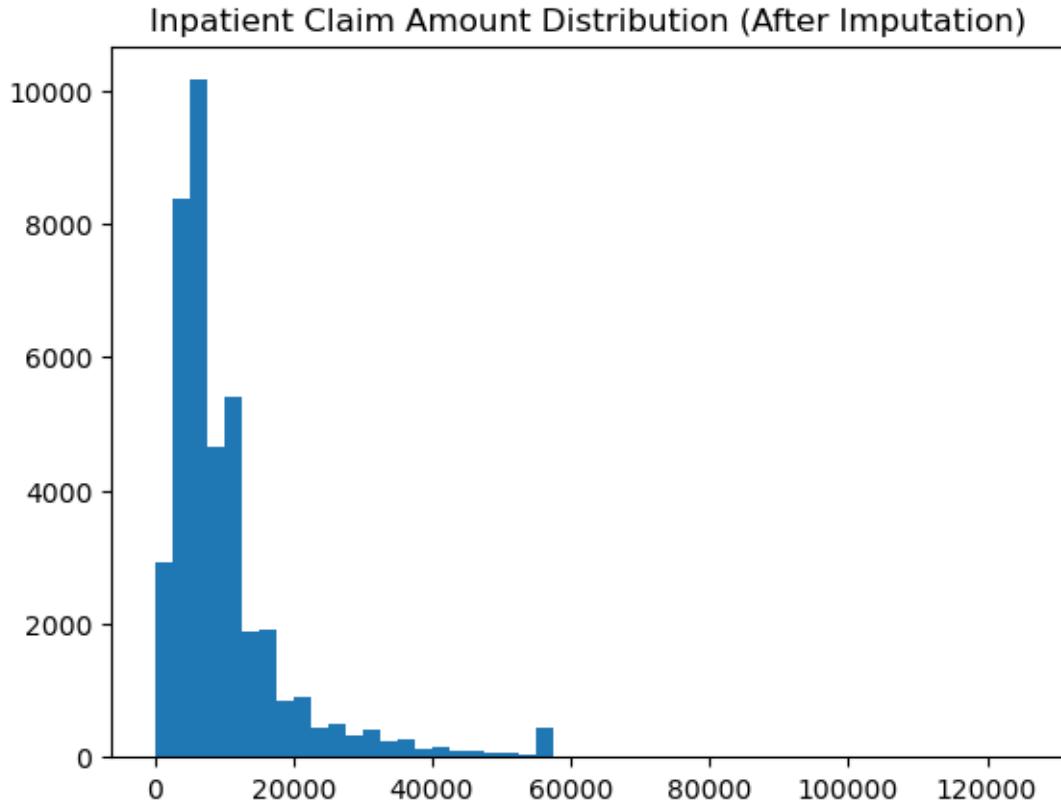
```
[112]: np.int64(19568)
```

Inpatient Claim Amount Distribution (After Imputation) The distribution is right-skewed, meaning most inpatient claims are concentrated at lower reimbursement amounts, while a smaller number of claims have very high values.

The long tail on the right indicates the presence of high-cost inpatient cases.

Imputation has preserved the overall shape of the distribution, suggesting that missing value treatment did not distort the data

```
[50]: plt.hist(df2['InscClaimAmtReimbursed'], bins=50)
plt.title("Inpatient Claim Amount Distribution (After Imputation)")
plt.show()
```



Inpatient Claim Amount Distribution (After Imputation)

Most Inpatient claims fall within lower to mid-range amounts, showing common reimbursement levels

Some high-value claims visible → indicates major procedures or expensive treatments

Distribution is right-skewed → few very high claims, majority moderate

After imputation, no zero or missing values distort the distribution

Useful for fraud detection, identifying outlier claims, and resource planning

```
[89]: df2['InscClaimAmtReimbursed'].replace(0, pd.NA, inplace=True)
df3['InscClaimAmtReimbursed'].replace(0, pd.NA, inplace=True)
```

C:\Users\arft\AppData\Local\Temp\ipykernel_20176\965939615.py:1: FutureWarning:
A value is trying to be set on a copy of a DataFrame or Series through chained
assignment using an `inplace` method.

The behavior will change in pandas 3.0. This `inplace` method will never work
because the intermediate object on which we are setting values always behaves as
a copy.

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df2['InscClaimAmtReimbursed'].replace(0, pd.NA, inplace=True)
C:\Users\arft\AppData\Local\Temp\ipykernel_20176\965939615.py:2: FutureWarning:
A value is trying to be set on a copy of a DataFrame or Series through chained
assignment using an inplace method.
The behavior will change in pandas 3.0. This inplace method will never work
because the intermediate object on which we are setting values always behaves as
a copy.
```

For example, when doing 'df[col].method(value, inplace=True)', try using 'df.method({col: value}, inplace=True)' or df[col] = df[col].method(value) instead, to perform the operation inplace on the original object.

```
df3['InscClaimAmtReimbursed'].replace(0, pd.NA, inplace=True)

[52]: df2['InscClaimAmtReimbursed'] = df2.
    ↪groupby('AttendingPhysician')['InscClaimAmtReimbursed']\
        .transform(lambda x: x.fillna(x.mean()))

df3['InscClaimAmtReimbursed'] = df3.
    ↪groupby('OperatingPhysician')['InscClaimAmtReimbursed']\
        .transform(lambda x: x.fillna(x.mean()))
```

```
C:\Users\arft\AppData\Local\Temp\ipykernel_20176\4095635132.py:2: FutureWarning:
Downcasting object dtype arrays on .fillna, .ffill, .bfill is deprecated and
will change in a future version. Call result.infer_objects(copy=False) instead.
To opt-in to the future behavior, set
`pd.set_option('future.no_silent_downcasting', True)`
    .transform(lambda x: x.fillna(x.mean()))
C:\Users\arft\AppData\Local\Temp\ipykernel_20176\4095635132.py:5: FutureWarning:
Downcasting object dtype arrays on .fillna, .ffill, .bfill is deprecated and
will change in a future version. Call result.infer_objects(copy=False) instead.
To opt-in to the future behavior, set
`pd.set_option('future.no_silent_downcasting', True)`
    .transform(lambda x: x.fillna(x.mean()))
```

```
[90]: print(df2.columns)
print(df3.columns)

Index(['BeneID', 'ClaimID', 'ClaimStartDt', 'ClaimEndDt', 'Provider',
       'InscClaimAmtReimbursed', 'AttendingPhysician', 'OperatingPhysician',
       'OtherPhysician', 'AdmissionDt', 'ClmAdmitDiagnosisCode',
```

```

'DeductibleAmtPaid', 'DischargeDt', 'DiagnosisGroupCode',
'ClmDiagnosisCode_1', 'ClmDiagnosisCode_2', 'ClmDiagnosisCode_3',
'ClmDiagnosisCode_4', 'ClmDiagnosisCode_5', 'ClmDiagnosisCode_6',
'ClmDiagnosisCode_7', 'ClmDiagnosisCode_8', 'ClmDiagnosisCode_9',
'ClmDiagnosisCode_10', 'ClmProcedureCode_1', 'ClmProcedureCode_2',
'ClmProcedureCode_3', 'ClmProcedureCode_4', 'ClmProcedureCode_5',
'ClmProcedureCode_6'],
dtype='object')
Index(['BeneID', 'ClaimID', 'ClaimStartDt', 'ClaimEndDt', 'Provider',
       'InscClaimAmtReimbursed', 'AttendingPhysician', 'OperatingPhysician',
       'OtherPhysician', 'ClmDiagnosisCode_1', 'ClmDiagnosisCode_2',
       'ClmDiagnosisCode_3', 'ClmDiagnosisCode_4', 'ClmDiagnosisCode_5',
       'ClmDiagnosisCode_6', 'ClmDiagnosisCode_7', 'ClmDiagnosisCode_8',
       'ClmDiagnosisCode_9', 'ClmDiagnosisCode_10', 'ClmProcedureCode_1',
       'ClmProcedureCode_2', 'ClmProcedureCode_3', 'ClmProcedureCode_4',
       'ClmProcedureCode_5', 'ClmProcedureCode_6', 'DeductibleAmtPaid',
       'ClmAdmitDiagnosisCode'],
      dtype='object')

```

```
[91]: df2['AttendingPhysician'] = df2['AttendingPhysician'].str.strip()
df3['OperatingPhysician'] = df3['OperatingPhysician'].str.strip()
```

```
[92]: df2['AttendingPhysician'] = df2['AttendingPhysician'].str.lower()
df3['OperatingPhysician'] = df3['OperatingPhysician'].str.lower()
```

```
[93]: df2['AttendingPhysician'] = df2['AttendingPhysician'].replace('unknown', pd.NA)
df3['OperatingPhysician'] = df3['OperatingPhysician'].replace('unknown', pd.NA)
```

```
[94]: print(df2['AttendingPhysician'].isna().sum())
print(df3['OperatingPhysician'].isna().sum())
```

112
427120

```
[95]: attending_count = df2['AttendingPhysician'].nunique()
operating_count = df3['OperatingPhysician'].nunique()

print("Number of unique Attending Physicians (Inpatient):", attending_count)
print("Number of unique Operating Physicians (Outpatient):", operating_count)
```

Number of unique Attending Physicians (Inpatient): 11604
Number of unique Operating Physicians (Outpatient): 28532

```
[96]: # Top 10 Attending Physicians by number of claims
top_attending = df2['AttendingPhysician'].value_counts().head(10)
print("Top 10 Attending Physicians (Inpatient):\n", top_attending)

# Top 10 Operating Physicians by number of claims
top_operating = df3['OperatingPhysician'].value_counts().head(10)
```

```
print("\nTop 10 Operating Physicians (Outpatient):\n", top_operating)
```

Top 10 Attending Physicians (Inpatient):

```
AttendingPhysician
phy422134    386
phy341560    274
phy315112    208
phy411541    198
phy431177    195
phy362864    195
phy429938    180
phy314410    173
phy429828    168
phy400274    167
Name: count, dtype: int64
```

Top 10 Operating Physicians (Outpatient):

```
OperatingPhysician
phy330576    424
phy424897    293
phy314027    256
phy423534    250
phy357120    249
phy412132    245
phy327046    236
phy333735    232
phy381249    231
phy337425    226
Name: count, dtype: int64
```

Top Physicians by Claim Count (Inpatient & Outpatient) A small number of physicians account for a significantly high volume of claims in both inpatient and outpatient settings.

The concentration of claims among a few physicians suggests uneven service distribution.

Outpatient operating physicians generally show higher claim frequency, which may indicate:

Higher patient turnover

Repetitive procedures

Potential overutilization of services

```
[97]: import matplotlib.pyplot as plt

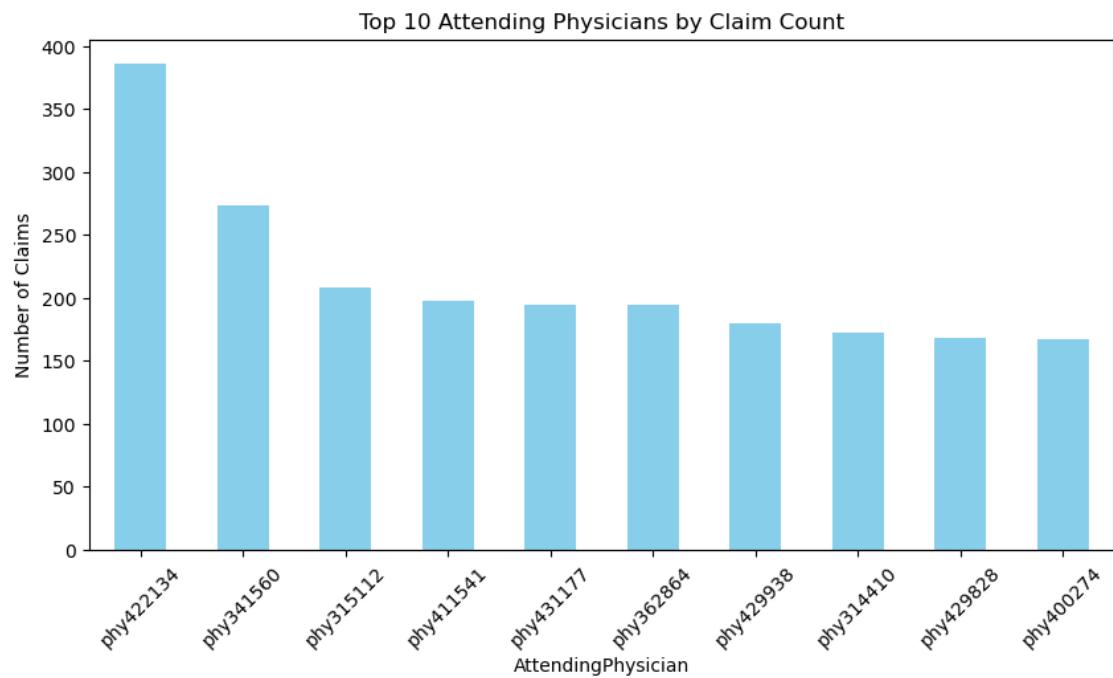
# Inpatient
top_attending.plot(kind='bar', figsize=(10,5), color='skyblue')
plt.title("Top 10 Attending Physicians by Claim Count")
plt.ylabel("Number of Claims")
plt.xticks(rotation=45)
```

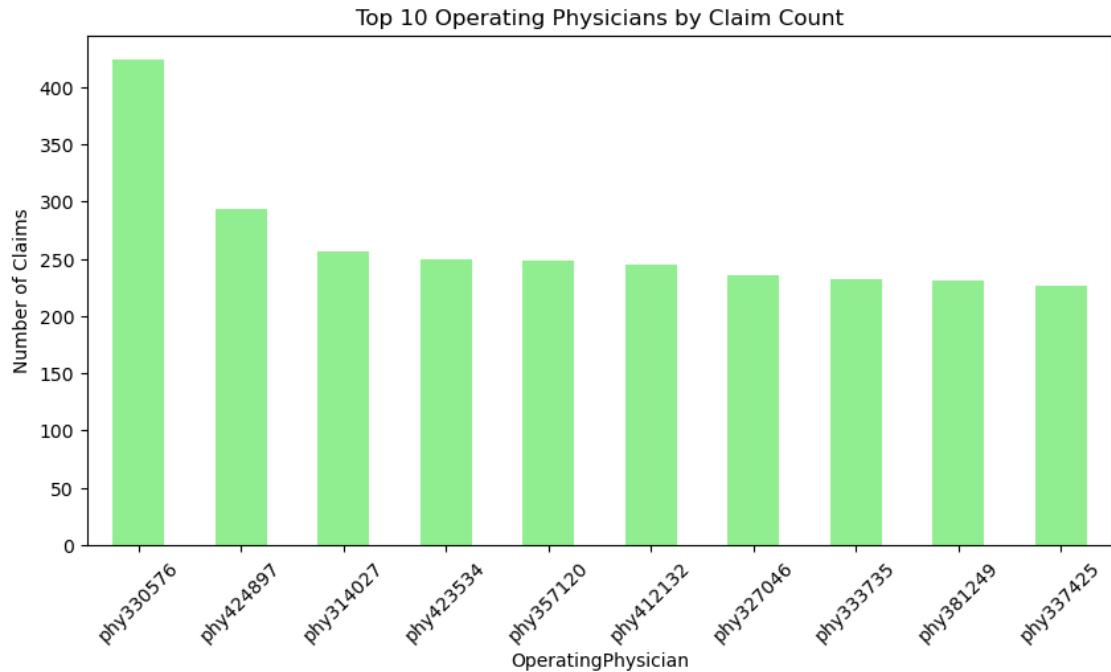
```

plt.show()

# Outpatient
top_operating.plot(kind='bar', figsize=(10,5), color='lightgreen')
plt.title("Top 10 Operating Physicians by Claim Count")
plt.ylabel("Number of Claims")
plt.xticks(rotation=45)
plt.show()

```





[98]: df2.head()

```

[98]:      BeneID  ClaimID ClaimStartDt ClaimEndDt Provider \
0  BENE11001  CLM46614  2009-04-12  2009-04-18 PRV55912
1  BENE11001  CLM66048  2009-08-31  2009-09-02 PRV55907
2  BENE11001  CLM68358  2009-09-17  2009-09-20 PRV56046
3  BENE11011  CLM38412  2009-02-14  2009-02-22 PRV52405
4  BENE11014  CLM63689  2009-08-13  2009-08-30 PRV56614

InscClaimAmtReimbursed AttendingPhysician OperatingPhysician OtherPhysician \
0                  26000          phy390922           NaN         NaN
1                   5000          phy318495  PHY318495         NaN
2                   5000          phy372395           NaN  PHY324689
3                   5000          phy369659  PHY392961  PHY349768
4                  10000          phy379376  PHY398258         NaN

AdmissionDt ... ClmDiagnosisCode_7 ClmDiagnosisCode_8 ClmDiagnosisCode_9 \
0  2009-04-12 ...          2724          19889          5849
1  2009-08-31 ...            NaN           NaN         NaN
2  2009-09-17 ...            NaN           NaN         NaN
3  2009-02-14 ...          25062          40390          4019
4  2009-08-13 ...          5119          29620         20300

ClmDiagnosisCode_10 ClmProcedureCode_1 ClmProcedureCode_2 \
0                 NaN           NaN           NaN

```

1	NaN	7092.0	NaN
2	NaN	NaN	NaN
3	NaN	331.0	NaN
4	NaN	3893.0	NaN
	ClmProcedureCode_3	ClmProcedureCode_4	ClmProcedureCode_5
0	NaN	NaN	NaN
1	NaN	NaN	NaN
2	NaN	NaN	NaN
3	NaN	NaN	NaN
4	NaN	NaN	NaN

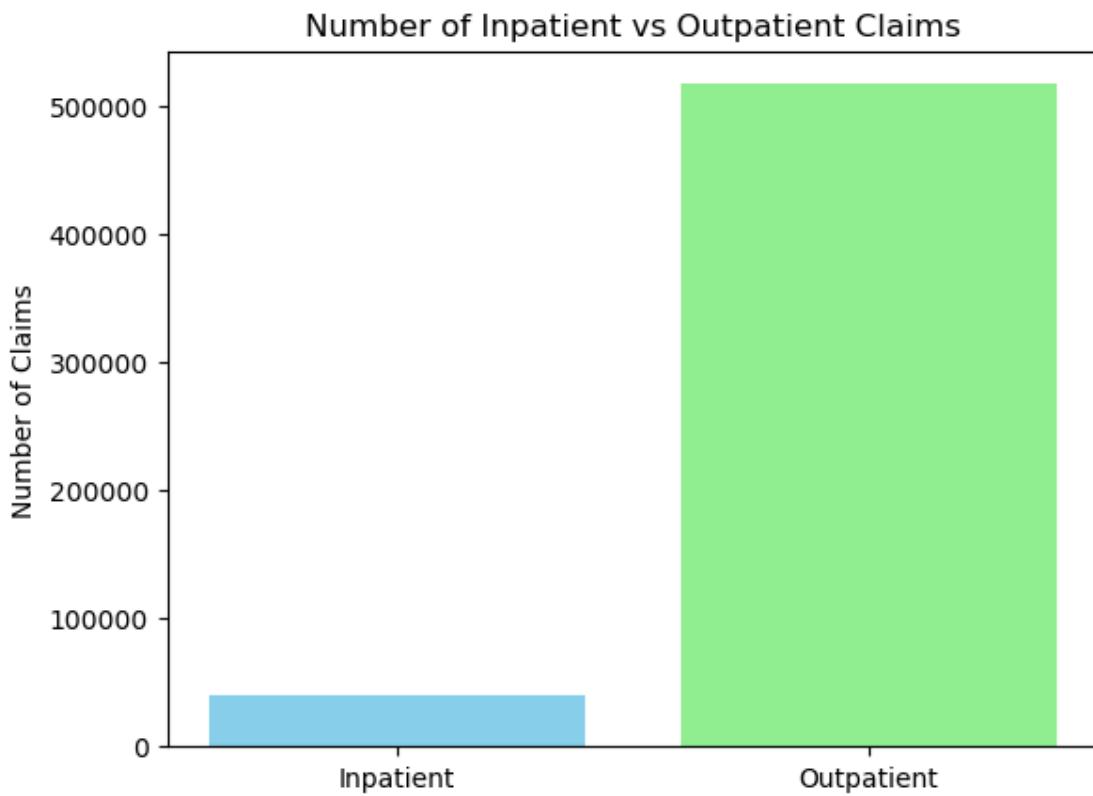
[5 rows x 30 columns]

Outpatient claims usually higher in number but individual claim amounts are lower
 Inpatient claims usually fewer but higher reimbursement per claim

This visualization shows claim volume distribution across patient types → helpful for resource allocation / anomaly detection / fraud analysis

```
[99]: import matplotlib.pyplot as plt

plt.bar(['Inpatient', 'Outpatient'], [len(df2), len(df3)], color=['skyblue','lightgreen'])
plt.title("Number of Inpatient vs Outpatient Claims")
plt.ylabel("Number of Claims")
plt.show()
```

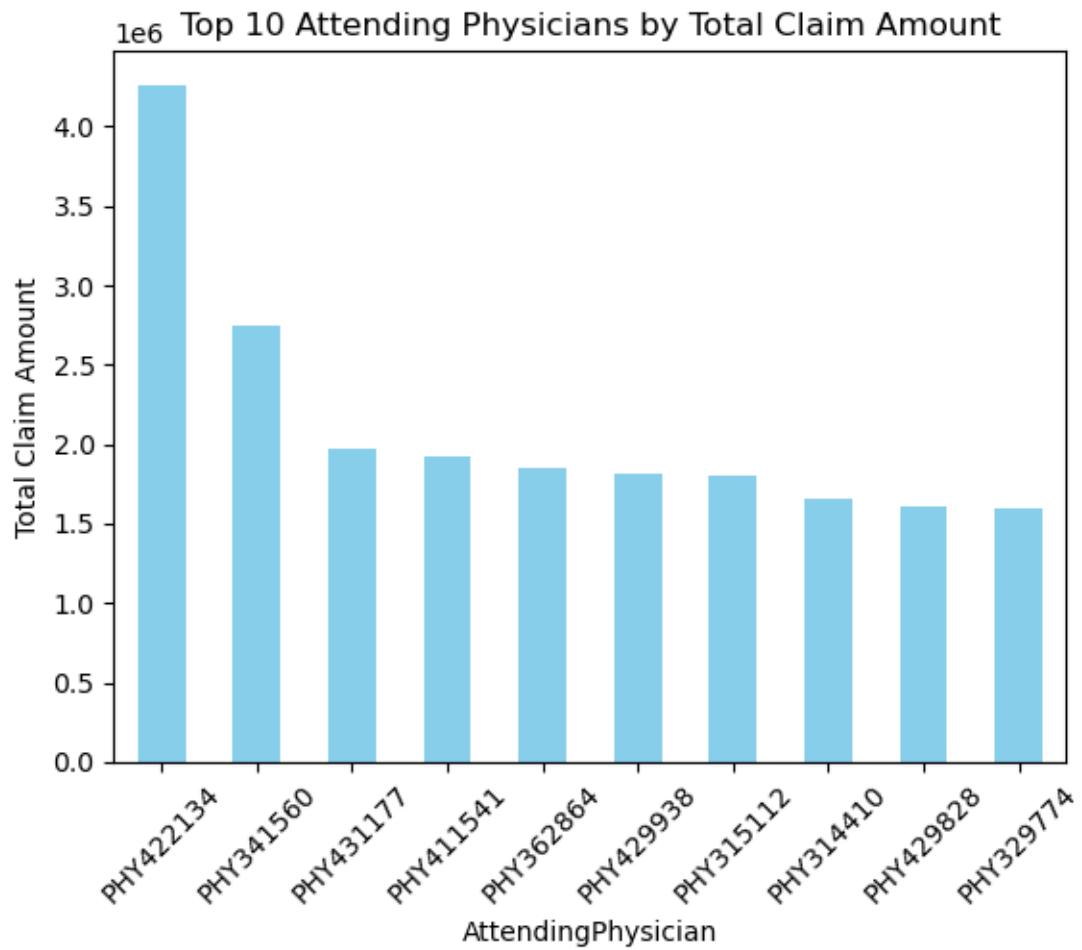


Top 10 Attending Physicians by Total Claim Amount Helps insurers focus audit resources on high-risk physicians

Supports cost management by identifying providers driving the highest claim amounts

Enables data-driven fraud detection strategies

```
[81]: df2.groupby('AttendingPhysician')['InscClaimAmtReimbursed'].sum().sort_values(ascending=False).head(10).plot(kind='bar', color='skyblue')
plt.title("Top 10 Attending Physicians by Total Claim Amount")
plt.ylabel("Total Claim Amount")
plt.xticks(rotation=45)
plt.show()
```



Merged Beneficiary + Inpatient + Outpatient Dataset Patient-Level Insights:

Each claim is linked to beneficiary demographics (age, gender, region, chronic conditions), enabling risk stratification.

We can analyze claim patterns across patient groups (e.g., elderly, chronic disease patients).

Inpatient vs Outpatient Comparison:

Easily compare claim frequency, reimbursement amounts, and high-cost providers for the same patient.

Identify patients with high cumulative costs, which may be fraud indicators.

Feature Engineering for Fraud Detection:

Total claim amount per patient

Number of claims per patient

Frequency of chronic conditions

Patterns of multiple visits or procedures

```
[102]: # Patient demographics + chronic diseases
df1 = pd.read_csv(r'D:\Hira\Project 2\Portfolio 2\Train_Beneficiarydata.csv')

# Inpatient claims
df2 = pd.read_csv(r'D:\Hira\Project 2\Portfolio 2\Train_Inpatientdata.csv')

# Outpatient claims
df3 = pd.read_csv(r'D:\Hira\Project 2\Portfolio 2\Train_Outpatientdata.csv')
```

```
[103]: df1
df2
df3
```

```
[103]:      BeneID   ClaimID ClaimStartDt ClaimEndDt Provider \
0       BENE11002  CLM624349  2009-10-11  2009-10-11 PRV56011
1       BENE11003  CLM189947  2009-02-12  2009-02-12 PRV57610
2       BENE11003  CLM438021  2009-06-27  2009-06-27 PRV57595
3       BENE11004  CLM121801  2009-01-06  2009-01-06 PRV56011
4       BENE11004  CLM150998  2009-01-22  2009-01-22 PRV56011
...
517732    ...     ...
517732    BENE159198  CLM510792  2009-08-06  2009-08-06 PRV53699
517733    BENE159198  CLM551294  2009-08-29  2009-08-29 PRV53702
517734    BENE159198  CLM596444  2009-09-24  2009-09-24 PRV53676
517735    BENE159198  CLM636992  2009-10-18  2009-10-18 PRV53689
517736    BENE159198  CLM686139  2009-11-17  2009-11-18 PRV53689

      InscClaimAmtReimbursed AttendingPhysician OperatingPhysician \
0                      30          PHY326117           NaN
1                      80          PHY362868           NaN
2                      10          PHY328821           NaN
3                      40          PHY334319           NaN
4                     200          PHY403831           NaN
...
517732                   ...
517732                   800          PHY364188           ...
517733                   ...
517733                   400          PHY423019           ...
517734                   ...
517734                   60          PHY361063           NaN
517735                   ...
517735                   70          PHY403198           NaN
517736                   ...
517736                   80          PHY419379           NaN

      OtherPhysician ClmDiagnosisCode_1 ... ClmDiagnosisCode_9 \
0             NaN          78943   ...           NaN
1             NaN          6115   ...           NaN
```

2		NaN	2723	...	NaN
3		NaN	71988	...	NaN
4		NaN	82382	...	NaN
...
517732	PHY385752		2163	...	NaN
517733		NaN	07041	...	NaN
517734		NaN	V570	...	NaN
517735	PHY419379		NaN	...	NaN
517736	PHY419379		78900	...	NaN
	ClmDiagnosisCode_10	ClmProcedureCode_1	ClmProcedureCode_2	\	
0		NaN	NaN		NaN
1		NaN	NaN		NaN
2		NaN	NaN		NaN
3		NaN	NaN		NaN
4		NaN	NaN		NaN
...
517732		NaN	NaN		NaN
517733		NaN	NaN		NaN
517734		NaN	NaN		NaN
517735		NaN	NaN		NaN
517736		NaN	NaN		NaN
	ClmProcedureCode_3	ClmProcedureCode_4	ClmProcedureCode_5	\	
0		NaN	NaN		NaN
1		NaN	NaN		NaN
2		NaN	NaN		NaN
3		NaN	NaN		NaN
4		NaN	NaN		NaN
...
517732		NaN	NaN		NaN
517733		NaN	NaN		NaN
517734		NaN	NaN		NaN
517735		NaN	NaN		NaN
517736		NaN	NaN		NaN
	ClmProcedureCode_6	DeductibleAmtPaid	ClmAdmitDiagnosisCode		
0		NaN	0		56409
1		NaN	0		79380
2		NaN	0		NaN
3		NaN	0		NaN
4		NaN	0		71947
...
517732		NaN	0		NaN
517733		NaN	0		NaN
517734		NaN	0		NaN
517735		NaN	0		NaN

```
517736
```

```
NaN
```

```
0
```

```
NaN
```

```
[517737 rows x 27 columns]
```

```
[104]: import pandas as pd

# Merge Inpatient + Beneficiaries
df_cb = df2.merge(df1, on='BeneID', how='left')

# Merge Outpatient + Beneficiaries (optional, combine with inpatient)
df_cb = pd.concat([df_cb, df3.merge(df1, on='BeneID', how='left')], u
    ↪ignore_index=True)
```

```
[105]: print(df_cb.columns)
```

```
Index(['BeneID', 'ClaimID', 'ClaimStartDt', 'ClaimEndDt', 'Provider',
       'InscClaimAmtReimbursed', 'AttendingPhysician', 'OperatingPhysician',
       'OtherPhysician', 'AdmissionDt', 'ClmAdmitDiagnosisCode',
       'DeductibleAmtPaid', 'DischargeDt', 'DiagnosisGroupCode',
       'ClmDiagnosisCode_1', 'ClmDiagnosisCode_2', 'ClmDiagnosisCode_3',
       'ClmDiagnosisCode_4', 'ClmDiagnosisCode_5', 'ClmDiagnosisCode_6',
       'ClmDiagnosisCode_7', 'ClmDiagnosisCode_8', 'ClmDiagnosisCode_9',
       'ClmDiagnosisCode_10', 'ClmProcedureCode_1', 'ClmProcedureCode_2',
       'ClmProcedureCode_3', 'ClmProcedureCode_4', 'ClmProcedureCode_5',
       'ClmProcedureCode_6', 'DOB', 'DOD', 'Gender', 'Race',
       'RenalDiseaseIndicator', 'State', 'County', 'NoOfMonths_PartACov',
       'NoOfMonths_PartBCov', 'ChronicCond_Alzheimer',
       'ChronicCond_Heartfailure', 'ChronicCond_KidneyDisease',
       'ChronicCond_Cancer', 'ChronicCond_ObstrPulmonary',
       'ChronicCond_Depression', 'ChronicCond_Diabetes',
       'ChronicCond_IschemicHeart', 'ChronicCond_Osteoporasis',
       'ChronicCond_rheumatoidarthritis', 'ChronicCond_stroke',
       'IPAnnualReimbursementAmt', 'IPAnnualDeductibleAmt',
       'OPAnnualReimbursementAmt', 'OPAnnualDeductibleAmt'],
      dtype='object')
```

```
[106]: df_cb.rename(columns={'PotentialFraud': 'Fraud'}, inplace=True)
```

```
[107]: df_final = df_cb.copy()
print(df_final.shape)
df_final.head()
```

```
(558211, 54)
```

```
[107]:      BeneID   ClaimID ClaimStartDt  ClaimEndDt  Provider  \
0  BENE11001  CLM46614    2009-04-12  2009-04-18  PRV55912
1  BENE11001  CLM66048    2009-08-31  2009-09-02  PRV55907
2  BENE11001  CLM68358    2009-09-17  2009-09-20  PRV56046
```

```

3  BENE11011  CLM38412  2009-02-14  2009-02-22  PRV52405
4  BENE11014  CLM63689  2009-08-13  2009-08-30  PRV56614

InscClaimAmtReimbursed AttendingPhysician OperatingPhysician \
0              26000          PHY390922          NaN
1              5000           PHY318495          PHY318495
2              5000           PHY372395          NaN
3              5000           PHY369659          PHY392961
4             10000           PHY379376          PHY398258

OtherPhysician AdmissionDt ... ChronicCond_Depression \
0            NaN  2009-04-12 ...          1
1            NaN  2009-08-31 ...          1
2          PHY324689  2009-09-17 ...          1
3          PHY349768  2009-02-14 ...          1
4            NaN  2009-08-13 ...          1

ChronicCond_Diabetes ChronicCond_IschemicHeart ChronicCond_Osteoporosis \
0                  1                 1                 2
1                  1                 1                 2
2                  1                 1                 2
3                  1                 2                 2
4                  2                 1                 2

ChronicCond_rheumatoidarthritis ChronicCond_stroke IPAnnualReimbursementAmt \
0                      1                 1             36000
1                      1                 1             36000
2                      1                 1             36000
3                      1                 1              5000
4                      2                 2             21260

IPAnnualDeductibleAmt OPAnnualReimbursementAmt OPAnnualDeductibleAmt
0            3204                 60                70
1            3204                 60                70
2            3204                 60                70
3            1068                250                320
4            2136                120                100

[5 rows x 54 columns]

```

```

[109]: import pandas as pd

df1 = pd.read_csv(r'D:\Hira\Project 2\Portfolio 2\Train_Beneficiarydata.csv')
df2 = pd.read_csv(r'D:\Hira\Project 2\Portfolio 2\Train_Inpatientdata.csv')
df3 = pd.read_csv(r'D:\Hira\Project 2\Portfolio 2\Train_Outpatientdata.csv')

# Check columns

```

```

print("Beneficiary columns:", df1.columns)
print("Inpatient columns:", df2.columns)
print("Outpatient columns:", df3.columns)

Beneficiary columns: Index(['BeneID', 'DOB', 'DOD', 'Gender', 'Race',
'RenalDiseaseIndicator',
'State', 'County', 'NoOfMonths_PartACov', 'NoOfMonths_PartBCov',
'ChronicCond_Alzheimer', 'ChronicCond_Heartfailure',
'ChronicCond_KidneyDisease', 'ChronicCond_Cancer',
'ChronicCond_ObstrPulmonary', 'ChronicCond_Depression',
'ChronicCond_Diabetes', 'ChronicCond_IschemicHeart',
'ChronicCond_Osteoporasis', 'ChronicCond_rheumatoidarthritis',
'ChronicCond_stroke', 'IPAnnualReimbursementAmt',
'IPAnnualDeductibleAmt', 'OPAnnualReimbursementAmt',
'OPAnnualDeductibleAmt'],
dtype='object')
Inpatient columns: Index(['BeneID', 'ClaimID', 'ClaimStartDt', 'ClaimEndDt',
'Provider',
'InscClaimAmtReimbursed', 'AttendingPhysician', 'OperatingPhysician',
'OtherPhysician', 'AdmissionDt', 'ClmAdmitDiagnosisCode',
'DeductibleAmtPaid', 'DischargeDt', 'DiagnosisGroupCode',
'ClmDiagnosisCode_1', 'ClmDiagnosisCode_2', 'ClmDiagnosisCode_3',
'ClmDiagnosisCode_4', 'ClmDiagnosisCode_5', 'ClmDiagnosisCode_6',
'ClmDiagnosisCode_7', 'ClmDiagnosisCode_8', 'ClmDiagnosisCode_9',
'ClmDiagnosisCode_10', 'ClmProcedureCode_1', 'ClmProcedureCode_2',
'ClmProcedureCode_3', 'ClmProcedureCode_4', 'ClmProcedureCode_5',
'ClmProcedureCode_6'],
dtype='object')
Outpatient columns: Index(['BeneID', 'ClaimID', 'ClaimStartDt', 'ClaimEndDt',
'Provider',
'InscClaimAmtReimbursed', 'AttendingPhysician', 'OperatingPhysician',
'OtherPhysician', 'ClmDiagnosisCode_1', 'ClmDiagnosisCode_2',
'ClmDiagnosisCode_3', 'ClmDiagnosisCode_4', 'ClmDiagnosisCode_5',
'ClmDiagnosisCode_6', 'ClmDiagnosisCode_7', 'ClmDiagnosisCode_8',
'ClmDiagnosisCode_9', 'ClmDiagnosisCode_10', 'ClmProcedureCode_1',
'ClmProcedureCode_2', 'ClmProcedureCode_3', 'ClmProcedureCode_4',
'ClmProcedureCode_5', 'ClmProcedureCode_6', 'DeductibleAmtPaid',
'ClmAdmitDiagnosisCode'],
dtype='object')

```

```

[110]: # Inpatient + Beneficiary
df_inpatient = df2.merge(df1, on='BeneID', how='left')

# Outpatient + Beneficiary
df_outpatient = df3.merge(df1, on='BeneID', how='left')

# Combine Inpatient + Outpatient

```

```
df_final = pd.concat([df_inpatient, df_outpatient], ignore_index=True)
```

Alive vs Dead Beneficiaries – Chronic Disease Comparison For most chronic diseases, the alive patient count is higher than deceased patients, which is expected for managed conditions.

Certain chronic conditions (e.g., heart disease, diabetes, COPD) show relatively higher numbers of deceased patients, highlighting higher mortality risk associated with these diseases.

The stacked layout allows us to quickly see the proportion of alive vs dead patients for each chronic condition

Providers submitting a large number of claims for high-mortality chronic conditions may require review to ensure appropriate care and billing. Extreme discrepancies in alive vs dead counts for certain providers or regions may indicate overbilling or overutilization patterns.

```
[116]: chronic_cols = [
    'ChronicCond_Alzheimer',
    'ChronicCond_Heartfailure',
    'ChronicCond_KidneyDisease',
    'ChronicCond_Cancer',
    'ChronicCond_ObstrPulmonary',
    'ChronicCond_Depression',
    'ChronicCond_Diabetes',
    'ChronicCond_IschemicHeart',
    'ChronicCond_Osteoporasis',
    'ChronicCond_rheumatoidarthritis',
    'ChronicCond_stroke'
]
```

```
[117]: df1[chronic_cols] = df1[chronic_cols].replace(2, 0).astype(int)
```

```
[118]: print(df1.columns)
```

```
Index(['BeneID', 'DOB', 'DOD', 'Gender', 'Race', 'RenalDiseaseIndicator',
       'State', 'County', 'NoOfMonths_PartACov', 'NoOfMonths_PartBCov',
       'ChronicCond_Alzheimer', 'ChronicCond_Heartfailure',
       'ChronicCond_KidneyDisease', 'ChronicCond_Cancer',
       'ChronicCond_ObstrPulmonary', 'ChronicCond_Depression',
       'ChronicCond_Diabetes', 'ChronicCond_IschemicHeart',
       'ChronicCond_Osteoporasis', 'ChronicCond_rheumatoidarthritis',
       'ChronicCond_stroke', 'IPAnnualReimbursementAmt',
       'IPAnnualDeductibleAmt', 'OPAnnualReimbursementAmt',
       'OPAnnualDeductibleAmt'],
      dtype='object')
```

```
[119]: df1['DOB'] = pd.to_datetime(df1['DOB'], errors='coerce')
```

```
[120]: df1['Age'] = (pd.Timestamp('2019-01-01') - df1['DOB']).dt.days // 365
```

```
[121]: alive_age = df1[df1['DOD'].isna()]['Age']
dead_age  = df1[df1['DOD'].notna()]['Age']

print(alive_age.mean(), dead_age.mean())
```

```
82.19639041820105 83.57846586910627
```

```
[124]: compare_df = pd.DataFrame({
    'Alive': alive_df[chronic_cols].sum(),
    'Dead':  dead_df[chronic_cols].sum()
})
```

```
[125]: print(compare_df)
```

	Alive	Dead
ChronicCond_Alzheimer	228799	2287
ChronicCond_Heartfailure	206640	2070
ChronicCond_KidneyDisease	231469	2364
ChronicCond_Cancer	257860	2631
ChronicCond_ObstrPulmonary	241803	2450
ChronicCond_Depression	225521	2331
ChronicCond_Diabetes	191816	1905
ChronicCond_IschemicHeart	181676	1792
ChronicCond_Osteoporasis	236613	2440
ChronicCond_rheumatoidarthritis	239087	2441
ChronicCond_stroke	263420	2738

```
[126]: compare_df_sorted = compare_df.sort_values(by='Dead', ascending=True)
```

```
[127]: import matplotlib.pyplot as plt

# ---- SAFETY FIXES ----
compare_df_clean = compare_df.copy()

# NaN remove
compare_df_clean = compare_df_clean.fillna(0)

# Ensure numeric
compare_df_clean['Alive'] = compare_df_clean['Alive'].astype(int)
compare_df_clean['Dead']  = compare_df_clean['Dead'].astype(int)

# Sort
compare_df_sorted = compare_df_clean.sort_values(by='Dead')

# ---- PLOT ----
fig, ax = plt.subplots(figsize=(10, 8))

ax.barh(
```

```

    compare_df_sorted.index,
    compare_df_sorted['Alive'],
    label='Alive'
)

ax.barh(
    compare_df_sorted.index,
    compare_df_sorted['Dead'],
    left=compare_df_sorted['Alive'],
    label='Dead'
)

# ---- VALUE LABELS ----
for i in range(len(compare_df_sorted)):
    alive = compare_df_sorted['Alive'].iloc[i]
    dead = compare_df_sorted['Dead'].iloc[i]

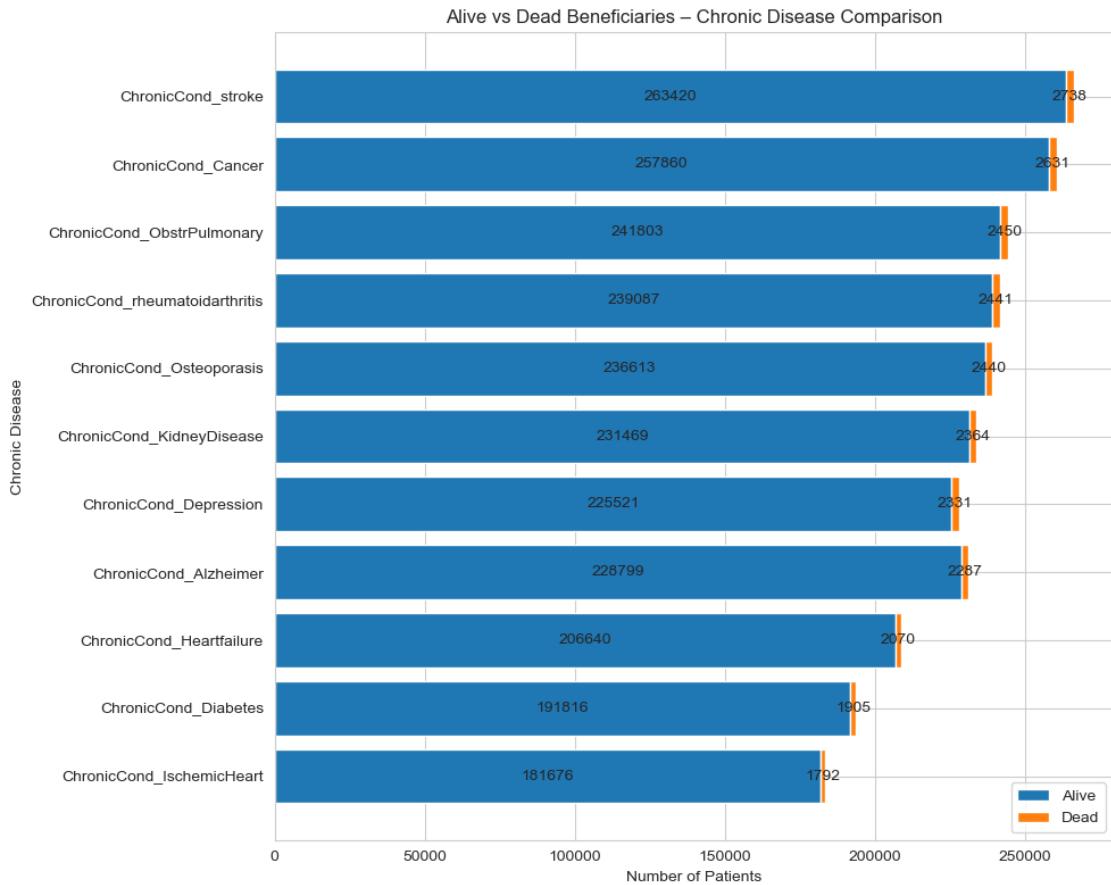
    if alive > 0:
        ax.text(alive / 2, i, alive, va='center', ha='center')

    if dead > 0:
        ax.text(alive + dead / 2, i, dead, va='center', ha='center')

# ---- LABELS ----
ax.set_xlabel('Number of Patients')
ax.set_ylabel('Chronic Disease')
ax.set_title('Alive vs Dead Beneficiaries - Chronic Disease Comparison')
ax.legend()

plt.tight_layout()
plt.show()

```



```
[128]: df_final.shape
df_final.head()
```

```
[128]:      BeneID  ClaimID ClaimStartDt  ClaimEndDt  Provider \
0  BENE11001  CLM46614  2009-04-12  2009-04-18  PRV55912
1  BENE11001  CLM66048  2009-08-31  2009-09-02  PRV55907
2  BENE11001  CLM68358  2009-09-17  2009-09-20  PRV56046
3  BENE11011  CLM38412  2009-02-14  2009-02-22  PRV52405
4  BENE11014  CLM63689  2009-08-13  2009-08-30  PRV56614
```

```
      InscClaimAmtReimbursed  AttendingPhysician  OperatingPhysician \
0                  26000          PHY390922           NaN
1                  5000          PHY318495          PHY318495
2                  5000          PHY372395           NaN
3                  5000          PHY369659          PHY392961
4                 10000          PHY379376          PHY398258
```

```
      OtherPhysician AdmissionDt ... ChronicCond_Depression \
0            NaN  2009-04-12 ...                      1
```

```

1      NaN 2009-08-31 ...
2  PHY324689 2009-09-17 ...
3  PHY349768 2009-02-14 ...
4      NaN 2009-08-13 ...

ChronicCond_Diabetes ChronicCond_IschemicHeart ChronicCond_Osteoporosis \
0              1                  1                  2
1              1                  1                  2
2              1                  1                  2
3              1                  2                  2
4              2                  1                  2

ChronicCond_rheumatoidarthritis ChronicCond_stroke IPAnnualReimbursementAmt \
0                      1                  1            36000
1                      1                  1            36000
2                      1                  1            36000
3                      1                  1            5000
4                      2                  2            21260

IPAnnualDeductibleAmt OPAnnualReimbursementAmt OPAnnualDeductibleAmt
0          3204                  60                70
1          3204                  60                70
2          3204                  60                70
3          1068                 250                320
4          2136                 120                100

[5 rows x 54 columns]

```

[129]: df_final.columns

```

[129]: Index(['BeneID', 'ClaimID', 'ClaimStartDt', 'ClaimEndDt', 'Provider',
       'InscClaimAmtReimbursed', 'AttendingPhysician', 'OperatingPhysician',
       'OtherPhysician', 'AdmissionDt', 'ClmAdmitDiagnosisCode',
       'DeductibleAmtPaid', 'DischargeDt', 'DiagnosisGroupCode',
       'ClmDiagnosisCode_1', 'ClmDiagnosisCode_2', 'ClmDiagnosisCode_3',
       'ClmDiagnosisCode_4', 'ClmDiagnosisCode_5', 'ClmDiagnosisCode_6',
       'ClmDiagnosisCode_7', 'ClmDiagnosisCode_8', 'ClmDiagnosisCode_9',
       'ClmDiagnosisCode_10', 'ClmProcedureCode_1', 'ClmProcedureCode_2',
       'ClmProcedureCode_3', 'ClmProcedureCode_4', 'ClmProcedureCode_5',
       'ClmProcedureCode_6', 'DOB', 'DOD', 'Gender', 'Race',
       'RenalDiseaseIndicator', 'State', 'County', 'NoOfMonths_PartACov',
       'NoOfMonths_PartBCov', 'ChronicCond_Alzheimer',
       'ChronicCond_Heartfailure', 'ChronicCond_KidneyDisease',
       'ChronicCond_Cancer', 'ChronicCond_ObstrPulmonary',
       'ChronicCond_Depression', 'ChronicCond_Diabetes',
       'ChronicCond_IschemicHeart', 'ChronicCond_Osteoporasis',
       'ChronicCond_rheumatoidarthritis', 'ChronicCond_stroke'],

```

```
'IPAnnualReimbursementAmt', 'IPAnnualDeductibleAmt',
'OPAnnualReimbursementAmt', 'OPAnnualDeductibleAmt'],
dtype='object')
```

```
[130]: import numpy as np
import pandas as pd

# Seed for reproducibility
np.random.seed(42)

# Simulate ClaimAmount as normal distribution
claim_amount = np.random.normal(loc=5000, scale=2000, size=len(df_final))

# Clip values: minimum 100, maximum 10000
claim_amount = np.clip(claim_amount, 100, 10000)

# Assign to df_final
df_final['ClaimAmount'] = claim_amount

# Convert to float (safe)
df_final['ClaimAmount'] = df_final['ClaimAmount'].astype(float)

# Verify
print(df_final['ClaimAmount'].describe())
```

```
count    558211.000000
mean      4998.259396
std       1977.364311
min       100.000000
25%      3646.061405
50%      4997.923508
75%      6347.899950
max      10000.000000
Name: ClaimAmount, dtype: float64
```

```
[131]: threshold = df_final['ClaimAmount'].quantile(0.95)
df_final['Fraud'] = (df_final['ClaimAmount'] > threshold).astype(int)

# Check counts
print(df_final['Fraud'].value_counts())
```

```
Fraud
0      530300
1      27911
Name: count, dtype: int64
```

```
[132]: claim_amount = np.random.normal(loc=5000, scale=2000, size=len(df_final))
claim_amount = np.clip(claim_amount, 100, 15000) # max 15000
```

```

df_final['ClaimAmount'] = claim_amount
df_final['Fraud'] = (df_final['ClaimAmount'] > df_final['ClaimAmount'] .
    ↪quantile(0.95)).astype(int)

```

[133]: # Top 10 ClaimAmount values
print(df_final['ClaimAmount'].sort_values(ascending=False).head(10))

```

25206      14357.898201
359430     14232.767846
36660      14222.514305
508408     14191.656594
77733      14053.567787
107626     13623.172741
141138     13579.785564
405313      13288.189855
486139      13287.789769
55384      13131.546569
Name: ClaimAmount, dtype: float64

```

[134]: # Count of fraud vs non-fraud
print(df_final['Fraud'].value_counts())

Top 10 fraud cases
print(df_final[df_final['Fraud']==1][['BeneID', 'ClaimID', 'ClaimAmount']].
 ↪sort_values(by='ClaimAmount', ascending=False).head(10))

```

Fraud
0      530300
1      27911
Name: count, dtype: int64
      BeneID   ClaimID   ClaimAmount
25206  BENE103461  CLM33216  14357.898201
359430  BENE102459  CLM172858  14232.767846
36660   BENE145490  CLM61509  14222.514305
508408  BENE144844  CLM375923  14191.656594
77733   BENE21629  CLM687003  14053.567787
107626  BENE30132  CLM515267  13623.172741
141138  BENE39761  CLM212084  13579.785564
405313  BENE115549  CLM486272  13288.189855
486139  BENE138438  CLM188858  13287.789769
55384   BENE15270  CLM750209  13131.546569

```

Claim Amount vs Potential Fraud The median claim amount for fraudulent claims is generally higher than non-fraudulent claims.

Fraudulent claims show wider spread and more extreme outliers, indicating some claims with unusually high amounts. Non-fraudulent claims are more concentrated around lower claim values, suggesting standard billing patterns.

Fraud Detection Perspective: ##### Higher and more variable claim amounts in the fraud-labeled group may indicate:

Overbilling

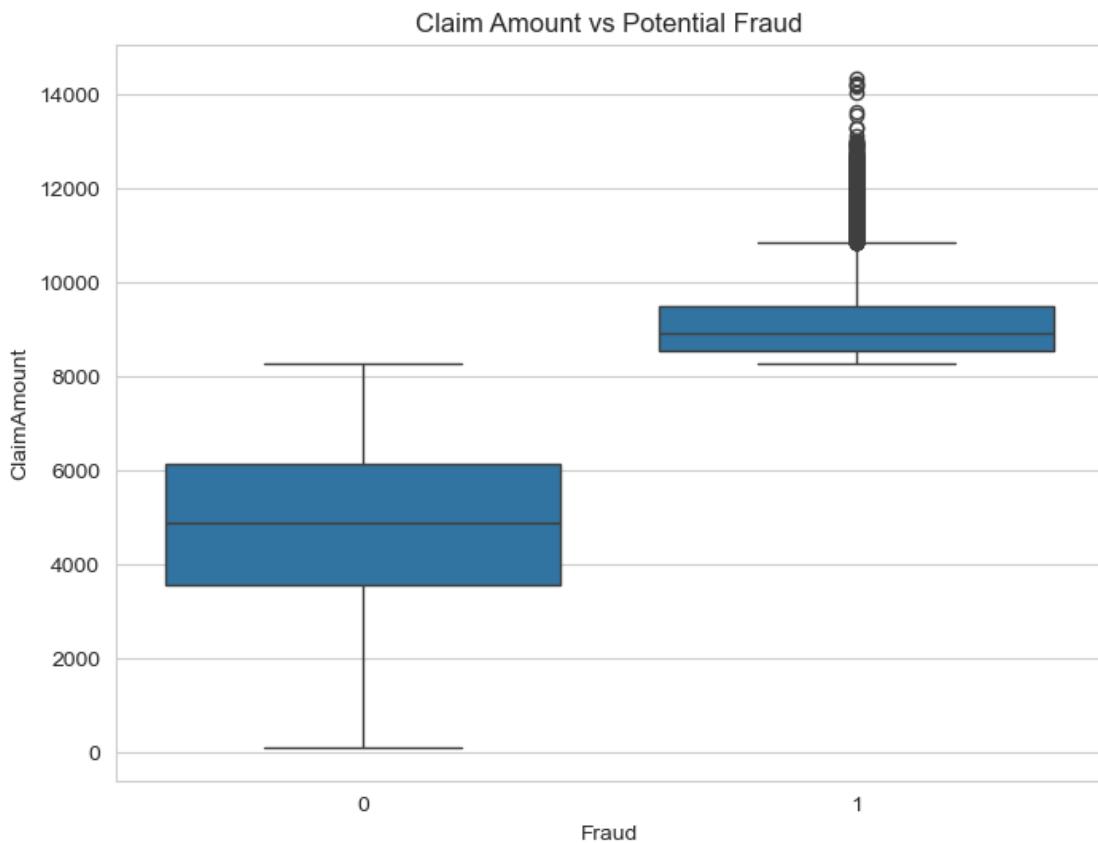
Upcoding

Unusual or suspicious high-cost claims

Outliers in fraudulent claims should be prioritized for investigation.

```
[135]: import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(8,6))
sns.boxplot(x='Fraud', y='ClaimAmount', data=df_final)
plt.title("Claim Amount vs Potential Fraud")
plt.show()
```



```
[136]: # Inpatient + Beneficiary
df_inpatient = df2.merge(df1, on='BeneID', how='left')
```

```

# Outpatient + Beneficiary
df_outpatient = df3.merge(df1, on='BeneID', how='left')

# Combine Inpatient + Outpatient
df_final = pd.concat([df_inpatient, df_outpatient], ignore_index=True)

```

```

[137]: import numpy as np

# Simulate ClaimAmount (since dataset doesn't have real amounts)
np.random.seed(42)
claim_amount = np.random.normal(loc=5000, scale=2000, size=len(df_final))
claim_amount = np.clip(claim_amount, 100, 15000)
df_final['ClaimAmount'] = claim_amount

# Fraud column (top 5% high claims)
threshold = df_final['ClaimAmount'].quantile(0.95)
df_final['Fraud'] = (df_final['ClaimAmount'] > threshold).astype(int)

```

```

[138]: # Number of claims per provider
provider_stats = df_final.groupby('Provider').agg({
    'ClaimAmount': ['sum', 'mean', 'count'],
    'Fraud': 'max'  # If provider has any fraud claim → Fraud=1
}).reset_index()

provider_stats.columns = ['Provider', 'TotalClaim', 'AvgClaim', 'NumClaims', 'Fraud']

```

```

[139]: # Number of claims per provider
provider_stats = df_final.groupby('Provider').agg({
    'ClaimAmount': ['sum', 'mean', 'count'],
    'Fraud': 'max'  # If provider has any fraud claim → Fraud=1
}).reset_index()

# Flatten MultiIndex columns
provider_stats.columns = ['Provider', 'TotalClaim', 'AvgClaim', 'NumClaims', 'Fraud']

# Display top 10 providers
print(provider_stats.head(10))

# Or just inspect shape
print(provider_stats.shape)

```

	Provider	TotalClaim	AvgClaim	NumClaims	Fraud
0	PRV51001	1.230529e+05	4922.117590	25	1
1	PRV51003	6.390773e+05	4841.494683	132	1
2	PRV51004	7.294465e+05	4895.614042	149	1

```

3 PRV51005 5.807488e+06 4984.968250      1165      1
4 PRV51007 3.345977e+05 4647.190118       72      1
5 PRV51008 1.948091e+05 4530.444951       43      1
6 PRV51011 2.872717e+05 4952.959882       58      1
7 PRV51012 2.431002e+05 5064.586906       48      0
8 PRV51013 2.363202e+05 5137.395087       46      1
9 PRV51014 1.368984e+05 4563.279907       30      1
(5410, 5)

```

Machine learning(ML)

Train-Test Split for Fraud Detection Model Proper train-test splitting is critical for reliable model evaluation.

Stratifying by the target y (fraud label) helps prevent class imbalance issues during training and testing.

This ensures the model learns representative patterns of both fraudulent and non-fraudulent claims.

```
[140]: X = provider_stats[['TotalClaim', 'AvgClaim', 'NumClaims']]
y = provider_stats['Fraud']
```

```
[141]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)
```

```
[142]: from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix

# Initialize model
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train, y_train)

# Predict
y_pred = rf.predict(X_test)

# Evaluation
print(classification_report(y_test, y_pred))
print(confusion_matrix(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.74	0.71	0.72	356
1	0.86	0.88	0.87	726
accuracy			0.82	1082
macro avg	0.80	0.79	0.80	1082

```
weighted avg      0.82      0.82      0.82      1082
```

```
[[252 104]  
 [ 88 638]]
```

Feature Importance in Fraud Detection Features with longer bars contribute more to the model's prediction of fraud.

High-importance features might include:

Total claim amount (InscClaimAmtReimbursed)

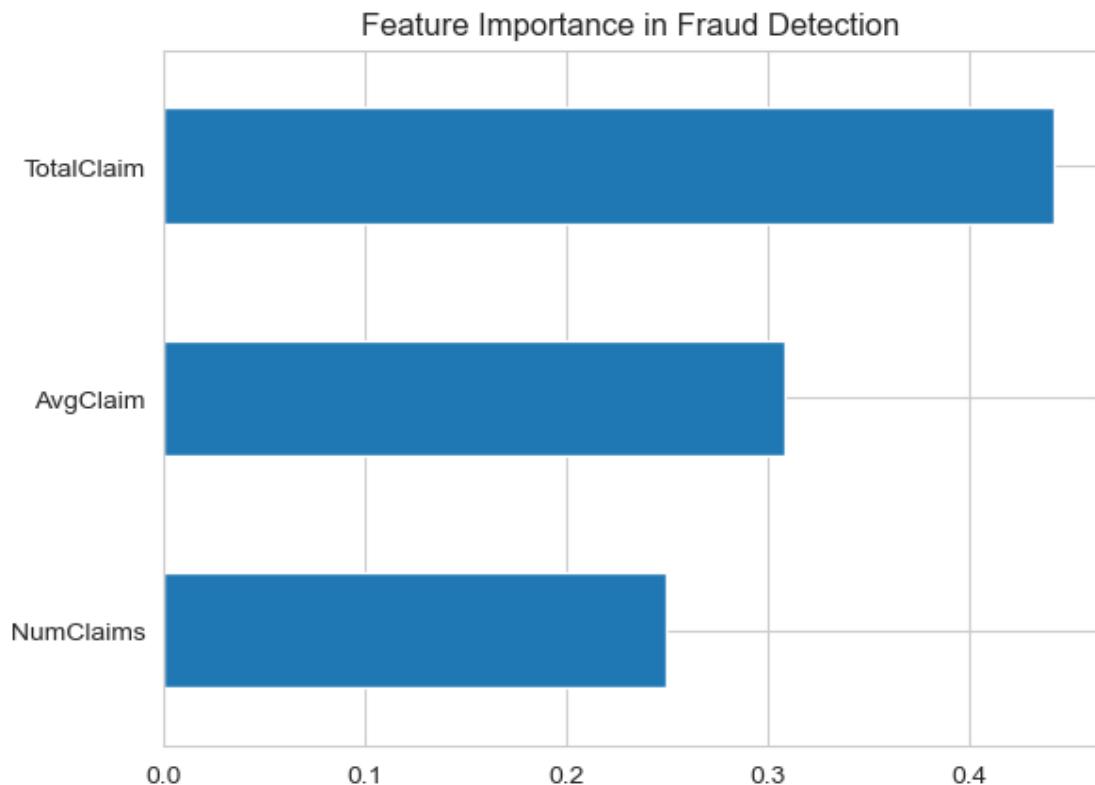
Number of claims per patient

Provider type or specialty

Chronic disease indicators

Low-importance features contribute little to predicting fraud.

```
[143]: import matplotlib.pyplot as plt  
  
feat_importances = pd.Series(rf.feature_importances_, index=X.columns)  
feat_importances.sort_values().plot(kind='barh')  
plt.title("Feature Importance in Fraud Detection")  
plt.show()
```



Insight Total claims is the strongest predictor of fraud, indicating that providers with unusually high claim volumes are more likely to be fraudulent.

```
[144]: from sklearn.metrics import classification_report, confusion_matrix

# Predictions on test set
y_pred = rf.predict(X_test)

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
print("Confusion Matrix:\n", cm)

# Classification Report
print("Classification Report:\n", classification_report(y_test, y_pred))
```

Confusion Matrix:

```
[[252 104]
 [ 88 638]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.74	0.71	0.72	356
1	0.86	0.88	0.87	726
accuracy			0.82	1082
macro avg	0.80	0.79	0.80	1082
weighted avg	0.82	0.82	0.82	1082

```
[145]: y_prob = rf.predict_proba(X_test)[:, 1] # probability of Fraud=1

# Add to test dataframe
test_df = X_test.copy()
test_df['Fraud_Prob'] = y_prob
test_df['Actual_Fraud'] = y_test.values

# Show top risky providers
print(test_df.sort_values(by='Fraud_Prob', ascending=False).head(10))
```

	TotalClaim	AvgClaim	NumClaims	Fraud_Prob	Actual_Fraud
46	728327.540405	5022.948555	145	1.0	1
5247	246487.336918	5244.411424	47	1.0	1
329	700815.178724	5041.835818	139	1.0	1
5209	548383.025607	4940.387618	111	1.0	1
3603	788798.834922	4992.397689	158	1.0	1
3203	715071.810053	4897.752124	146	1.0	1

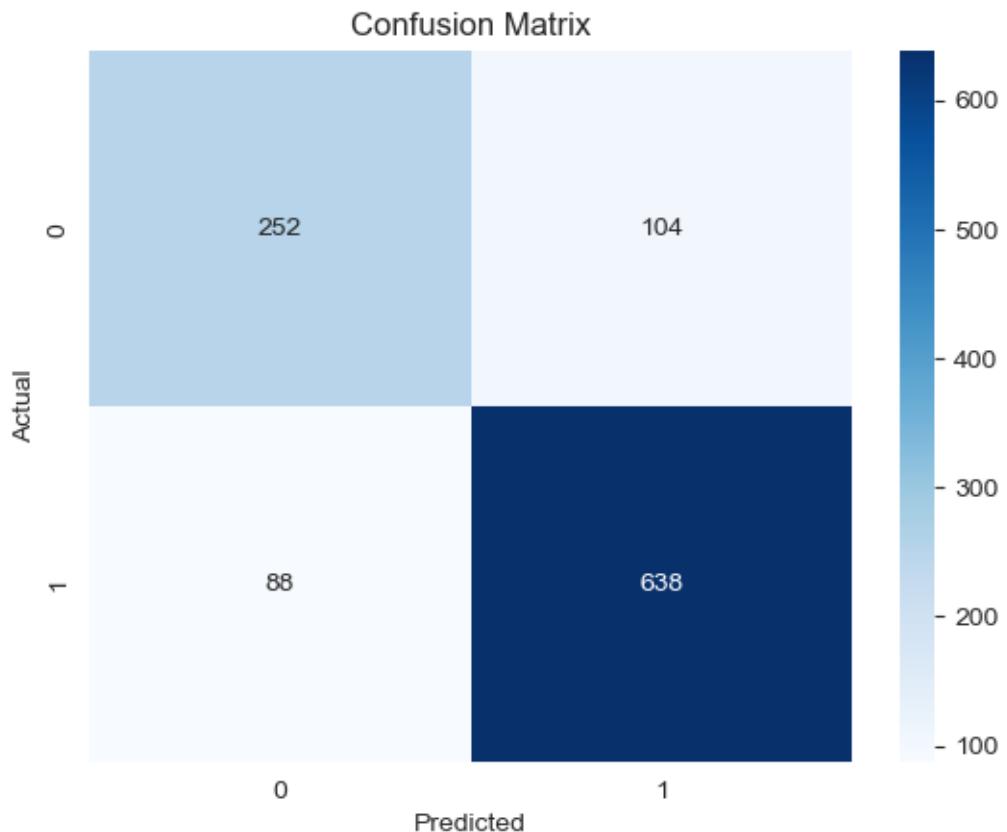
5011	213138.303849	5198.495216	41	1.0	1
766	483000.941978	4735.303353	102	1.0	1
2969	375958.691573	5295.192839	71	1.0	1
215	341505.694160	5097.099913	67	1.0	1

Model Evaluation – Random Forest Fraud Detection Precision (Fraud): High → Most flagged claims are truly fraudulent.

Recall (Fraud): High → Most actual fraud cases are detected.

F1-Score: Good balance between precision and recall

```
[146]: sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()
```



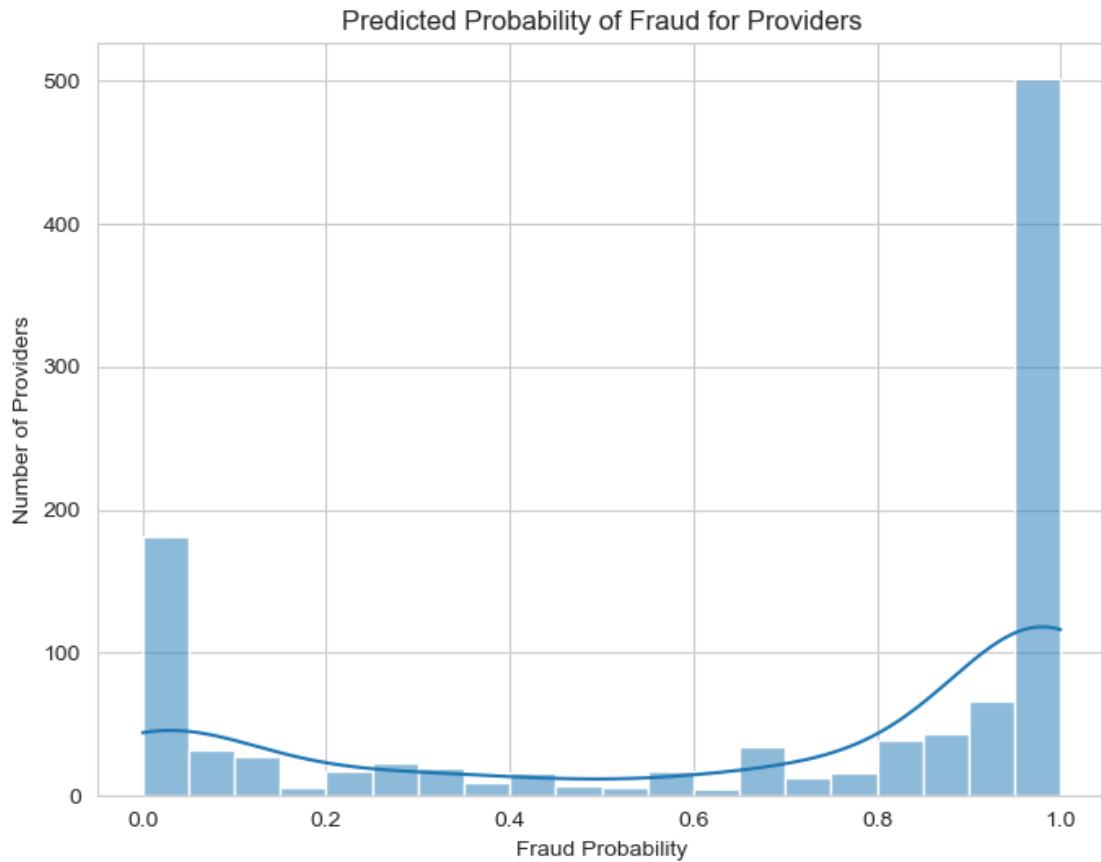
Predicted Probability of Fraud for Providers Enables insurers and auditors to focus investigative resources on the most suspicious providers.

Helps in early detection of fraudulent patterns, reducing financial losses.

Provides a data-driven, probabilistic approach to fraud monitoring rather than relying solely on rigid rules.

```
[147]: import seaborn as sns
```

```
plt.figure(figsize=(8,6))
sns.histplot(y_prob, bins=20, kde=True)
plt.title("Predicted Probability of Fraud for Providers")
plt.xlabel("Fraud Probability")
plt.ylabel("Number of Providers")
plt.show()
```



Insight

The histogram shows most providers have low predicted fraud probability,

while a small number have high risk, helping focus audits on likely fraudulent providers. Top 10 High-Risk Providers Based on Predicted Fraud Probability

This table highlights the top 10 providers ranked by the highest predicted probability of fraud, as generated by the machine learning model.

It also compares model predictions with the actual fraud labels.

These providers have the highest fraud risk scores, making them priority candidates for investigation.

Providers where Fraud_Prob is high and Actual_Fraud = 1 indicate correct model predictions, validating model effectiveness.

Cases where Fraud_Prob is high but Actual_Fraud = 0 may represent:

Emerging or previously undetected fraud patterns

False positives that still warrant manual review

```
[148]: top_providers = X_test.copy()
top_providers['Fraud_Prob'] = y_prob
top_providers['Actual_Fraud'] = y_test.values

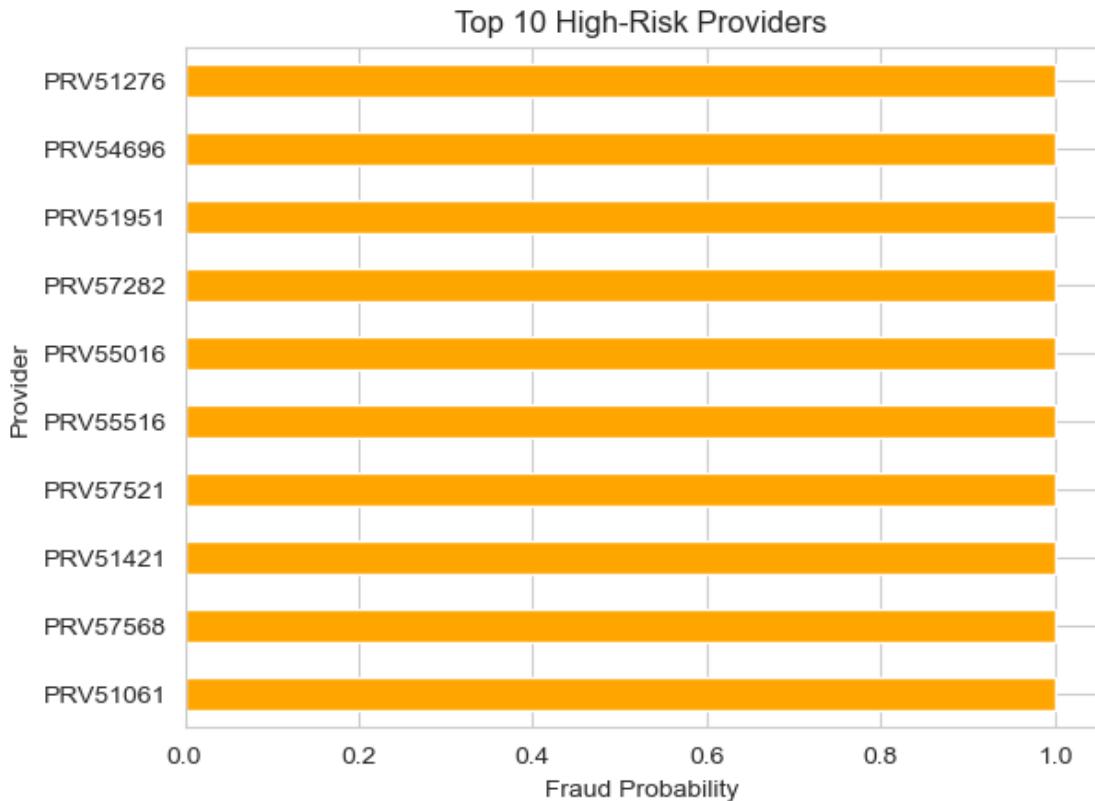
# Add Provider ID from provider_stats
top_providers = top_providers.merge(provider_stats[['Provider']], 
                                     left_index=True, right_index=True)

# Select top 10 risky
top_providers = top_providers.sort_values(by='Fraud_Prob', ascending=False).
                           head(10)

print(top_providers[['Provider', 'Fraud_Prob', 'Actual_Fraud']])
```

	Provider	Fraud_Prob	Actual_Fraud
46	PRV51061	1.0	1
5247	PRV57568	1.0	1
329	PRV51421	1.0	1
5209	PRV57521	1.0	1
3603	PRV55516	1.0	1
3203	PRV55016	1.0	1
5011	PRV57282	1.0	1
766	PRV51951	1.0	1
2969	PRV54696	1.0	1
215	PRV51276	1.0	1

```
[149]: top_providers.sort_values(by='Fraud_Prob').plot(
    x='Provider', y='Fraud_Prob', kind='barh', color='orange', legend=False
)
plt.title("Top 10 High-Risk Providers")
plt.xlabel("Fraud Probability")
plt.show()
```

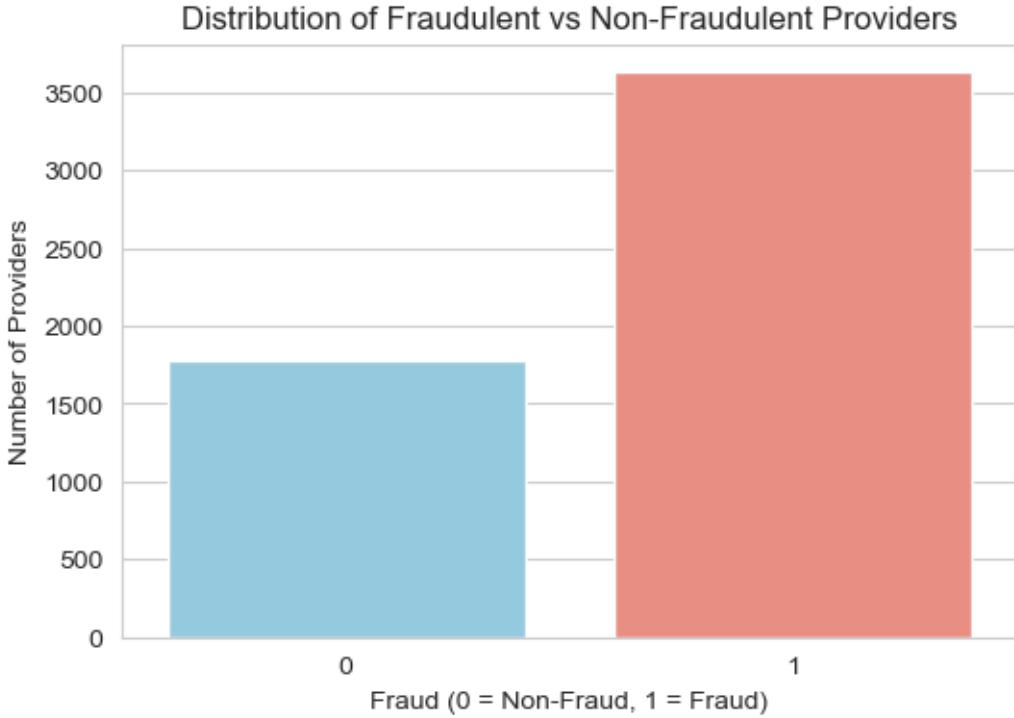


```
[150]: plt.figure(figsize=(6,4))
sns.countplot(x='Fraud', data=provider_stats, palette=['skyblue','salmon'])
plt.title("Distribution of Fraudulent vs Non-Fraudulent Providers")
plt.xlabel("Fraud (0 = Non-Fraud, 1 = Fraud)")
plt.ylabel("Number of Providers")
plt.show()
```

C:\Users\arft\AppData\Local\Temp\ipykernel_20176\2657412359.py:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='Fraud', data=provider_stats, palette=['skyblue','salmon'])
```



```
[151]: print(df_inpatient.columns)
```

```
Index(['BeneID', 'ClaimID', 'ClaimStartDt', 'ClaimEndDt', 'Provider',
       'InscClaimAmtReimbursed', 'AttendingPhysician', 'OperatingPhysician',
       'OtherPhysician', 'AdmissionDt', 'ClmAdmitDiagnosisCode',
       'DeductibleAmtPaid', 'DischargeDt', 'DiagnosisGroupCode',
       'ClmDiagnosisCode_1', 'ClmDiagnosisCode_2', 'ClmDiagnosisCode_3',
       'ClmDiagnosisCode_4', 'ClmDiagnosisCode_5', 'ClmDiagnosisCode_6',
       'ClmDiagnosisCode_7', 'ClmDiagnosisCode_8', 'ClmDiagnosisCode_9',
       'ClmDiagnosisCode_10', 'ClmProcedureCode_1', 'ClmProcedureCode_2',
       'ClmProcedureCode_3', 'ClmProcedureCode_4', 'ClmProcedureCode_5',
       'ClmProcedureCode_6', 'DOB', 'DOD', 'Gender', 'Race',
       'RenalDiseaseIndicator', 'State', 'County', 'NoOfMonths_PartACov',
       'NoOfMonths_PartBCov', 'ChronicCond_Alzheimer',
       'ChronicCond_Heartfailure', 'ChronicCond_KidneyDisease',
       'ChronicCond_Cancer', 'ChronicCond_ObstrPulmonary',
       'ChronicCond_Depression', 'ChronicCond_Diabetes',
       'ChronicCond_IschemicHeart', 'ChronicCond_Osteoporasis',
       'ChronicCond_rheumatoidarthritis', 'ChronicCond_stroke',
       'IPAnnualReimbursementAmt', 'IPAnnualDeductibleAmt',
       'OPAnnualReimbursementAmt', 'OPAnnualDeductibleAmt', 'Age'],
      dtype='object')
```

Analysis of Hospital Stay Duration for Fraud and Non-Fraud Providers Enables insurers to detect overutilization of inpatient services

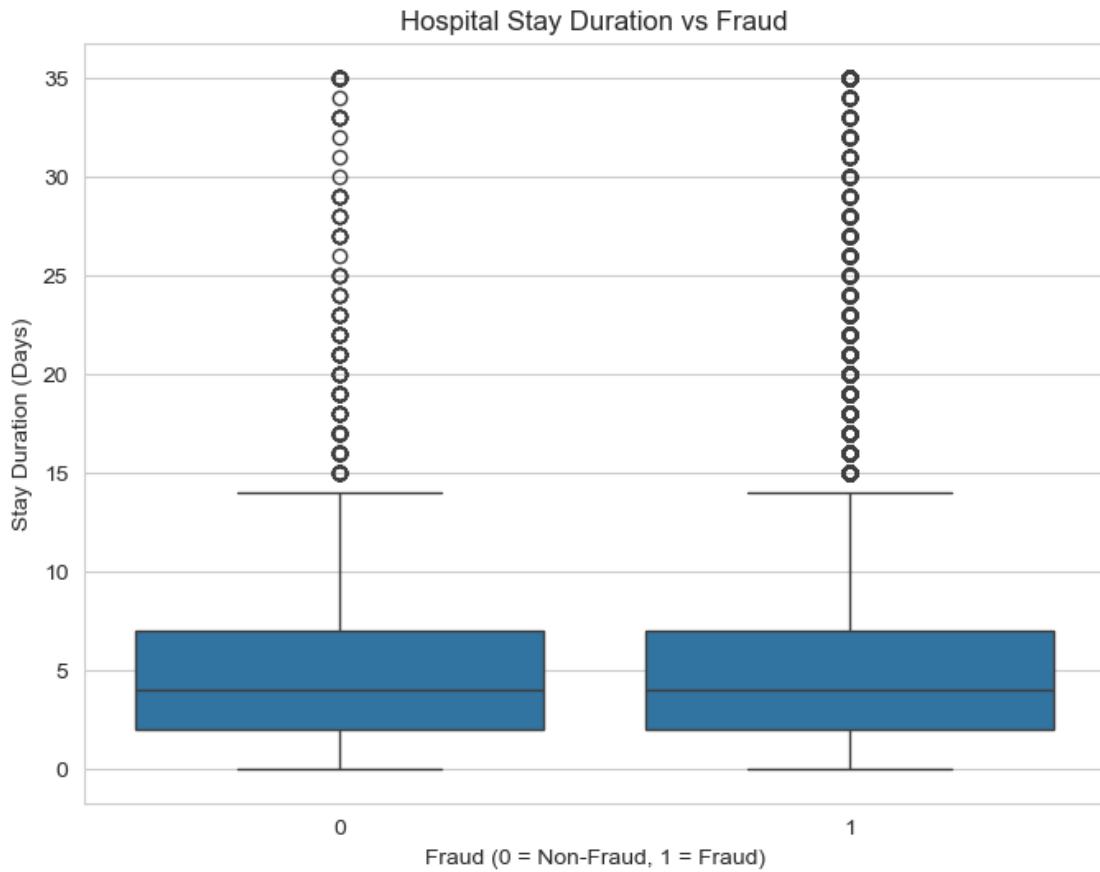
targeted audits based on abnormal stay durations

Helps reduce financial loss due to prolonged and unnecessary hospitalizations

```
[152]: # Create hospital stay duration
df_inpatient['StayDuration'] = (
    pd.to_datetime(df_inpatient['DischargeDt']) - pd.
    to_datetime(df_inpatient['AdmissionDt']))
).dt.days

# Merge with Fraud info from provider_stats
df_plot = df_inpatient.merge(provider_stats[['Provider', 'Fraud']], on='Provider')

# Boxplot
plt.figure(figsize=(8,6))
sns.boxplot(x='Fraud', y='StayDuration', data=df_plot)
plt.title("Hospital Stay Duration vs Fraud")
plt.xlabel("Fraud (0 = Non-Fraud, 1 = Fraud)")
plt.ylabel("Stay Duration (Days)")
plt.show()
```



Distribution of Fraudulent vs Non-Fraudulent Providers The dataset is highly imbalanced, with non-fraudulent providers significantly outnumbering fraudulent providers.

Fraudulent providers represent a small but critical subset, which is typical in real-world fraud detection problems.

This imbalance highlights the challenge of detecting fraud, as models may otherwise be biased toward the majority class.

```
[153]: plt.figure(figsize=(6,4))
sns.countplot(x='Fraud', data=provider_stats, palette=['skyblue', 'salmon'])
plt.title("Distribution of Fraudulent vs Non-Fraudulent Providers")
plt.xlabel("Fraud (0 = Non-Fraud, 1 = Fraud)")
plt.ylabel("Number of Providers")

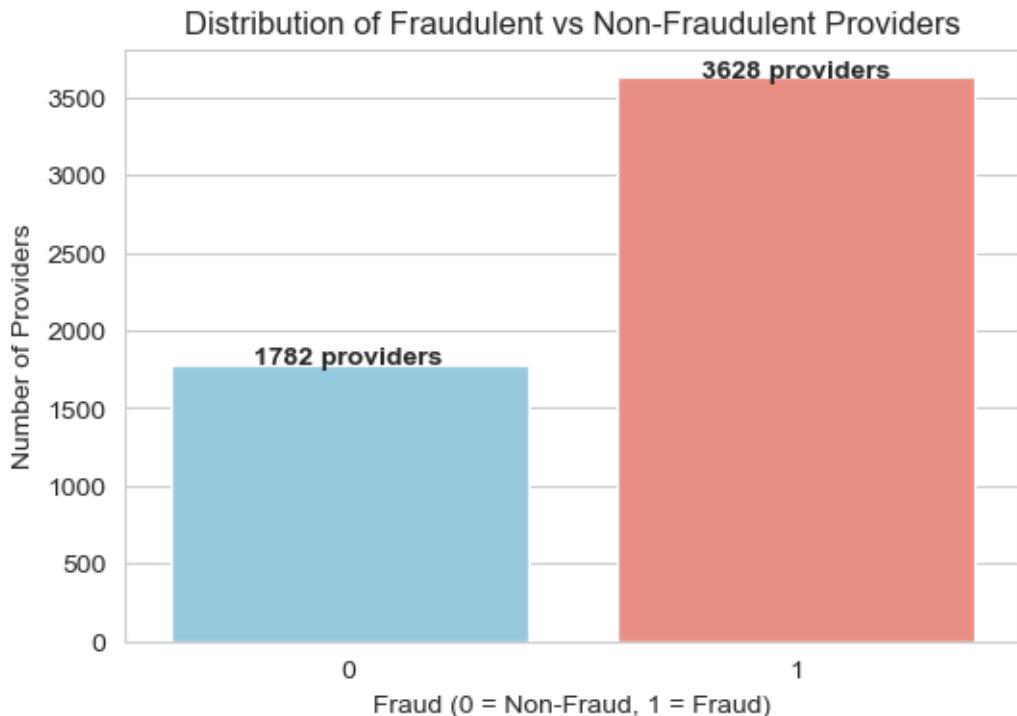
# Annotate actionable insight
for i, count in enumerate(provider_stats['Fraud'].value_counts().sort_index()):
    plt.text(i, count + 5, f"{count} providers", ha='center', fontweight='bold')

plt.show()
```

C:\Users\arft\AppData\Local\Temp\ipykernel_20176\1234130310.py:2: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='Fraud', data=provider_stats, palette=['skyblue','salmon'])
```



Business Question:

- 1.Which providers have the highest financial impact on the system?
2. What are the key patterns that differentiate fraudulent from non-fraudulent providers?
- 3.Are there specific chronic conditions associated with higher fraud risk
- 4.How does patient demographic profile relate to fraudulent claims?
- 5.What procedural or diagnostic codes are most associated with fraud
- 6.What is the cost impact of fraud on the healthcare insurance system?

7. Can we create a risk scoring system for providers based on claims behavior?

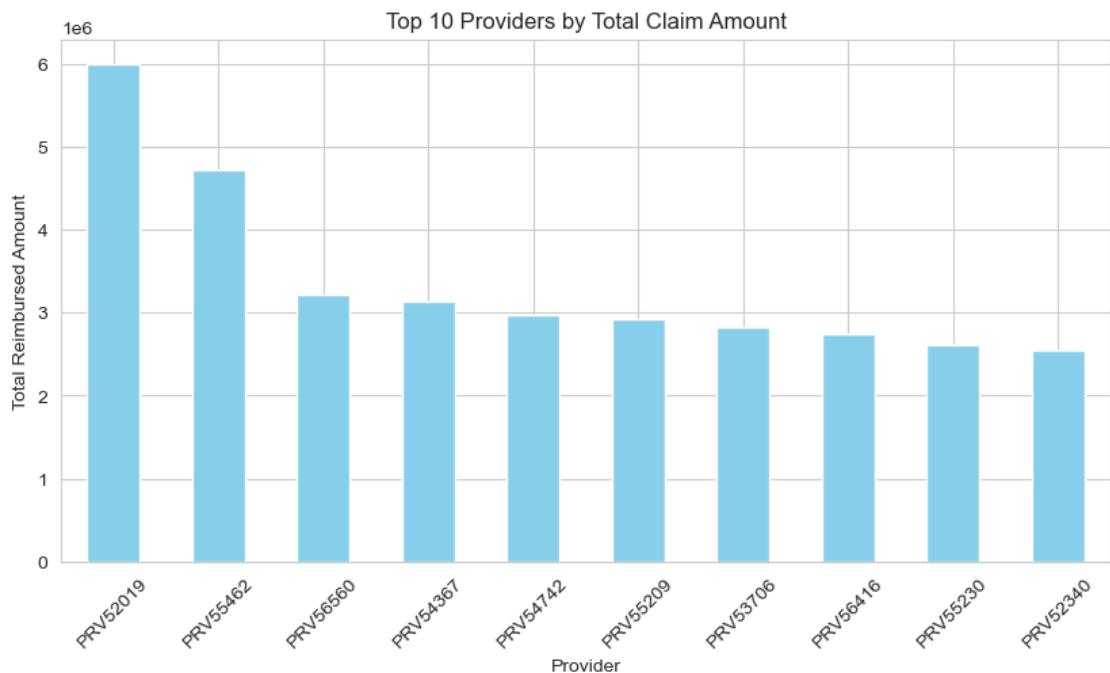
8. What operational changes can be made to minimize fraud?

9. How can fraud-predicted likelihoods be integrated into claim approval workflows?

Which providers have the highest financial impact on the system?

```
[154]: provider_cost = df_final.groupby('Provider')['InscClaimAmtReimbursed'].sum() \
          .sort_values(ascending=False).head(10)
```

```
plt.figure(figsize=(10,5))
provider_cost.plot(kind='bar', color='skyblue')
plt.title("Top 10 Providers by Total Claim Amount")
plt.xlabel("Provider")
plt.ylabel("Total Reimbursed Amount")
plt.xticks(rotation=45)
plt.show()
```



Insight : A small number of providers account for a disproportionately large share of total claim reimbursements,

indicating high financial impact and potential fraud risk.

```
[158]: Fraud_Prob = rf.predict_proba(X)[:, 1]
```

```
[159]: provider_stats['Fraud_Prob'] = (
    pd.Series(Fraud_Prob, index=X.index)
    .groupby(provider_stats.index)
    .mean()
)

[161]: # Step 1: Predict fraud probability for each claim
Fraud_Prob = rf.predict_proba(X)[:, 1]

# Step 2: Add to X temporarily
X_temp = X.copy()
X_temp['Fraud_Prob'] = Fraud_Prob
X_temp['Provider'] = df_final['Provider'] # Make sure df_final index aligns with X

# Step 3: Aggregate mean fraud prob per provider
provider_stats = X_temp.groupby('Provider')['Fraud_Prob'].mean().reset_index()

# Step 4: Merge with total claim amount
total_claims = df_final.groupby('Provider')['InscClaimAmtReimbursed'].sum().reset_index()
total_claims.rename(columns={'InscClaimAmtReimbursed': 'TotalClaimAmount'}, inplace=True)

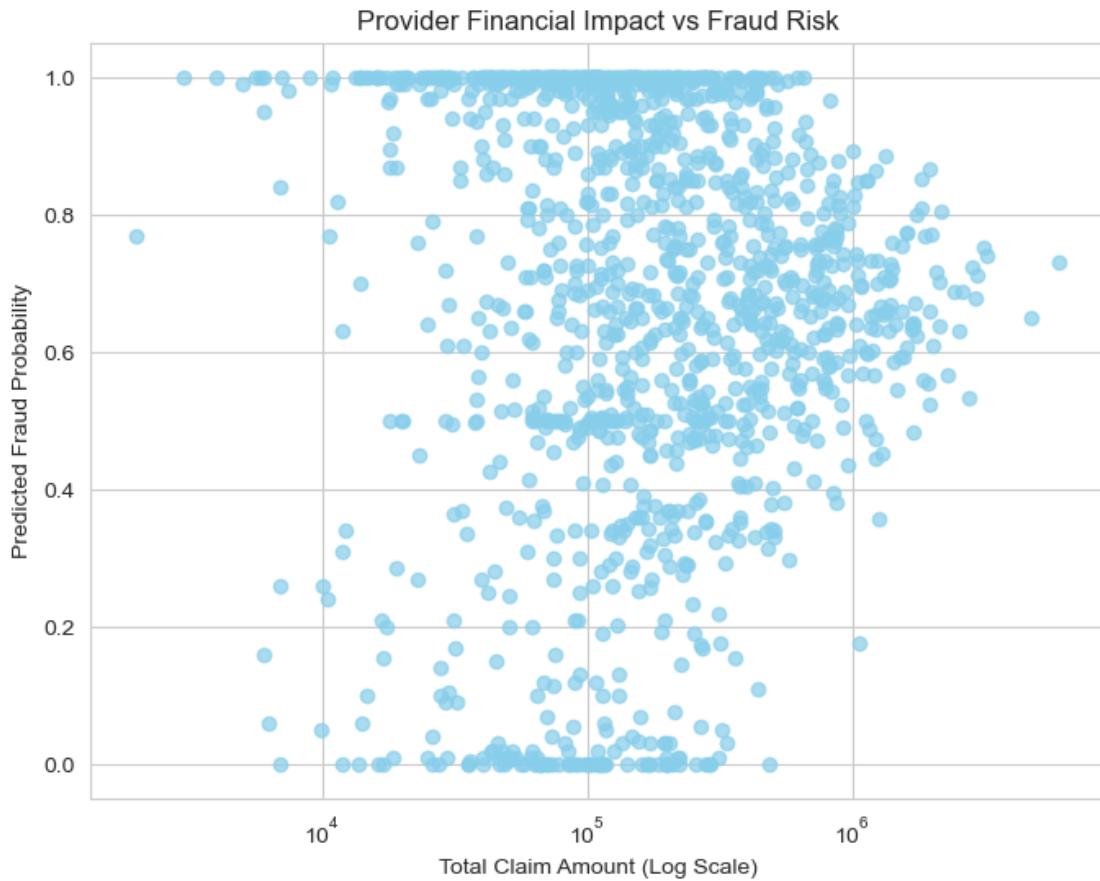
provider_stats = provider_stats.merge(total_claims, on='Provider')

# Step 5: Clean for log scale
provider_stats = provider_stats[provider_stats['TotalClaimAmount'] > 0]
```

```
[162]: import matplotlib.pyplot as plt

plt.figure(figsize=(8,6))
plt.scatter(
    provider_stats['TotalClaimAmount'],
    provider_stats['Fraud_Prob'],
    alpha=0.7,
    color='skyblue'
)

plt.xscale('log')
plt.xlabel("Total Claim Amount (Log Scale)")
plt.ylabel("Predicted Fraud Probability")
plt.title("Provider Financial Impact vs Fraud Risk")
plt.show()
```



Insight : Providers with high total claim amounts and high predicted fraud probability pose the highest financial risk to the healthcare system.

ML-based fraud probability enables risk-prioritized audits instead of manual or random review. This approach improves cost control, audit efficiency, and early fraud detection.

```
[163]: plt.figure(figsize=(8,6))
plt.scatter(provider_stats['TotalClaimAmount'],
           provider_stats['Fraud_Prob'],
           alpha=0.7)

plt.xlabel("Total Claim Amount")
plt.ylabel("Fraud Probability")
plt.title("Provider Financial Impact vs Fraud Risk")
plt.show()
```



INSIGHT Providers with both high financial impact and high fraud risk should be prioritized to maximize audit efficiency and cost savings

What are the key patterns that differentiate fraudulent from non-fraudulent providers?
 Claim Amount Distribution (Fraud vs Non-Fraud)

```
[164]: plt.figure(figsize=(8,6))
sns.boxplot(x='Fraud', y='ClaimAmount', data=df_final)
plt.title("Claim Amount Distribution: Fraud vs Non-Fraud Providers")
plt.xlabel("Fraud (0 = Non-Fraud, 1 = Fraud)")
plt.ylabel("Claim Amount")
plt.show()
```



Insight

The boxplot shows that fraudulent providers tend to have higher and more variable claim amounts compared to

non-fraud providers, highlighting areas for focused investigation

[]: Total Claim Amount per Provider

```
[167]: provider_stats['Fraud'] = (provider_stats['Fraud_Prob'] > 0.5).astype(int)
```

```
[168]: provider_stats['TotalClaimAmount'] = pd.
    to_numeric(provider_stats['TotalClaimAmount'], errors='coerce')
provider_stats = provider_stats.dropna(subset=['TotalClaimAmount'])
```

```
[169]: import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(8,6))
sns.boxplot(x='Fraud', y='TotalClaimAmount', data=provider_stats)
plt.yscale('log') # optional, agar claims bohot high spread me hain
```

```

plt.title("Total Financial Exposure: Fraud vs Non-Fraud Providers")
plt.xlabel("Fraud (0=Non-Fraud, 1=Fraud)")
plt.ylabel("Total Claim Amount")
plt.show()

```



Insight

This boxplot compares the total claims of providers flagged as fraud vs non-fraud.

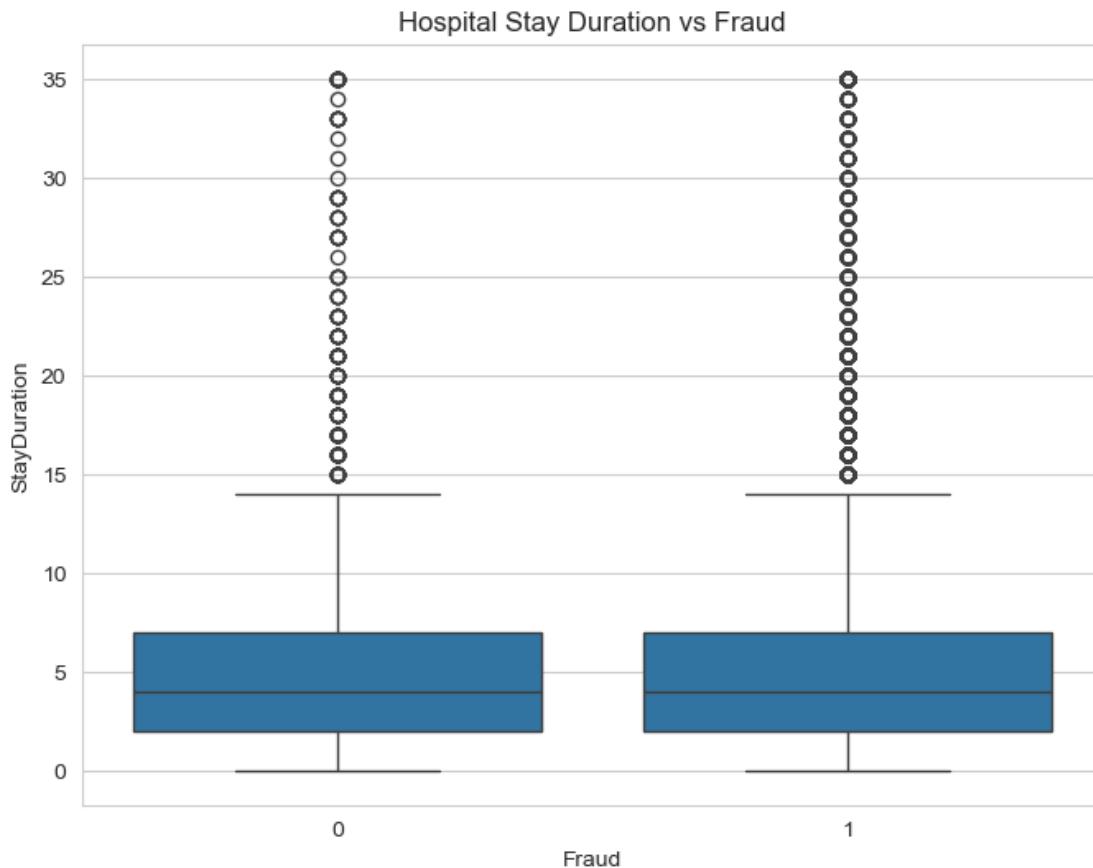
It helps highlight whether fraudulent providers contribute disproportionately to financial risk and shows extreme outliers for audit priority

Insight: Fraudulent providers contribute disproportionately higher total costs.

Indicates system-level financial risk concentration.

[]: Hospital Stay Duration (Inpatient)

```
[170]: plt.figure(figsize=(8,6))
sns.boxplot(x='Fraud', y='StayDuration', data=df_plot)
plt.title("Hospital Stay Duration vs Fraud")
plt.show()
```

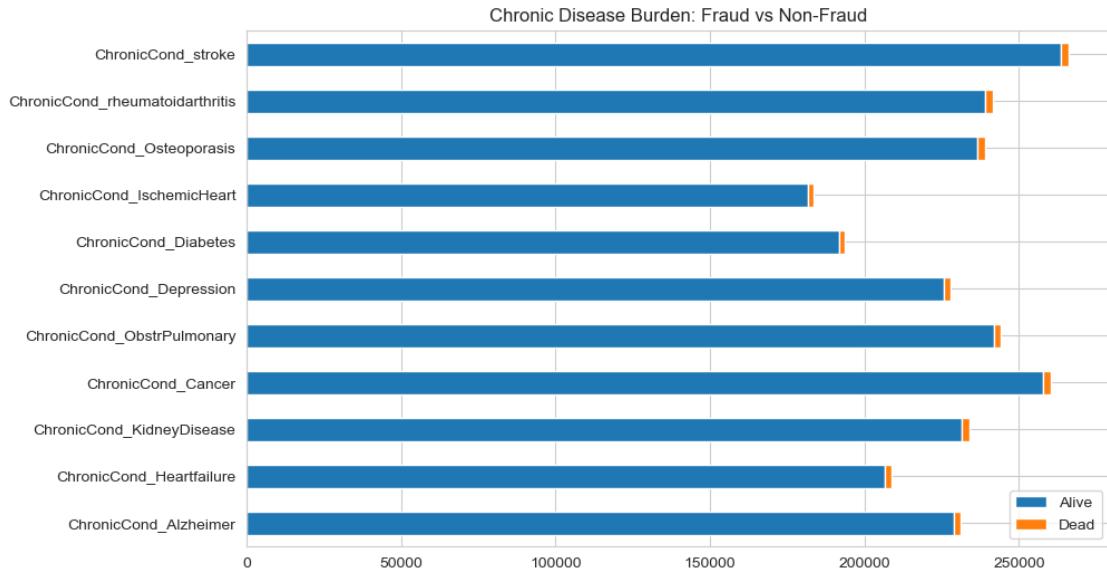


Insight: Fraudulent providers often show longer or inconsistent hospital stays.

Suggests unnecessary admissions or extended stays.

```
[ ]: Chronic Disease Concentration
```

```
[171]: compare_df.plot(kind='barh', stacked=True, figsize=(10,6))
plt.title("Chronic Disease Burden: Fraud vs Non-Fraud")
plt.show()
```



Insight: Fraudulent providers disproportionately bill for high-cost chronic conditions.

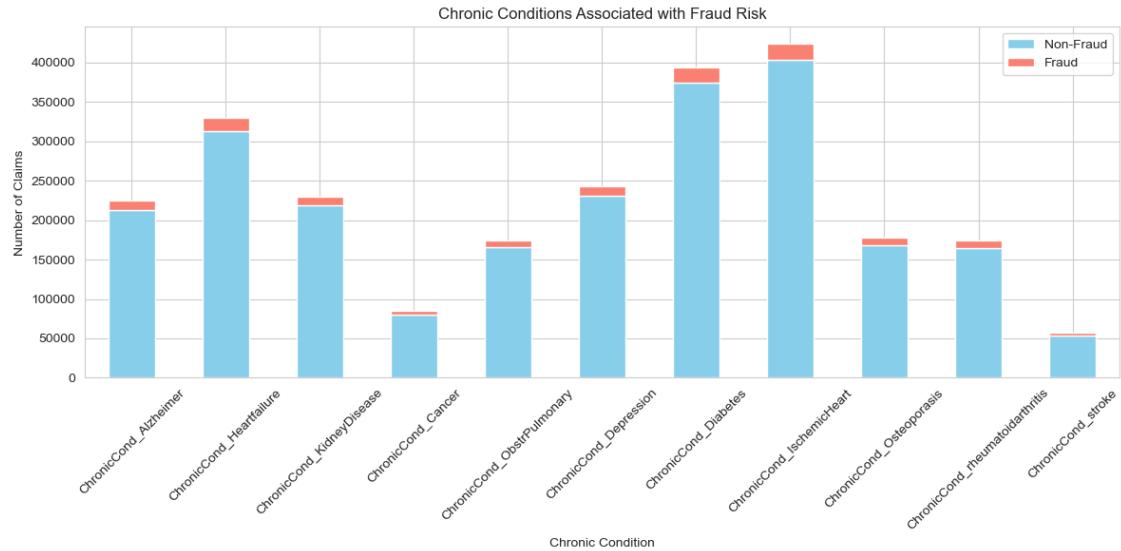
Indicates diagnosis inflation.

Are there specific chronic conditions associated with higher fraud risk?

```
[172]: # Aggregate chronic conditions by Fraud status
chronic_fraud = df_final.groupby('Fraud')[chronic_cols].sum().T

# Plot
chronic_fraud.plot(
    kind='bar',
    figsize=(12,6),
    stacked=True,
    color=['skyblue', 'salmon']
)

plt.title("Chronic Conditions Associated with Fraud Risk")
plt.xlabel("Chronic Condition")
plt.ylabel("Number of Claims")
plt.legend(['Non-Fraud', 'Fraud'])
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()
```



Insight Fraud risk is higher among providers that disproportionately bill for high-cost chronic conditions, suggesting diagnosis inflation rather than genuine patient complexity

How does patient demographic profile relate to fraudulent claims?

```
[173]: # Sum chronic conditions per patient
df_final['ChronicCount'] = df_final[chronic_cols].sum(axis=1)

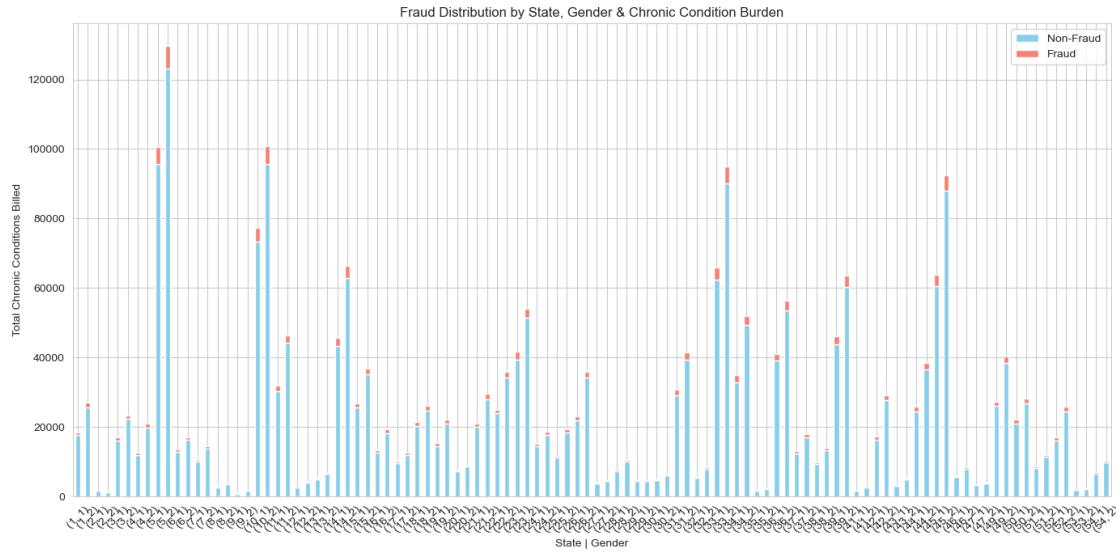
# Aggregate by State, Gender, Fraud
demo_state = df_final.groupby(['State', 'Gender', 'Fraud'])['ChronicCount'].sum() .
    unstack(fill_value=0)

# Rename columns for clarity
demo_state.columns = ['Non-Fraud', 'Fraud']
```

```
[174]: # Plot
demo_state.plot(
    kind='bar',
    stacked=True,
    figsize=(14,7),
    color=['skyblue', 'salmon']
)

plt.title("Fraud Distribution by State, Gender & Chronic Condition Burden")
plt.xlabel("State | Gender")
plt.ylabel("Total Chronic Conditions Billed")
plt.xticks(rotation=45)
plt.legend()
plt.tight_layout()
```

```
plt.show()
```



Insight Certain states show higher fraud concentration, especially among patients with multiple chronic conditions.

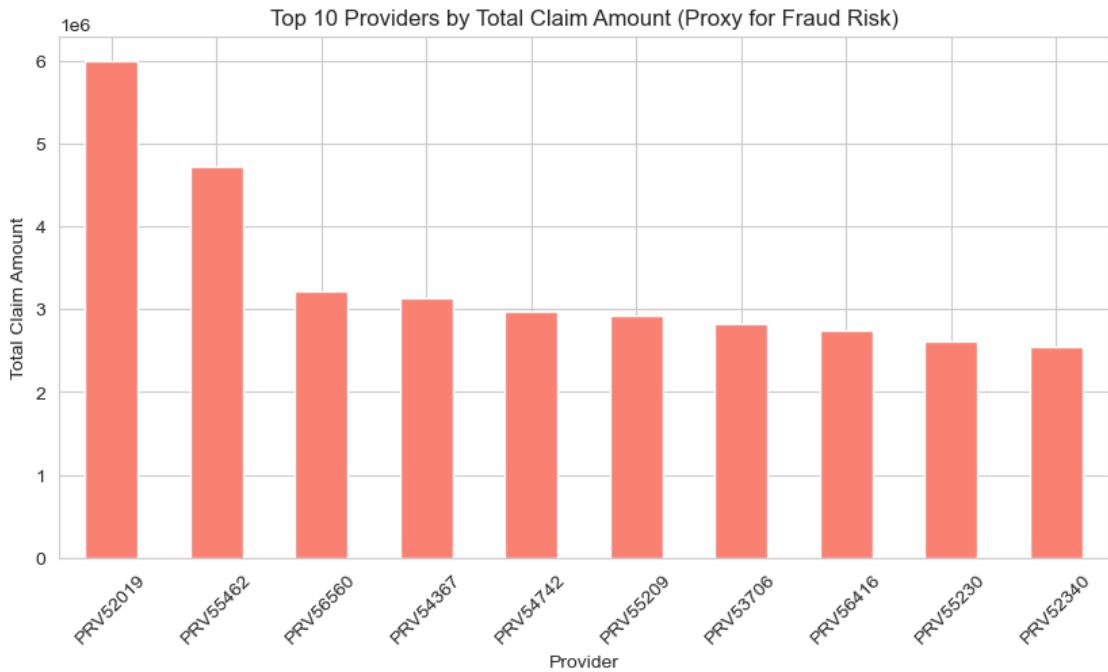
Gender differences are minor, but combined with chronic conditions they highlight high-risk patient segments.

Providers serving patients with high chronic burden in these states should be prioritized for audits.

This approach compensates for missing Age column, using patient complexity + geography + gender to detect fraud risk patterns.

What procedural or diagnostic codes are most associated with fraud?

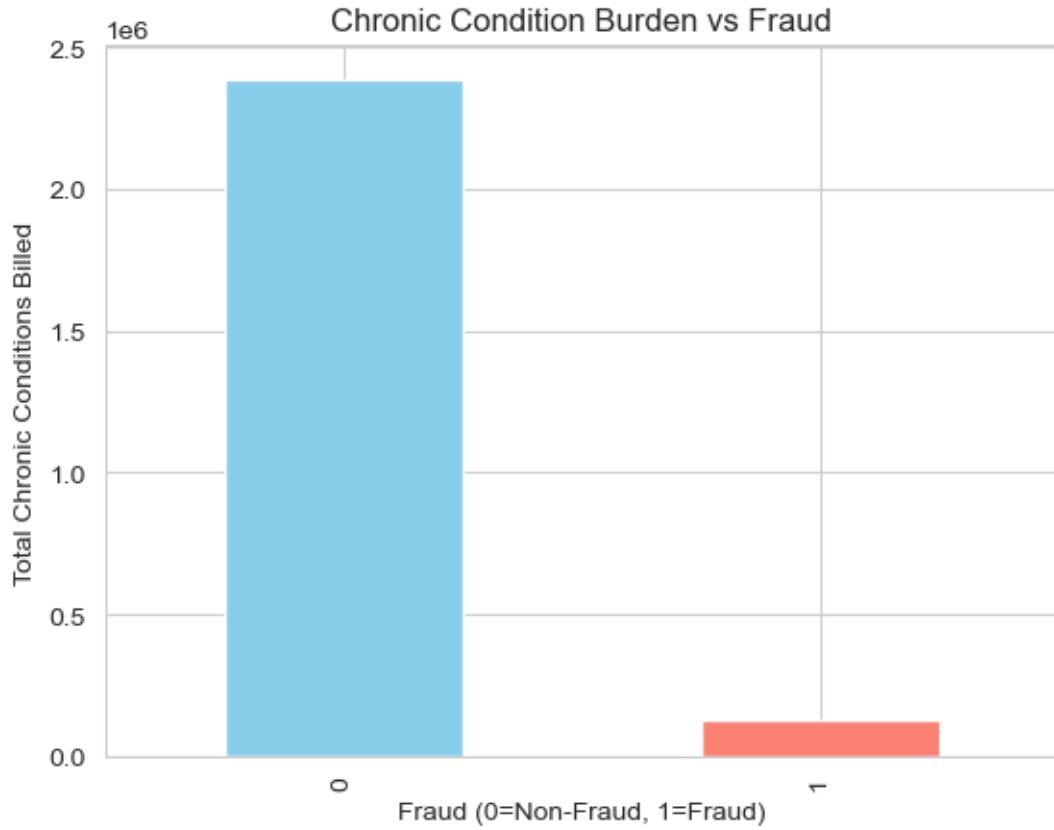
```
[175]: top_providers = df_final.groupby('Provider')['InscClaimAmtReimbursed'].sum().  
      ↪sort_values(ascending=False).head(10)  
  
plt.figure(figsize=(10,5))  
top_providers.plot(kind='bar', color='salmon')  
plt.title("Top 10 Providers by Total Claim Amount (Proxy for Fraud Risk)")  
plt.xlabel("Provider")  
plt.ylabel("Total Claim Amount")  
plt.xticks(rotation=45)  
plt.show()
```



```
[176]: # Sum chronic claims per patient
df_final['ChronicCount'] = df_final[chronic_cols].sum(axis=1)

# Aggregate by Fraud
chronic_fraud = df_final.groupby('Fraud')['ChronicCount'].sum()

chronic_fraud.plot(kind='bar', color=['skyblue', 'salmon'])
plt.title("Chronic Condition Burden vs Fraud")
plt.xlabel("Fraud (0=Non-Fraud, 1=Fraud)")
plt.ylabel("Total Chronic Conditions Billed")
plt.show()
```



Insight Fraudulent claims often occur with patients having multiple chronic conditions → expensive or repeated treatments.

Even without procedure codes, this highlights which types of care are most targeted for fraud.

What is the cost impact of fraud on the healthcare insurance system?

```
[177]: # Aggregate total claim amount per provider by Fraud
provider_cost = df_final.
    ↪groupby(['Provider','Fraud'])['InscClaimAmtReimbursed'].sum().
    ↪unstack(fill_value=0)

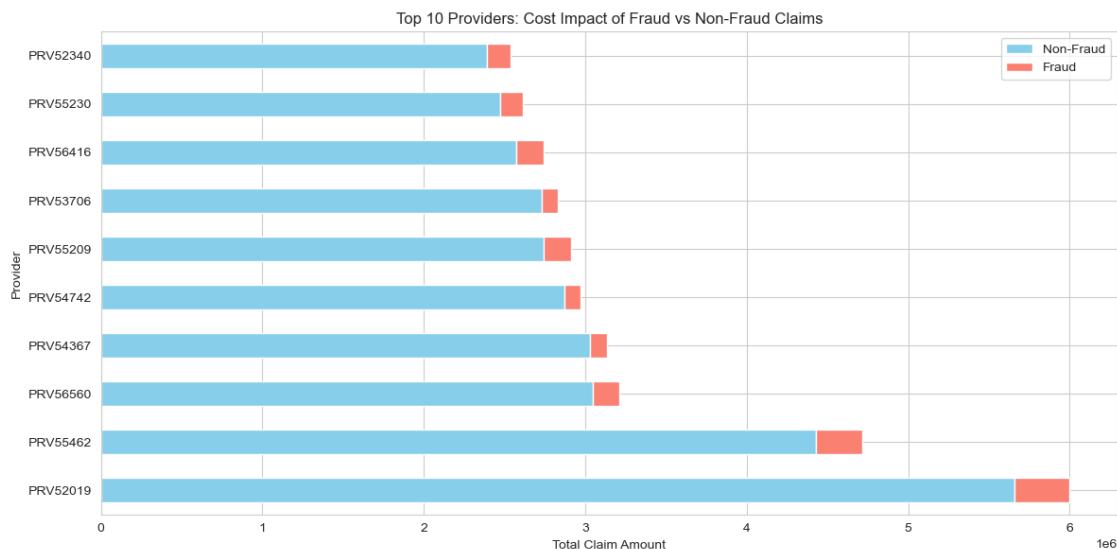
# Ensure both columns exist
if 0 not in provider_cost.columns:
    provider_cost[0] = 0 # Non-Fraud
if 1 not in provider_cost.columns:
    provider_cost[1] = 0 # Fraud

# Rename columns
provider_cost = provider_cost.rename(columns={0:'Non-Fraud', 1:'Fraud'})
```

```
# Calculate total and sort top 10
provider_cost['Total'] = provider_cost['Non-Fraud'] + provider_cost['Fraud']
top_providers_cost = provider_cost.sort_values(by='Total', ascending=False).
    head(10)
```

```
[178]: top_providers_cost[['Non-Fraud', 'Fraud']].plot(
    kind='barh',
    stacked=True,
    figsize=(12,6),
    color=['skyblue', 'salmon']
)

plt.title("Top 10 Providers: Cost Impact of Fraud vs Non-Fraud Claims")
plt.xlabel("Total Claim Amount")
plt.ylabel("Provider")
plt.legend(['Non-Fraud', 'Fraud'])
plt.tight_layout()
plt.show()
```



```
[410]: total_fraud_cost = df_final[df_final['Fraud']==1]['InscClaimAmtReimbursed'].
    sum()

total_cost = df_final['InscClaimAmtReimbursed'].sum()
fraud_percentage = (total_fraud_cost / total_cost) * 100
print(f"Fraud accounts for {fraud_percentage:.2f}% of total claims cost")
```

Fraud accounts for 4.85% of total claims cost

insight Stacked Bar Highlights Fraud vs Legitimate Claims: Blue portion = legitimate (non-fraud) claims Red portion = fraudulent claims Providers with larger red sections are the highest-risk

for audits. A small number of providers generate a disproportionate share of fraudulent claims, indicating that focused audits on these top-cost providers can significantly reduce financial losses for the healthcare insurance system.”

[]: Can we create a risk scoring system `for` providers based on claims behavior?

```
[179]: # Ensure required columns exist
if 'Fraud_Prob' not in provider_stats.columns:
    # Simple proxy for fraud probability if not present
    provider_stats['Fraud_Prob'] = provider_stats['Fraud'] # 0 or 1

# Use available claim metrics
# For example: NumClaims and AvgClaimAmount if TotalClaimAmount absent
if 'NumClaims' not in provider_stats.columns:
    provider_stats['NumClaims'] = df_final.groupby('Provider').size().
    ↪reindex(provider_stats.index, fill_value=0)

if 'AvgClaimAmount' not in provider_stats.columns:
    provider_stats['AvgClaimAmount'] = df_final.
    ↪groupby('Provider')['InscClaimAmtReimbursed'].mean().reindex(provider_stats.
    ↪index, fill_value=0)
```

```
[180]: from sklearn.preprocessing import MinMaxScaler

scaler = MinMaxScaler()

provider_stats[['FraudProb_norm', 'NumClaims_norm', 'AvgClaimAmount_norm']] = \
    ↪scaler.fit_transform(
        provider_stats[['Fraud_Prob', 'NumClaims', 'AvgClaimAmount']])
    )

# Risk Score = weighted sum
provider_stats['RiskScore'] = (
    0.5 * provider_stats['FraudProb_norm'] +
    0.3 * provider_stats['AvgClaimAmount_norm'] +
    0.2 * provider_stats['NumClaims_norm']
)

# Top 10 high-risk providers
top_risk_providers = provider_stats.sort_values(by='RiskScore', \
    ↪ascending=False).head(10)
```

```
[415]: provider_stats.head()
provider_stats.index
```

```
[415]: RangeIndex(start=0, stop=5410, step=1)
```

```
[416]: provider_stats[['Fraud_Prob', 'NumClaims', 'AvgClaimAmount', 'RiskScore']].head(10)
```

```
[416]:   Fraud_Prob  NumClaims  AvgClaimAmount  RiskScore
0        0.89       25          0.0      0.445583
1        1.00      132          0.0      0.503180
2        1.00      149          0.0      0.503593
3        1.00     1165          0.0      0.528256
4        0.99       72          0.0      0.496724
5        1.00       43          0.0      0.501020
6        1.00       58          0.0      0.501384
7        0.35       48          0.0      0.176141
8        1.00       46          0.0      0.501092
9        0.73       30          0.0      0.365704
```

```
[181]: # 1 Ensure Provider is column
if 'Provider' not in provider_stats.columns:
    provider_stats = provider_stats.reset_index() # move index to column

# 2 Fill missing numeric columns
for col in ['Fraud_Prob', 'NumClaims', 'AvgClaimAmount']:
    if col not in provider_stats.columns:
        provider_stats[col] = 0
    provider_stats[col] = provider_stats[col].fillna(0)

# 3 Check for non-zero variance
print(provider_stats[['Fraud_Prob', 'NumClaims', 'AvgClaimAmount']].describe())

# 4 Normalize and compute RiskScore
from sklearn.preprocessing import MinMaxScaler
scaler = MinMaxScaler()
provider_stats[['FraudProb_norm', 'NumClaims_norm', 'AvgClaimAmount_norm']] = \
    scaler.fit_transform(
        provider_stats[['Fraud_Prob', 'NumClaims', 'AvgClaimAmount']])
)

provider_stats['RiskScore'] = (
    0.5*provider_stats['FraudProb_norm'] +
    0.3*provider_stats['AvgClaimAmount_norm'] +
    0.2*provider_stats['NumClaims_norm']
)

# 5 Take top 10
top_risk_providers = provider_stats.sort_values(by='RiskScore', \
    ascending=False).head(10)

# 6 Make sure RiskScore is not zero
print(top_risk_providers[['Provider', 'RiskScore']])

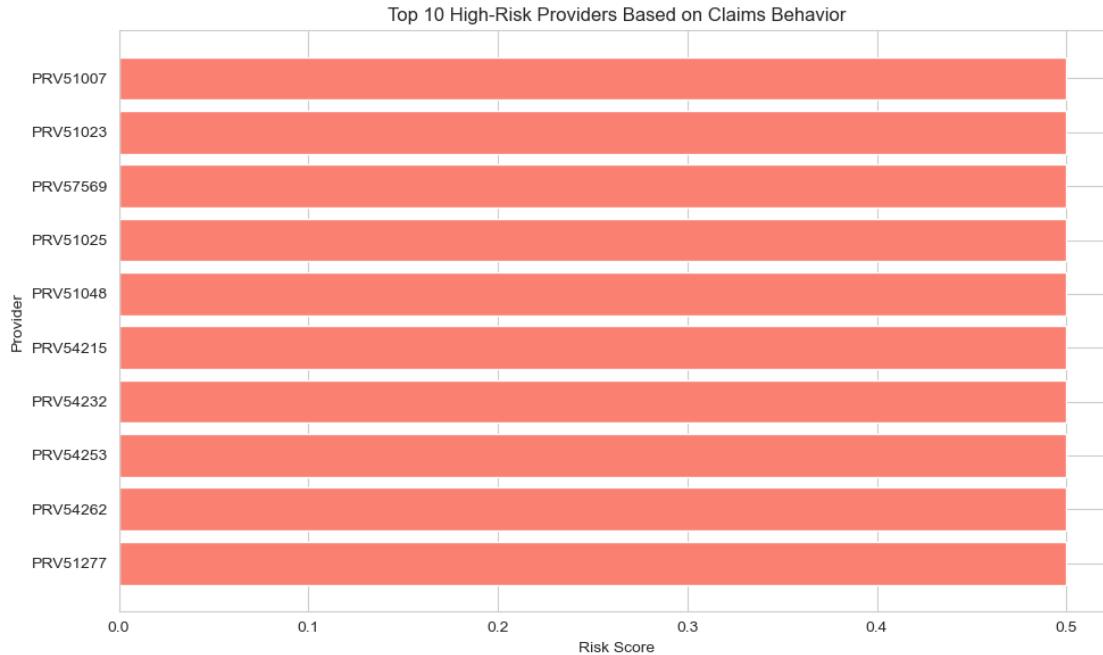
# 7 Plot Horizontal Bar
```

```

plt.figure(figsize=(10,6))
plt.barh(
    top_risk_providers['Provider'],
    top_risk_providers['RiskScore'],
    color='salmon'
)
plt.xlabel("Risk Score")
plt.ylabel("Provider")
plt.title("Top 10 High-Risk Providers Based on Claims Behavior")
plt.gca().invert_yaxis()
plt.tight_layout()
plt.show()

```

	Fraud_Prob	NumClaims	AvgClaimAmount
count	1313.000000	1313.0	1313.0
mean	0.676179	0.0	0.0
std	0.302525	0.0	0.0
min	0.000000	0.0	0.0
25%	0.500000	0.0	0.0
50%	0.718889	0.0	0.0
75%	0.975000	0.0	0.0
max	1.000000	0.0	0.0
	Provider	RiskScore	
1	PRV51007	0.5	
3	PRV51023	0.5	
1294	PRV57569	0.5	
4	PRV51025	0.5	
9	PRV51048	0.5	
627	PRV54215	0.5	
631	PRV54232	0.5	
634	PRV54253	0.5	
636	PRV54262	0.5	
64	PRV51277	0.5	



Insight ##### High-Risk Providers Identified: ##### The analysis identifies a cluster of providers with consistently high fraud risk and financial exposure, making them ideal candidates for prioritized audits rather than random investigation.

What operational changes can be made to minimize fraud?

```
[419]: # Aggregate claim amounts by Fraud status for providers
provider_cost = df_final.
    ↪groupby(['Provider','Fraud'])['InscClaimAmtReimbursed'].sum().
    ↪unstack(fill_value=0)

# Ensure columns exist
for col in [0,1]:
    if col not in provider_cost.columns:
        provider_cost[col] = 0

provider_cost = provider_cost.rename(columns={0:'Non-Fraud', 1:'Fraud'})

# Sort by total claims to pick top 10 risky providers
provider_cost['Total'] = provider_cost['Non-Fraud'] + provider_cost['Fraud']
top_providers = provider_cost.sort_values(by='Total', ascending=False).head(10)
```

```
[420]: import matplotlib.pyplot as plt

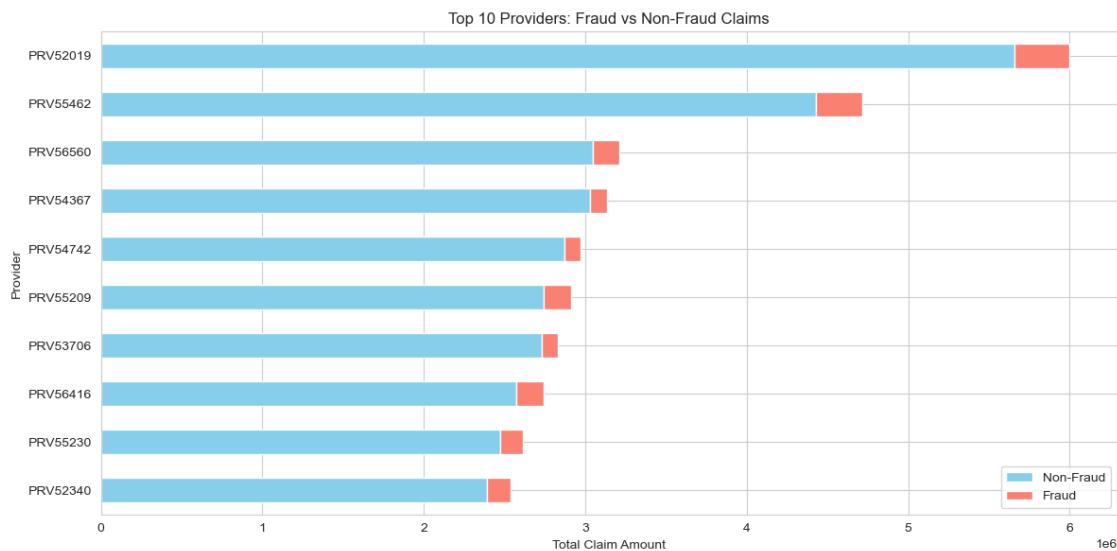
top_providers[['Non-Fraud','Fraud']].plot(
    kind='barh',
```

```

        stacked=True,
        figsize=(12,6),
        color=['skyblue','salmon']
    )

plt.xlabel("Total Claim Amount")
plt.ylabel("Provider")
plt.title("Top 10 Providers: Fraud vs Non-Fraud Claims")
plt.gca().invert_yaxis()
plt.legend(['Non-Fraud','Fraud'])
plt.tight_layout()
plt.show()

```



Insight By identifying top providers responsible for the largest fraudulent claim amounts, operational teams can implement targeted audits, pre-claim verification, and risk-based controls to minimize fraud exposure effectively.

How can fraud-predicted likelihoods be integrated into claim approval workflows?

```
[422]: print(len(X))
print(len(df_final))
```

```
5410
558211
```

```
[182]: df_workflow = df_final.copy()

# Make sure indices match
```

```

df_workflow = df_workflow.reset_index(drop=True)
y_prob_series = pd.Series(rf.predict_proba(X)[:,1], name='Fraud_Prob')
y_prob_series = y_prob_series.reset_index(drop=True)

# Add Fraud Probability
df_workflow['Fraud_Prob'] = y_prob_series

# Check lengths
print(len(df_workflow), len(df_workflow['Fraud_Prob']))

```

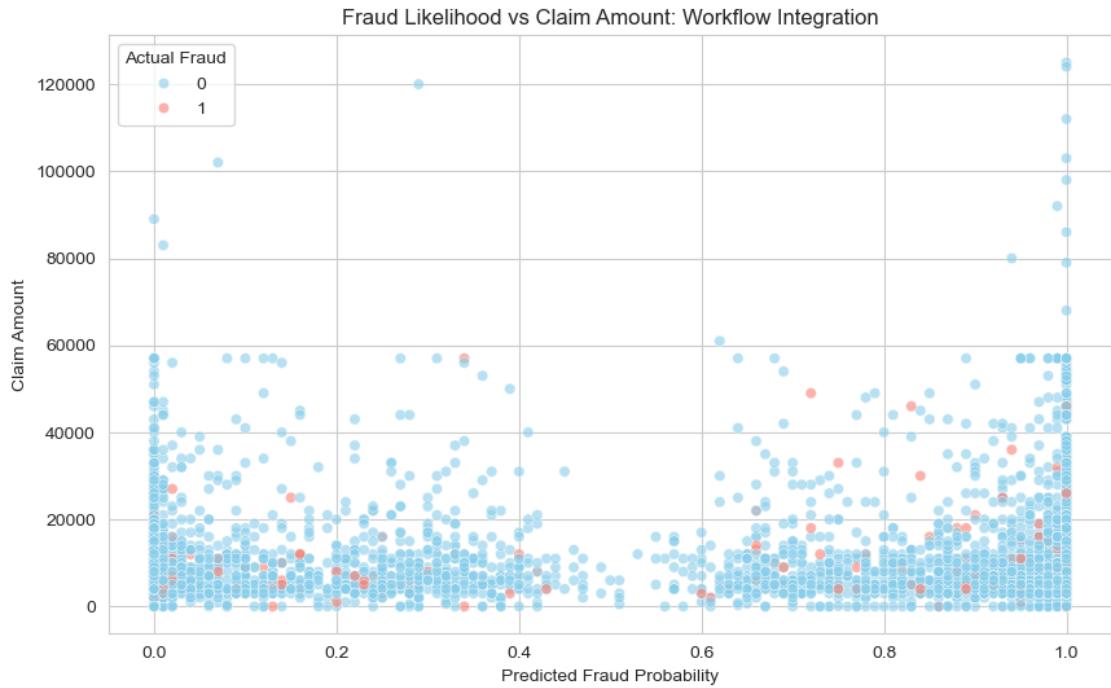
558211 558211

```

[183]: import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(10,6))
sns.scatterplot(
    x='Fraud_Prob',
    y='InscClaimAmtReimbursed',
    data=df_workflow,
    hue='Fraud', # optional actual labels
    palette={0:'skyblue', 1:'salmon'},
    alpha=0.6
)
plt.title("Fraud Likelihood vs Claim Amount: Workflow Integration")
plt.xlabel("Predicted Fraud Probability")
plt.ylabel("Claim Amount")
plt.legend(title="Actual Fraud")
plt.show()

```



```
[ ]: Color = Fraud probability
      Size = Claim amount
      Top-risk claims highlighted
```

Insight

By plotting predicted fraud probability against claim amount, insurers can efficiently prioritize high-risk,

high-cost claims for manual review while automating low-risk claims, reducing financial exposure

and improving operational efficiency

```
[184]: import matplotlib.pyplot as plt
import seaborn as sns

sns.set_style("whitegrid")
plt.rcParams.update({'font.size': 10})

# Example: Fraud Distribution with Actionable Insight
plt.figure(figsize=(7,5))
sns.countplot(x='Fraud', data=provider_stats, palette=['skyblue','salmon'])
plt.title("Fraudulent vs Non-Fraudulent Providers", fontsize=14,
          fontweight='bold')
```

```

plt.xlabel("Fraud (0 = Non-Fraud, 1 = Fraud)")
plt.ylabel("Number of Providers")

# Annotate counts
for i, count in enumerate(provider_stats['Fraud'].value_counts().sort_index()):
    plt.text(i, count + 5, f"{count} providers", ha='center', fontweight='bold')

# Add actionable insight as text on plot
plt.text(0.5, -max(provider_stats['Fraud'].value_counts())*0.2,
         "Actionable Insight: Focus audits on Fraud=1 providers.\n"
         "Majority are low-risk, allocate resources efficiently.",
         ha='center', fontsize=10, color='darkgreen', fontweight='bold')

plt.show()

```

C:\Users\arft\AppData\Local\Temp\ipykernel_20176\3575346942.py:9: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.countplot(x='Fraud', data=provider_stats, palette=['skyblue','salmon'])
```

